

Yandex

MapReduce

Streaming

MapReduce in Python

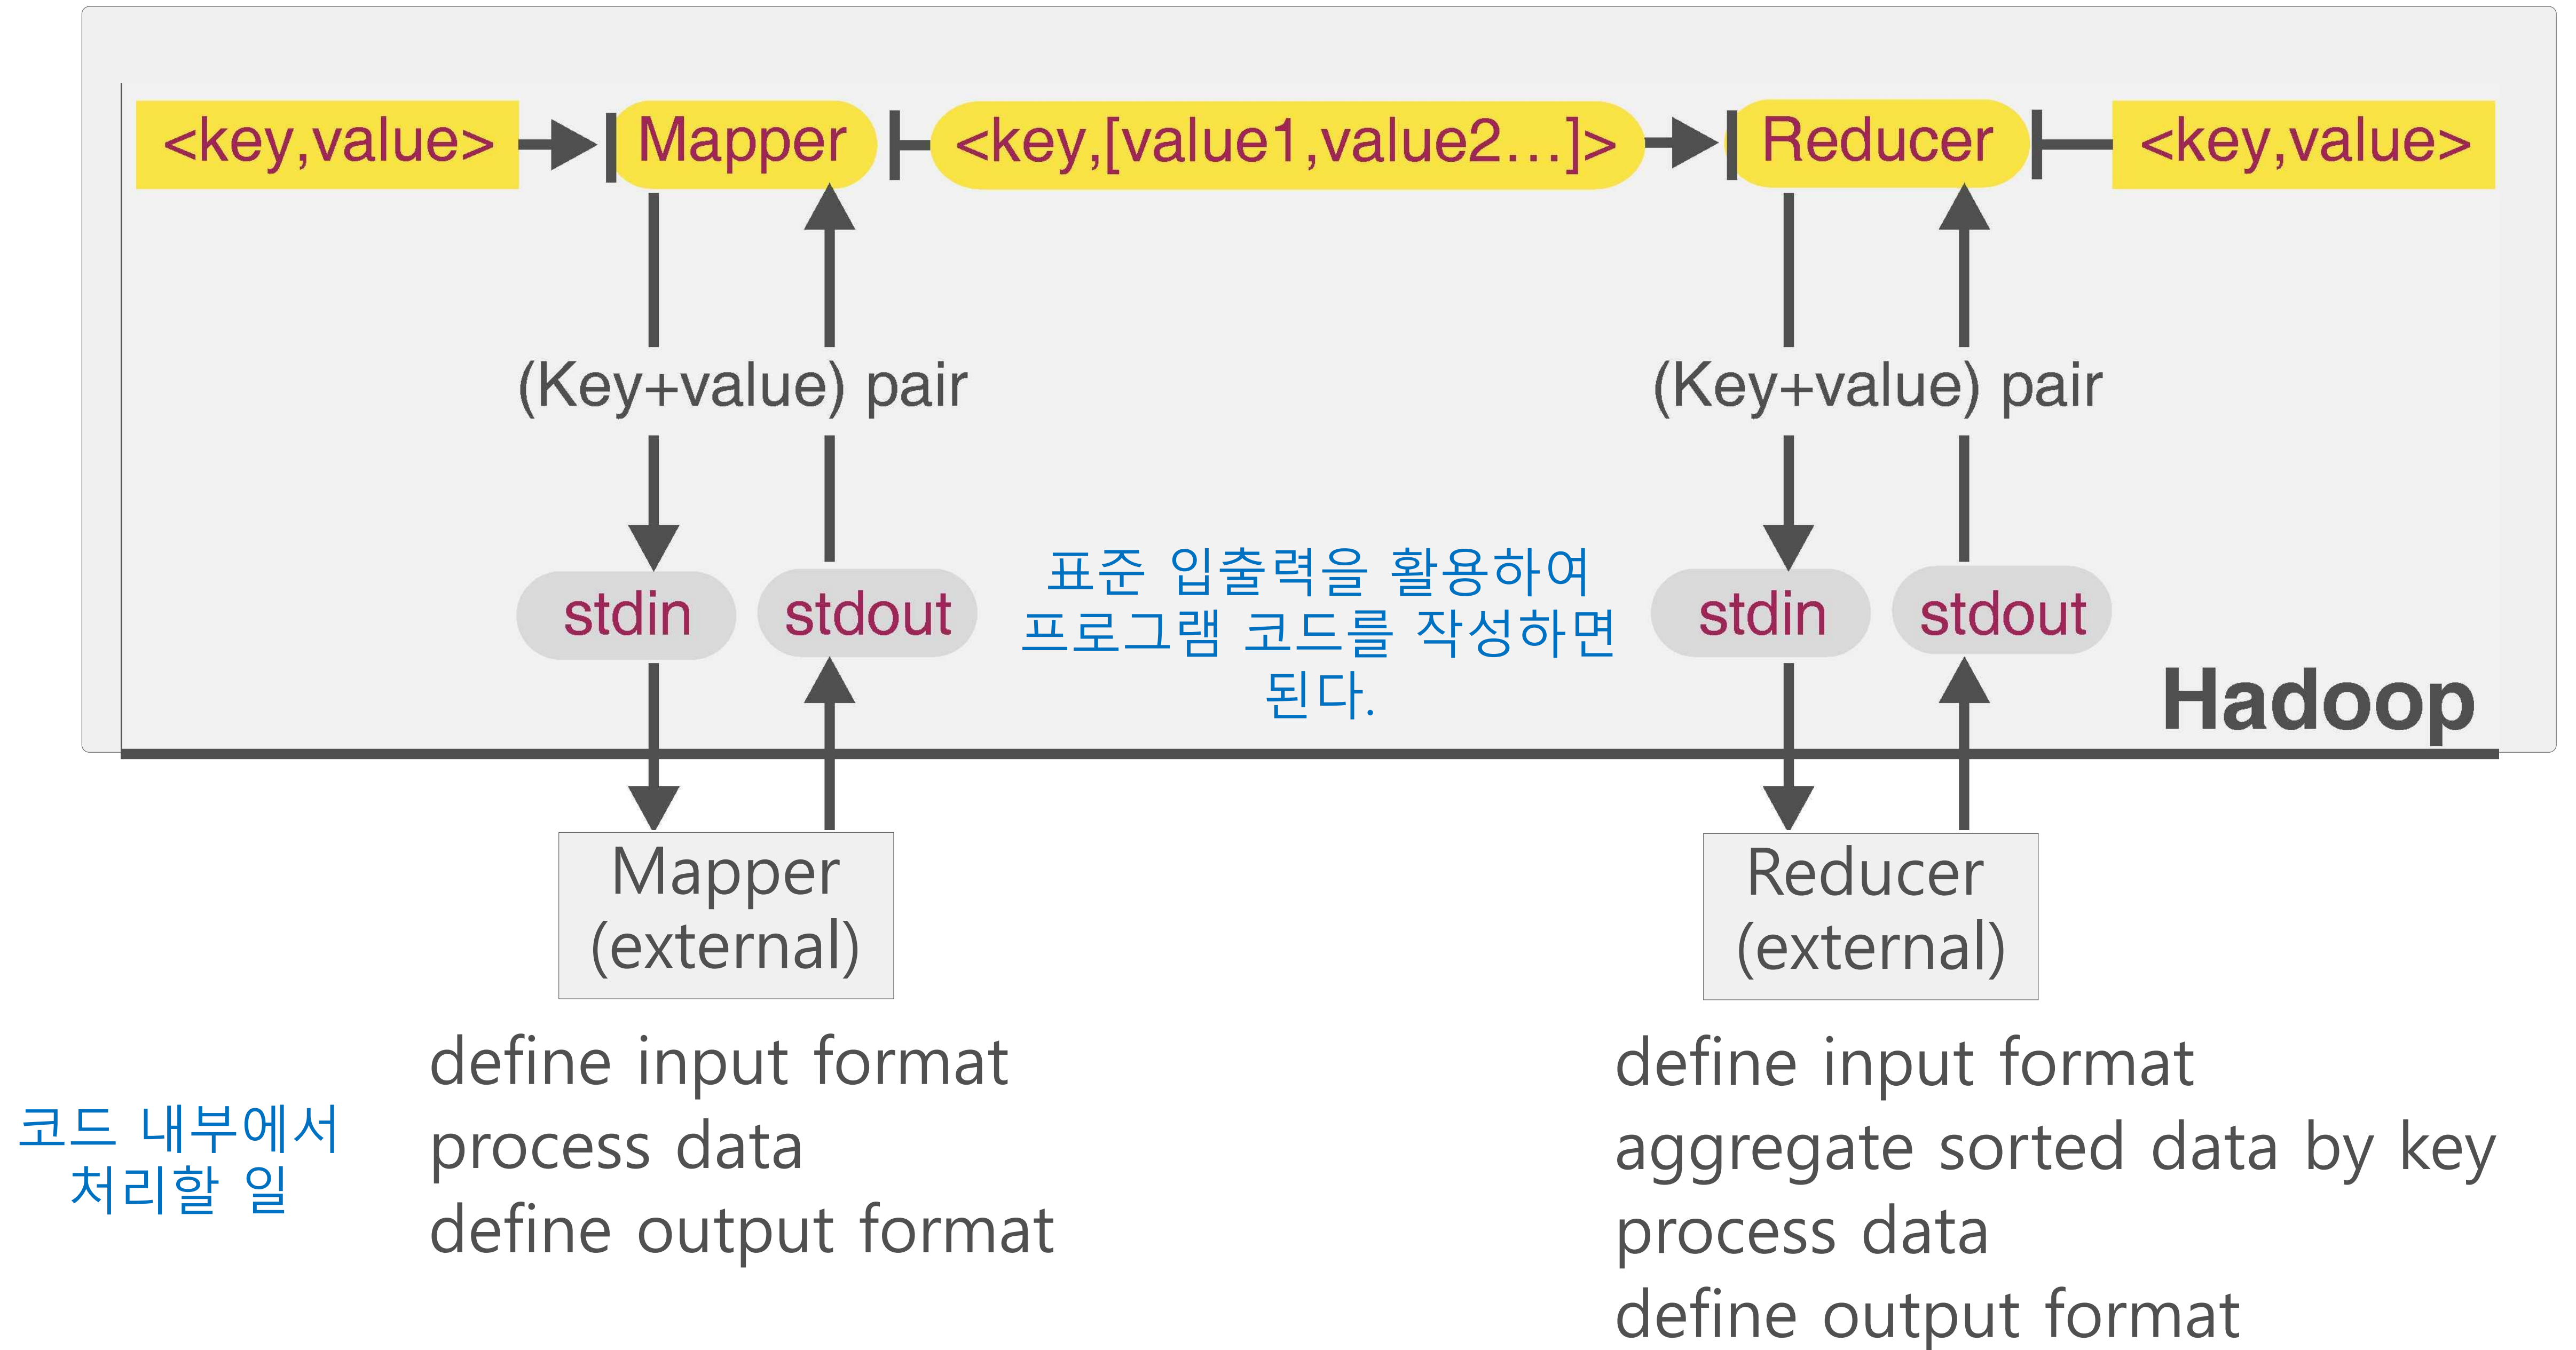
2가지 방법이 있다.



Scipy와 Numpy같은 C로 만든 라이브러리와 호환이 불가능하다.




하둡에서 다양한 언어를 지원

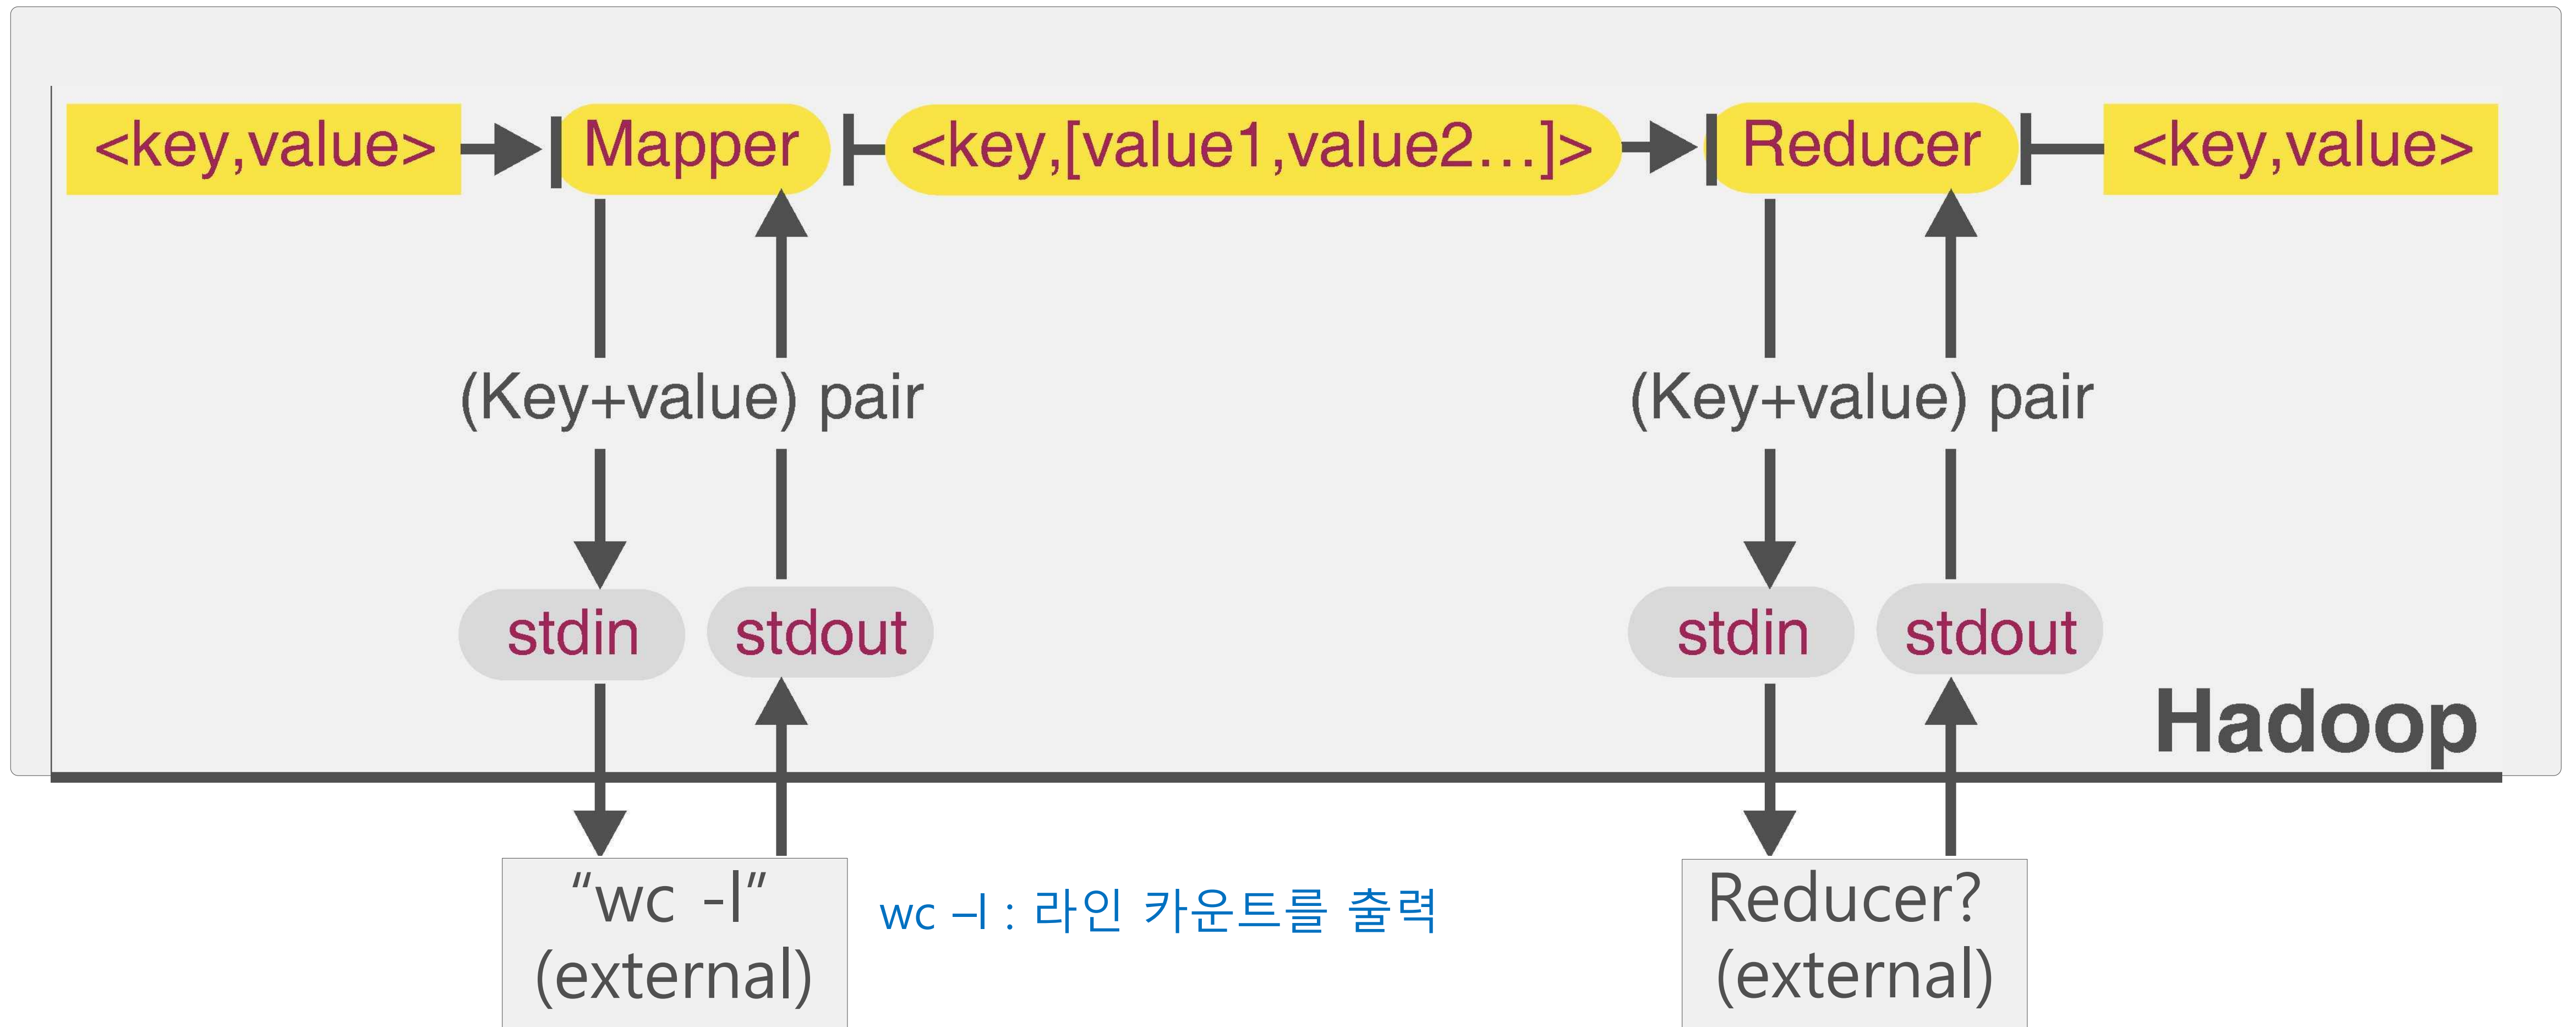


Distributed Shell: **wc**

파일의 총 라인 수를 파악하는 문제 with
Shell Programming

wc는 UNIX의 wordcount 명령어





이 문제는 각 mapper에서 line수를 구한 결과를 단순히 더하면 되기 때문에 reducer가 필요없다.

yarn 실행을 위해서 Hadoop-streaming.jar 파일의 위치를 알아야 한다.
(locate 활용가능)

/opt/cloudera/parcels/CDH-5.9.0-1.cdh5.9.0.p0.23/lib/hadoop-mapreduce/hadoop-streaming.jar



```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
```

```
yarn jar $HADOOP_STREAMING_JAR \  
    -mapper 'wc -l' \  
    -numReduceTasks 0 \ reducer를 실행하지 않기 위해 0으로 지정  
    -input /data/wiki/en_articles \  
    -output wc_mr (폴더명)
```

```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
```

```
yarn jar $HADOOP_STREAMING_JAR \  
    -mapper 'wc -l' \  
    -numReduceTasks 0 \  
    -input /data/wiki/en_articles \  
    -output wc_mr
```

ERROR streaming.StreamJob: Error Launching job : Output directory
hdfs://virtual-master.atp-fivt.org:8020/user/adral/**wc_mr already exists**
Streaming Command Failed!

존재하는 폴더에는 작업을 수행할 수 없다.

```
$ hdfs dfs -rm -r wc_mr
```



```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
```

```
yarn jar $HADOOP_STREAMING_JAR \  
    -mapper 'wc -l' \  
    -numReduceTasks 0 \  
    -input /data/wiki/en_articles \  
    -output wc_mr
```

```
$ hdfs dfs -ls wc_mr
```

 자동으로 mapper가 2개 실행되었다.

```
Found 3 items
```

| | | | | | | | |
|------------|---|-------|-------|---|------------|-------|------------------|
| -rw-r--r-- | 3 | adral | adral | 0 | 2017-03-21 | 14:48 | wc_mr/_SUCCESS |
| -rw-r--r-- | 3 | adral | adral | 6 | 2017-03-21 | 14:48 | wc_mr/part-00000 |
| -rw-r--r-- | 3 | adral | adral | 6 | 2017-03-21 | 14:48 | wc_mr/part-00001 |

```
$ hdfs dfs -text wc_mr/*
```

```
1986
```

```
2114
```

1968 + 2114 = 4100

```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
```

```
yarn jar $HADOOP_STREAMING_JAR \
```

```
-mapper 'wc -l' \
```

awk : 쉘 프로그래밍 언어

```
-reducer "awk '{line_count += $1} END { print line_count
```

awk를 활용한 reducer를
추가한 예제

```
-numReduceTasks 1 \
```

```
-input /data/wiki/en_articles \
```

```
-output wc_mr
```

```
$ hdfs dfs -ls wc_mr_with_reducer
```

```
Found 2 items
```

```
-rw-r--r-- 3 adral adral
```

```
wc_mr_with_reducer/_SUCCESS
```

```
-rw-r--r-- 3 adral adral
```

```
wc_mr_with_reducer/part-00000
```

```
$ hdfs dfs -text wc_mr_with_reducer/*
```

```
4100
```



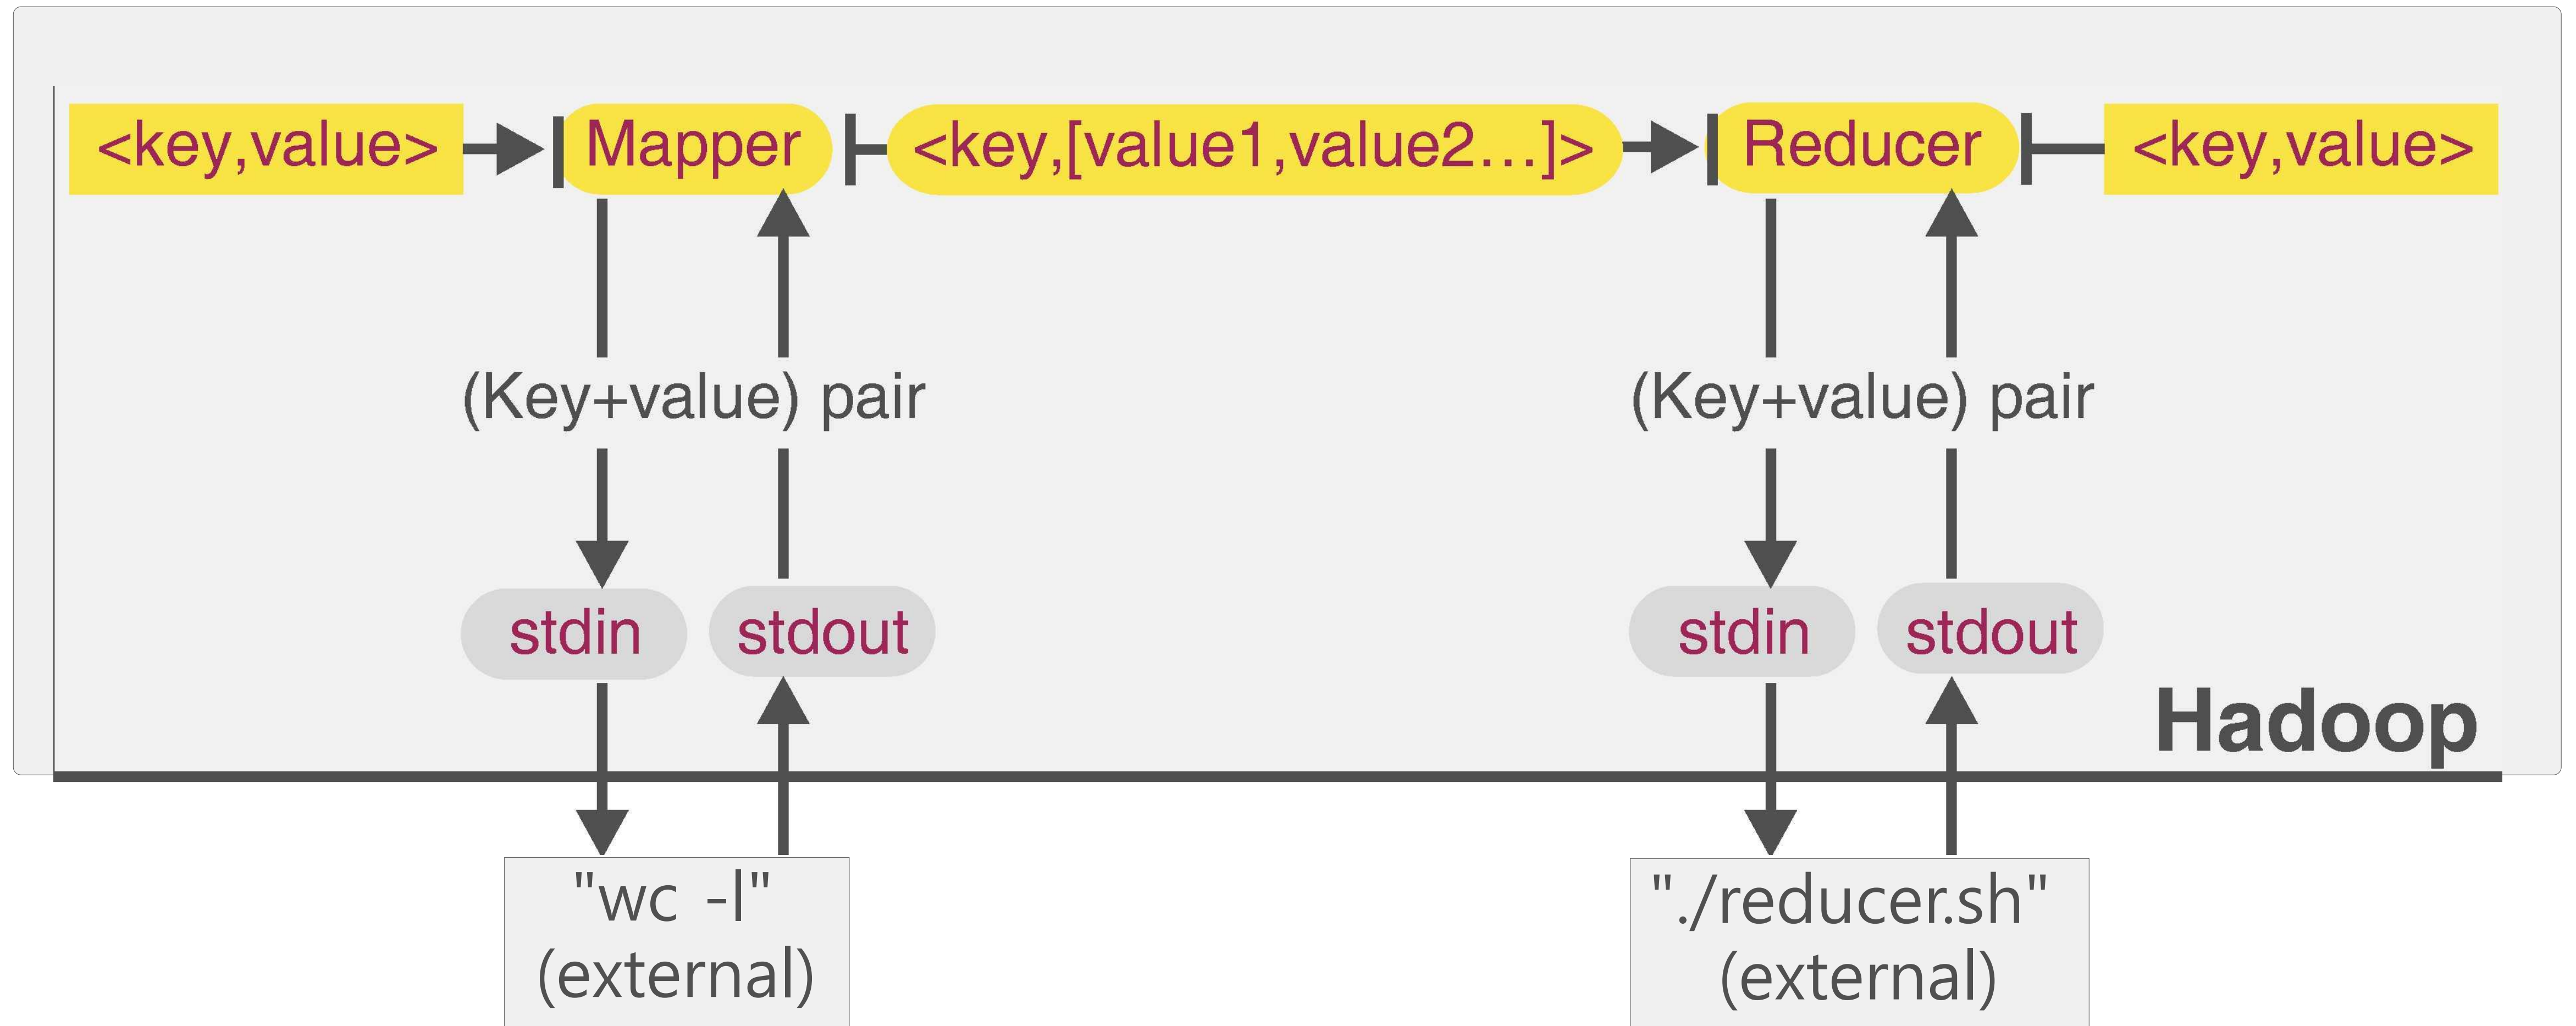
reducer.sh (sh 파일을 직접 만들어 node에 업로드 후 사용하기)

```
#!/usr/bin/env bash
```

```
awk '{line_count += $1} END { print line_count }'
```

```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
```

```
yarn jar $HADOOP_STREAMING_JAR \  
    -mapper 'wc -l' \  
    -reducer './reducer.sh' \  
    -file reducer.sh \  
    -numReduceTasks 1 \  
    -input /data/wiki/en_articles \  
    -output wc_mr_with_reducer
```

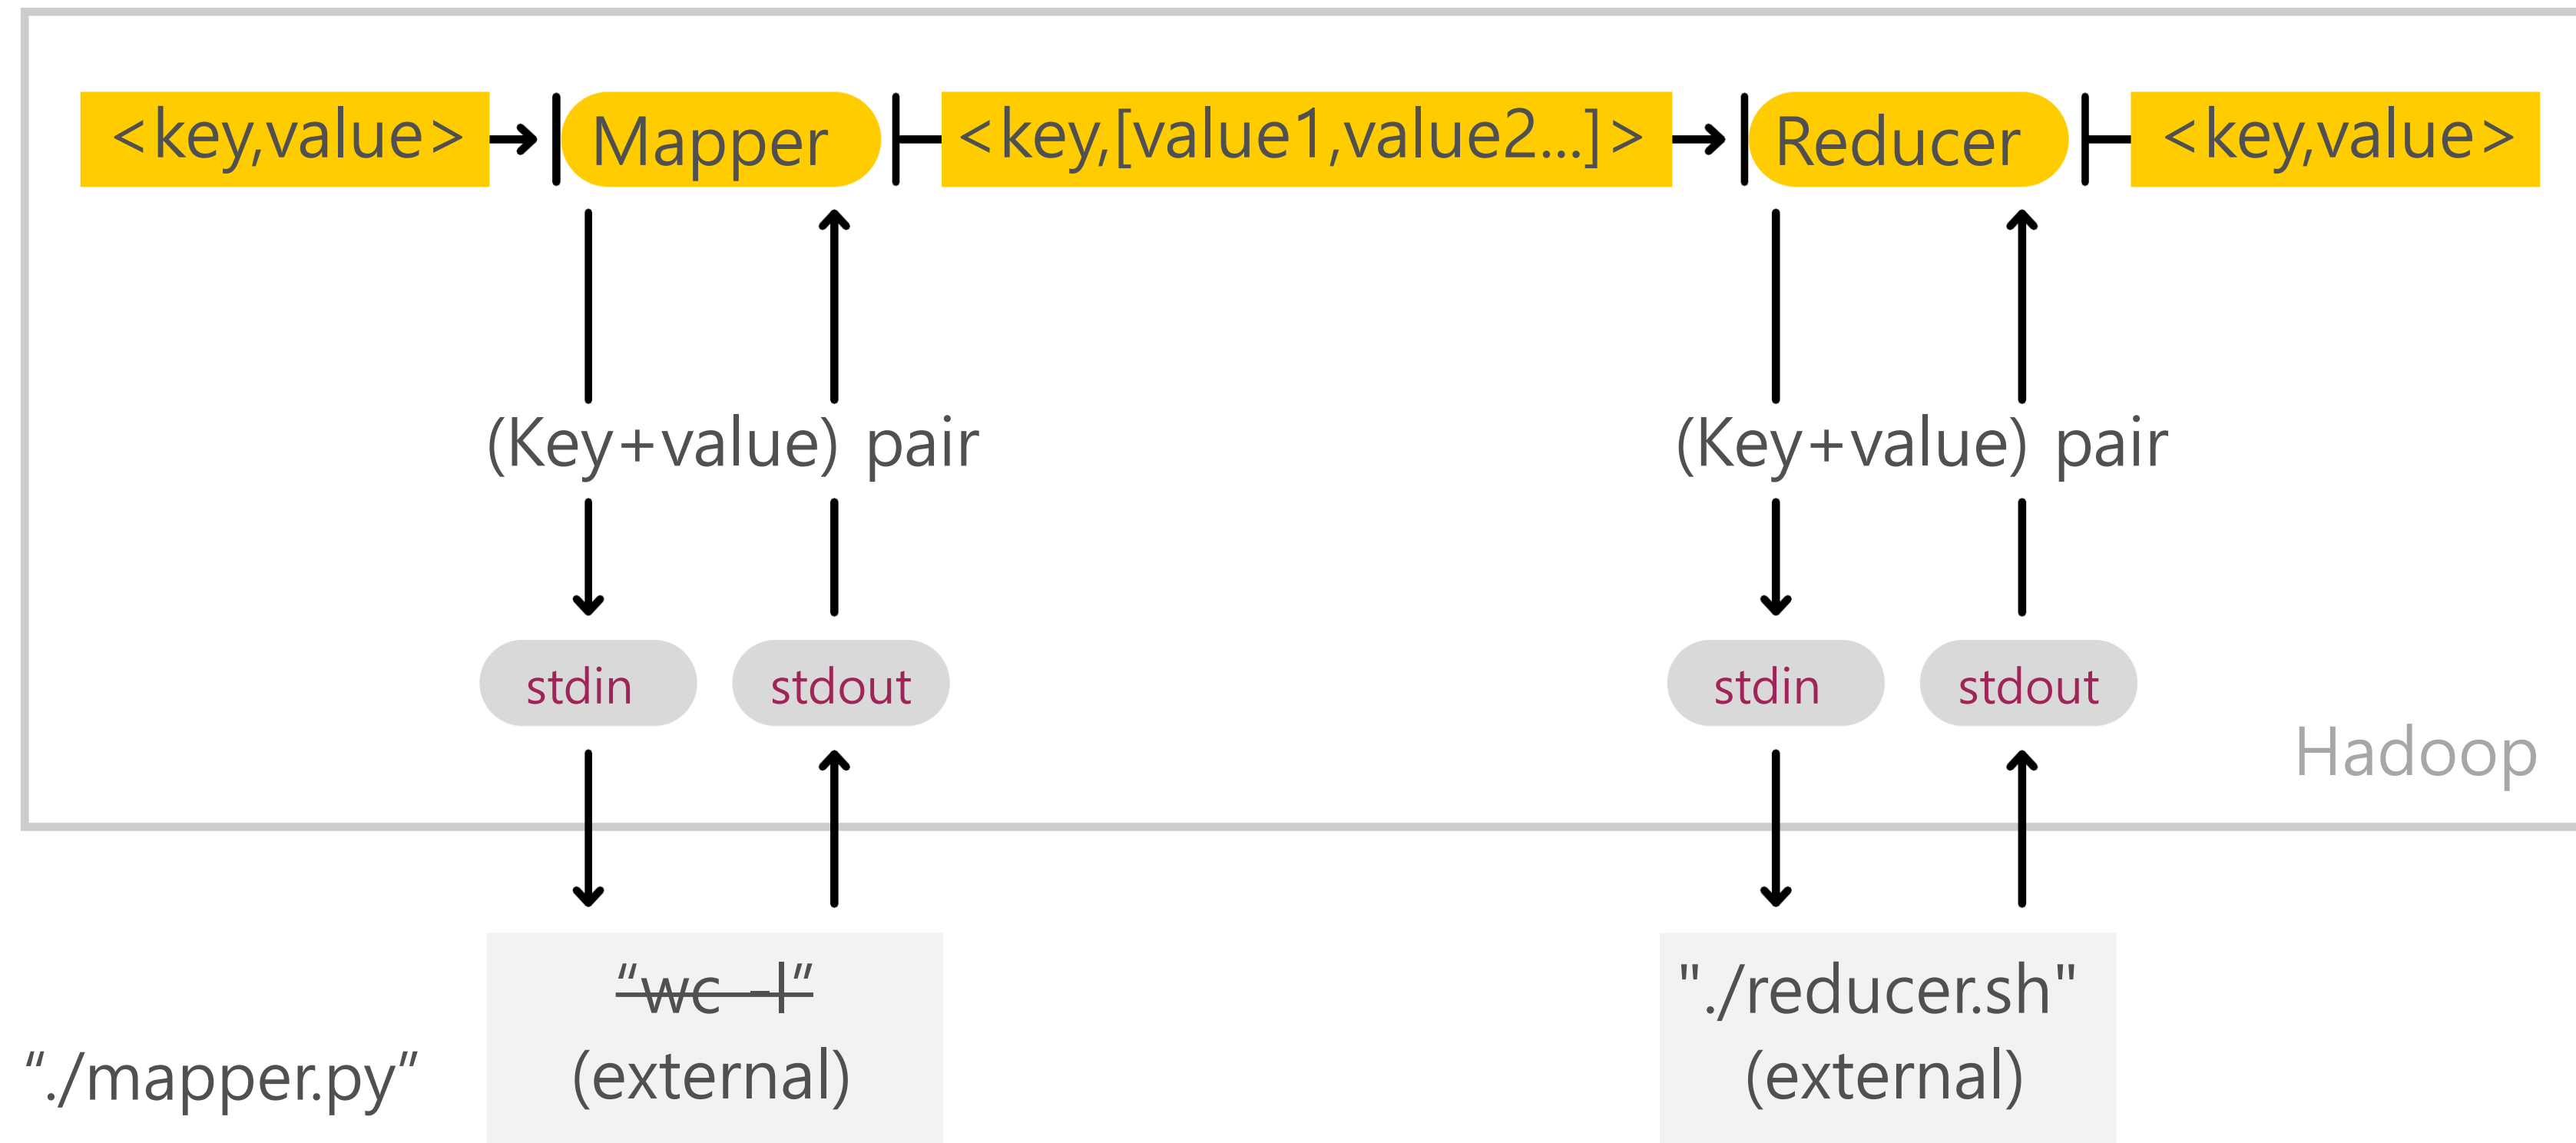


MapReduce

Streaming in Python

line 수를 구하는 문제 풀이

py 파일을 활용



stdin

sys.stdin을 통해 데이터를 라인별로 받을 수 있다.

Mapper (Python): mapper.py

```
from __future__ import print_function  python2 와 python3 호환을 위함
import sys

line_count = 0
for line in sys.stdin:
    pass_count += 1

print(line_count)
```

stdout

출력결과는 print를 통해서 내보낸다.

```
HADOOP_STREAMING_JAR="/path/to/hadoop-streaming.jar"
yarn jar $HADOOP_STREAMING_JAR \
    -files mapper.py, reducer.sh \
    -mapper 'python mapper.py' \
    -reducer './reducer.sh' \
    -numReduceTasks 1 \
    -input /data/wiki/en_articles \
    -output wc_mr_with_reducer
```

The general command line syntax is

bin/hadoop command [**genericOptions**] [commandOptions]

-conf <configuration file>

-D <property=value>

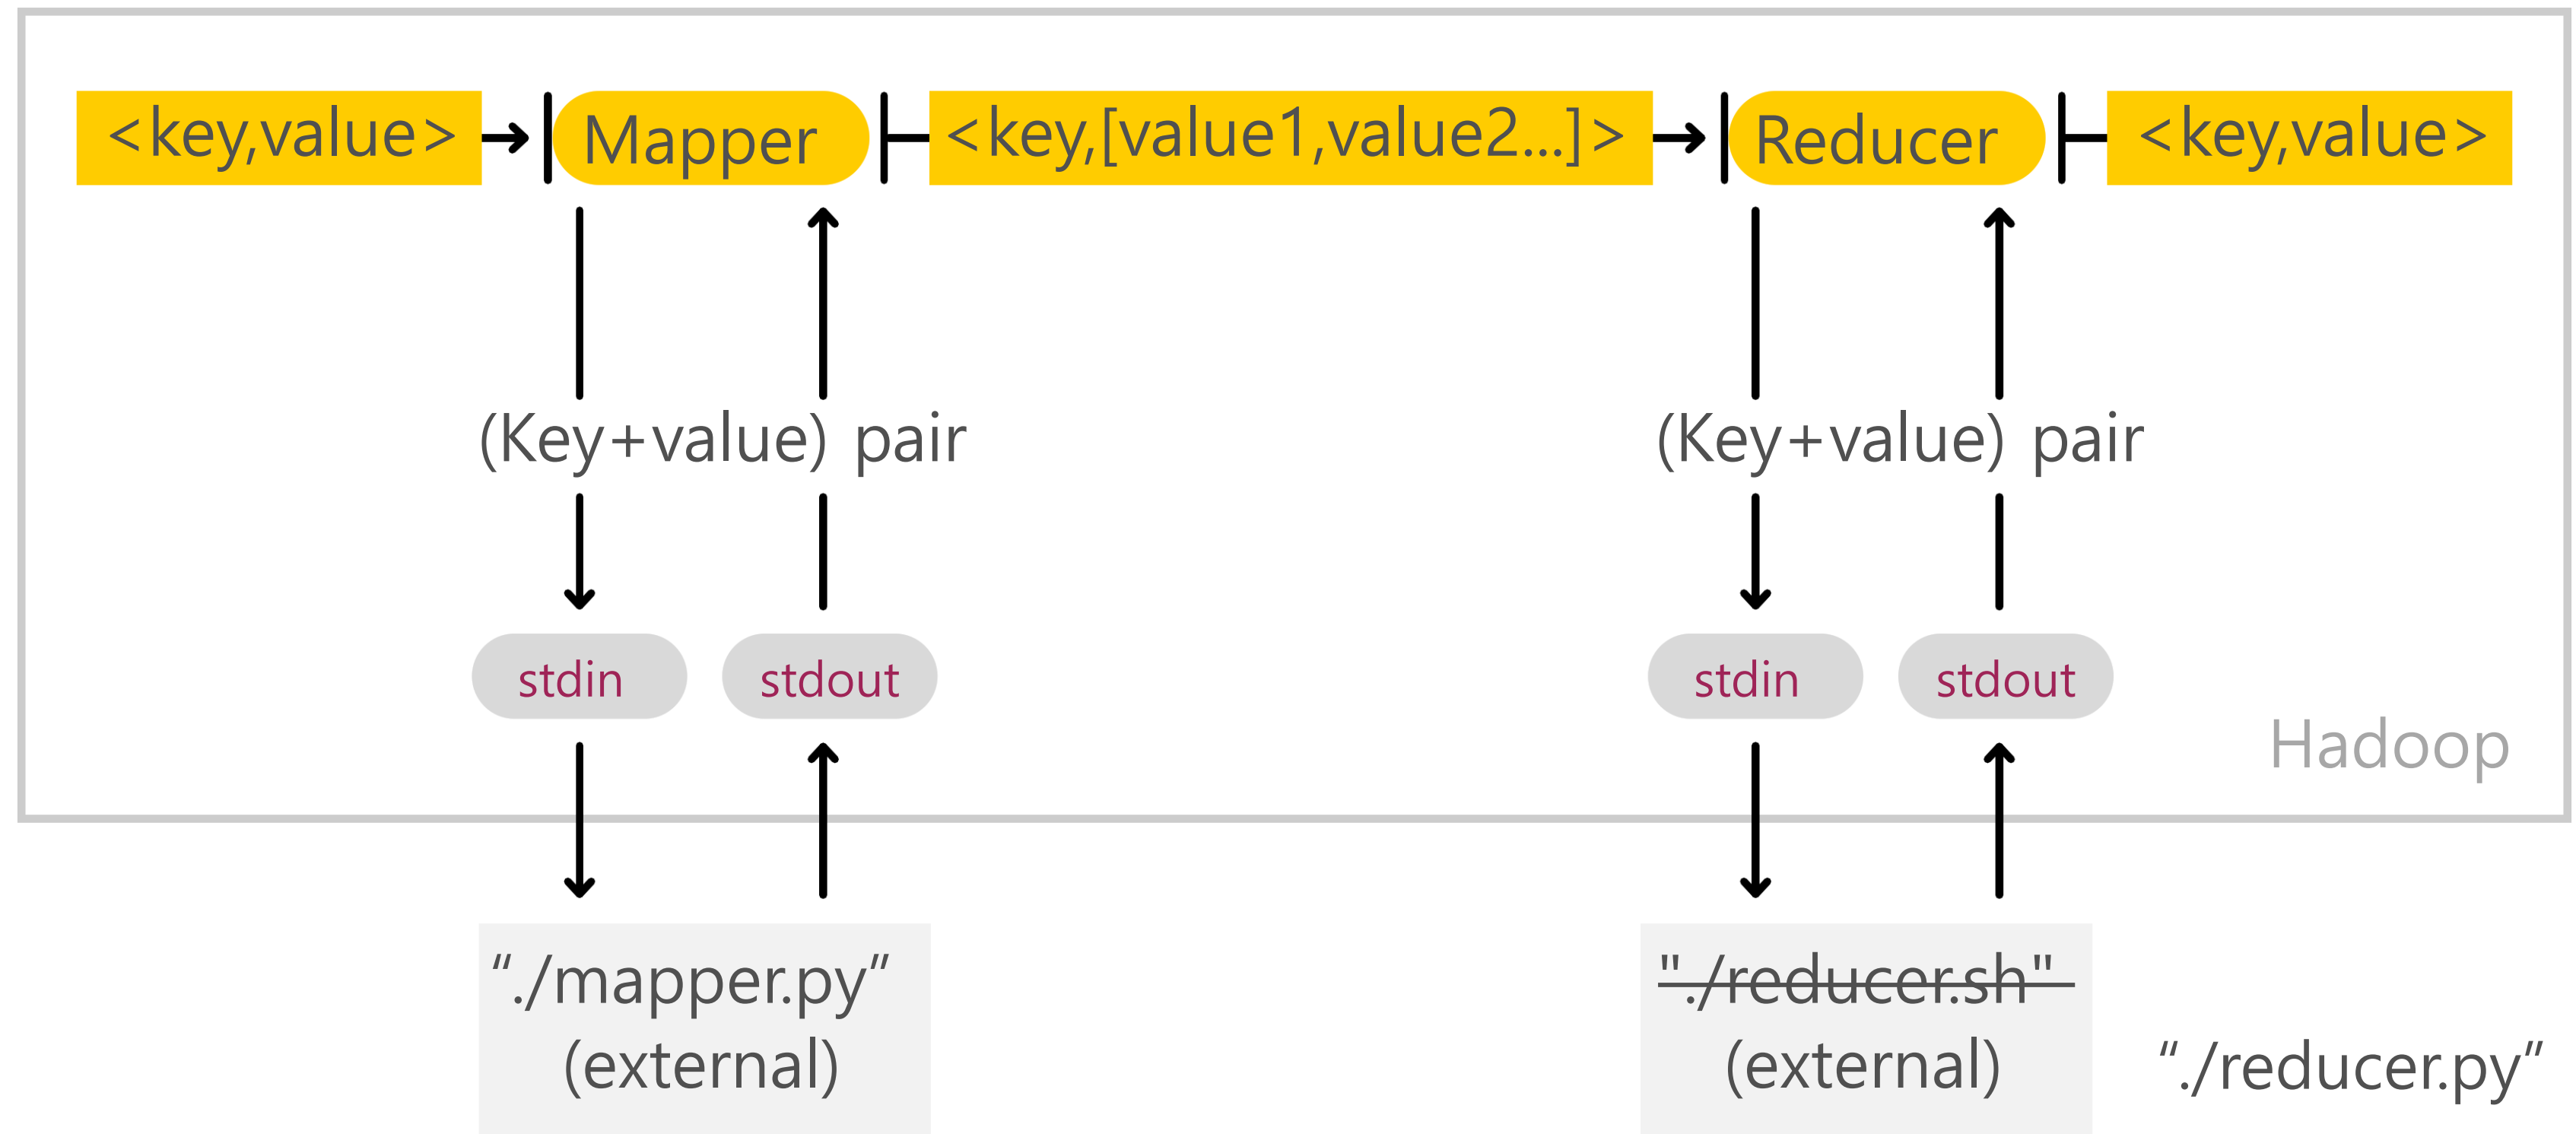
-fs <local|namenode:port>

-jt <local|resourcemanager:port>

-files <comma separated list of files>

-libjars <comma separated list of jars>

-archives <comma separated list of archives>



reducer.py

stdin



```
from __future__ import print_function
import sys

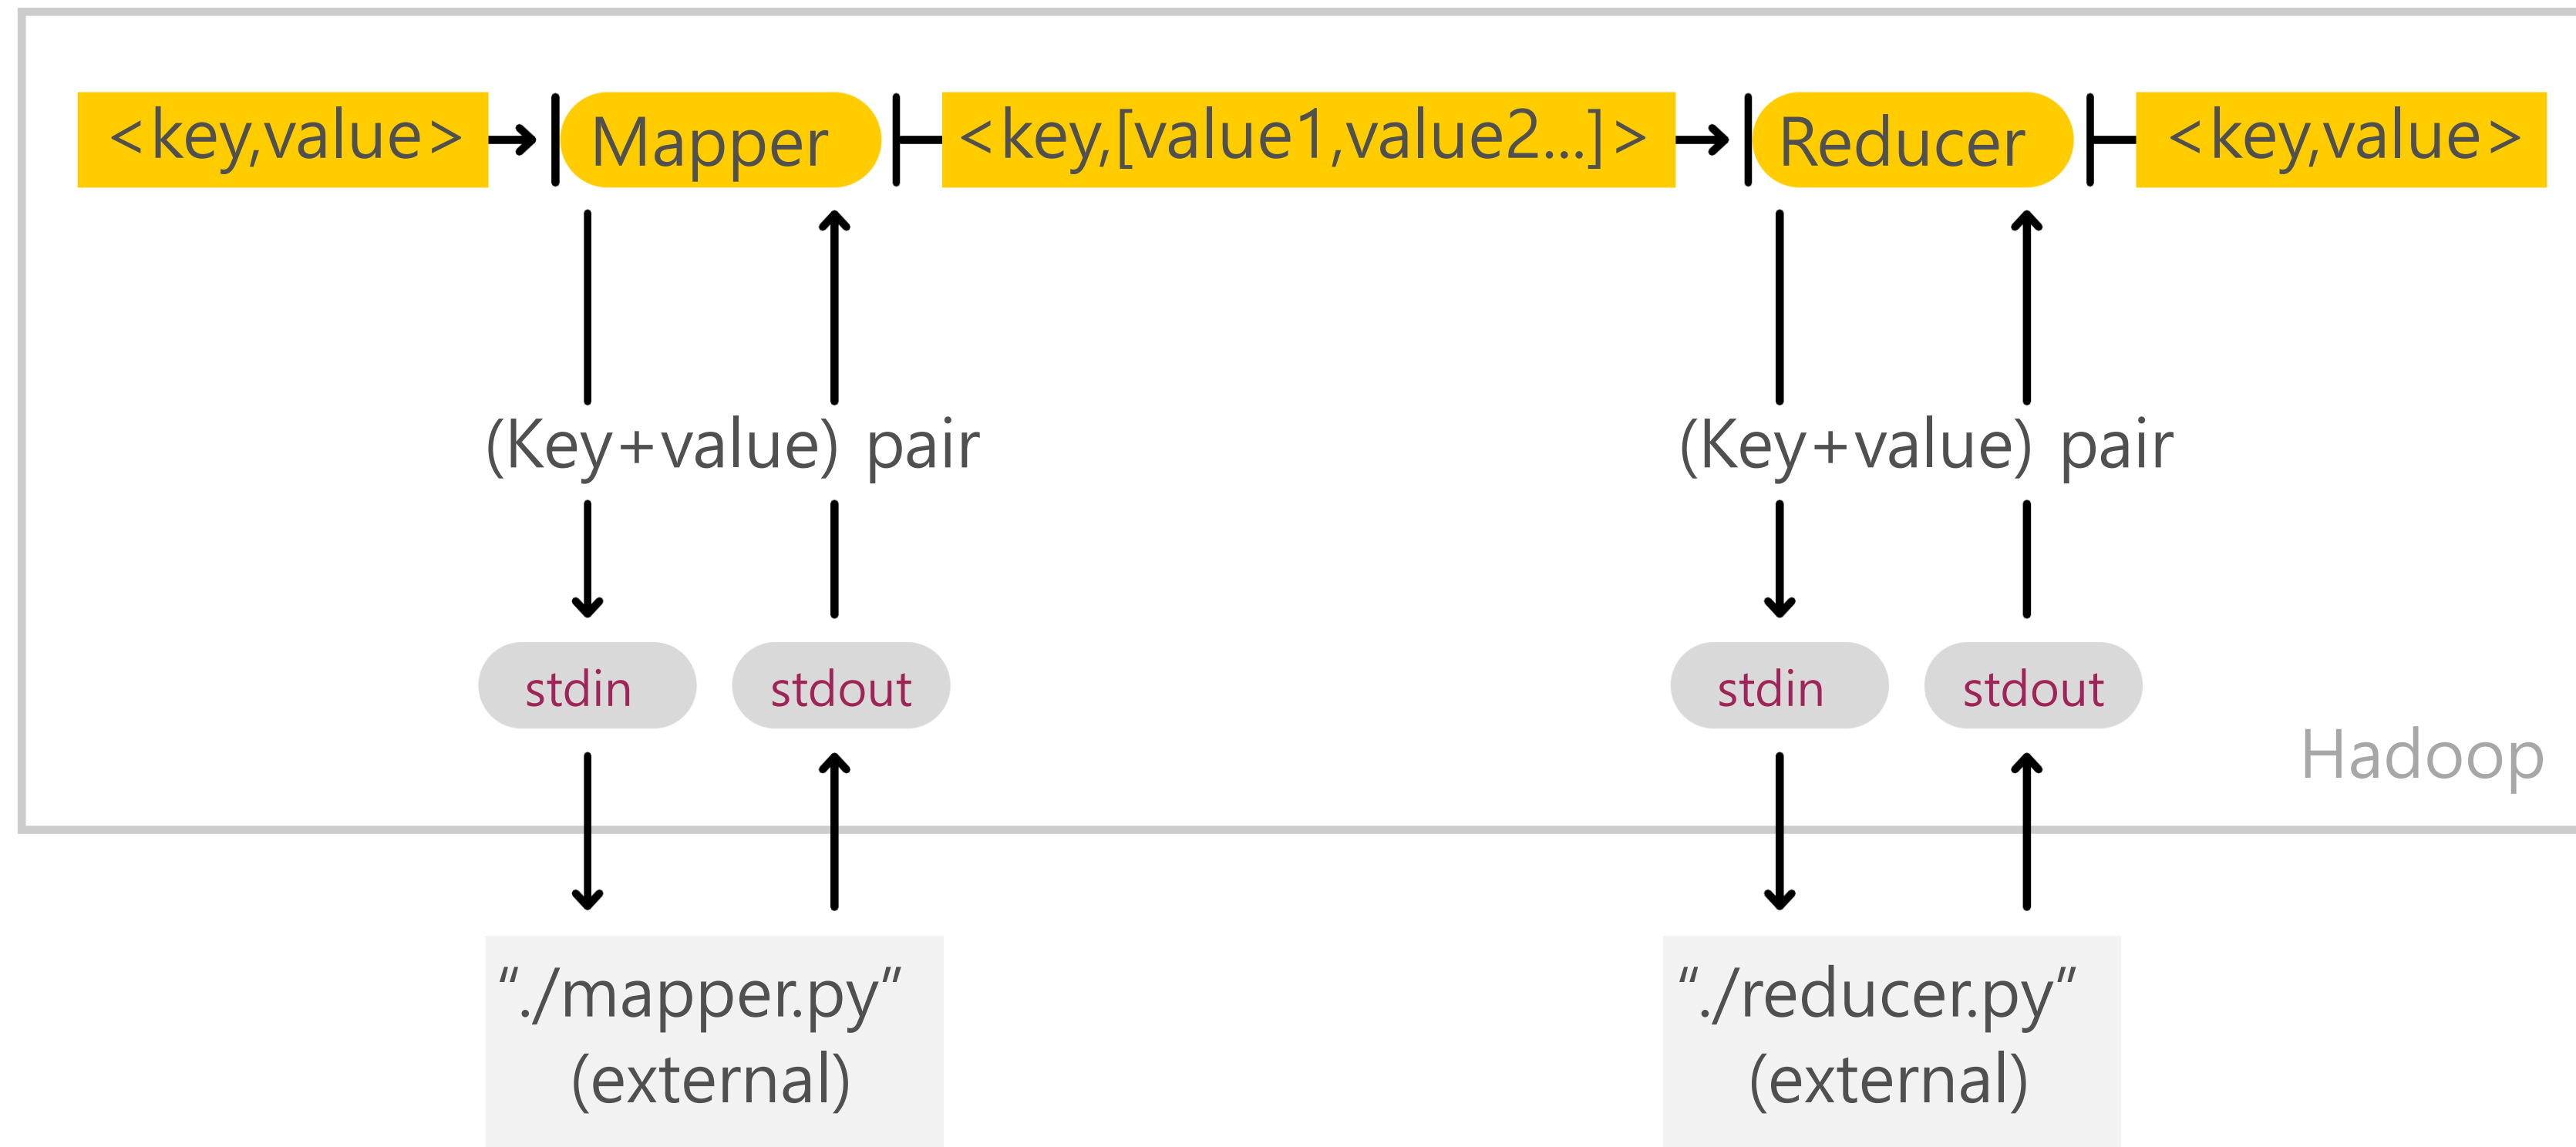
line_count = sum(
    int(value) for value in sys.stdin
)

print(line_count)
```



stdout

mapper에서 온
line 수 결과를 합한다.

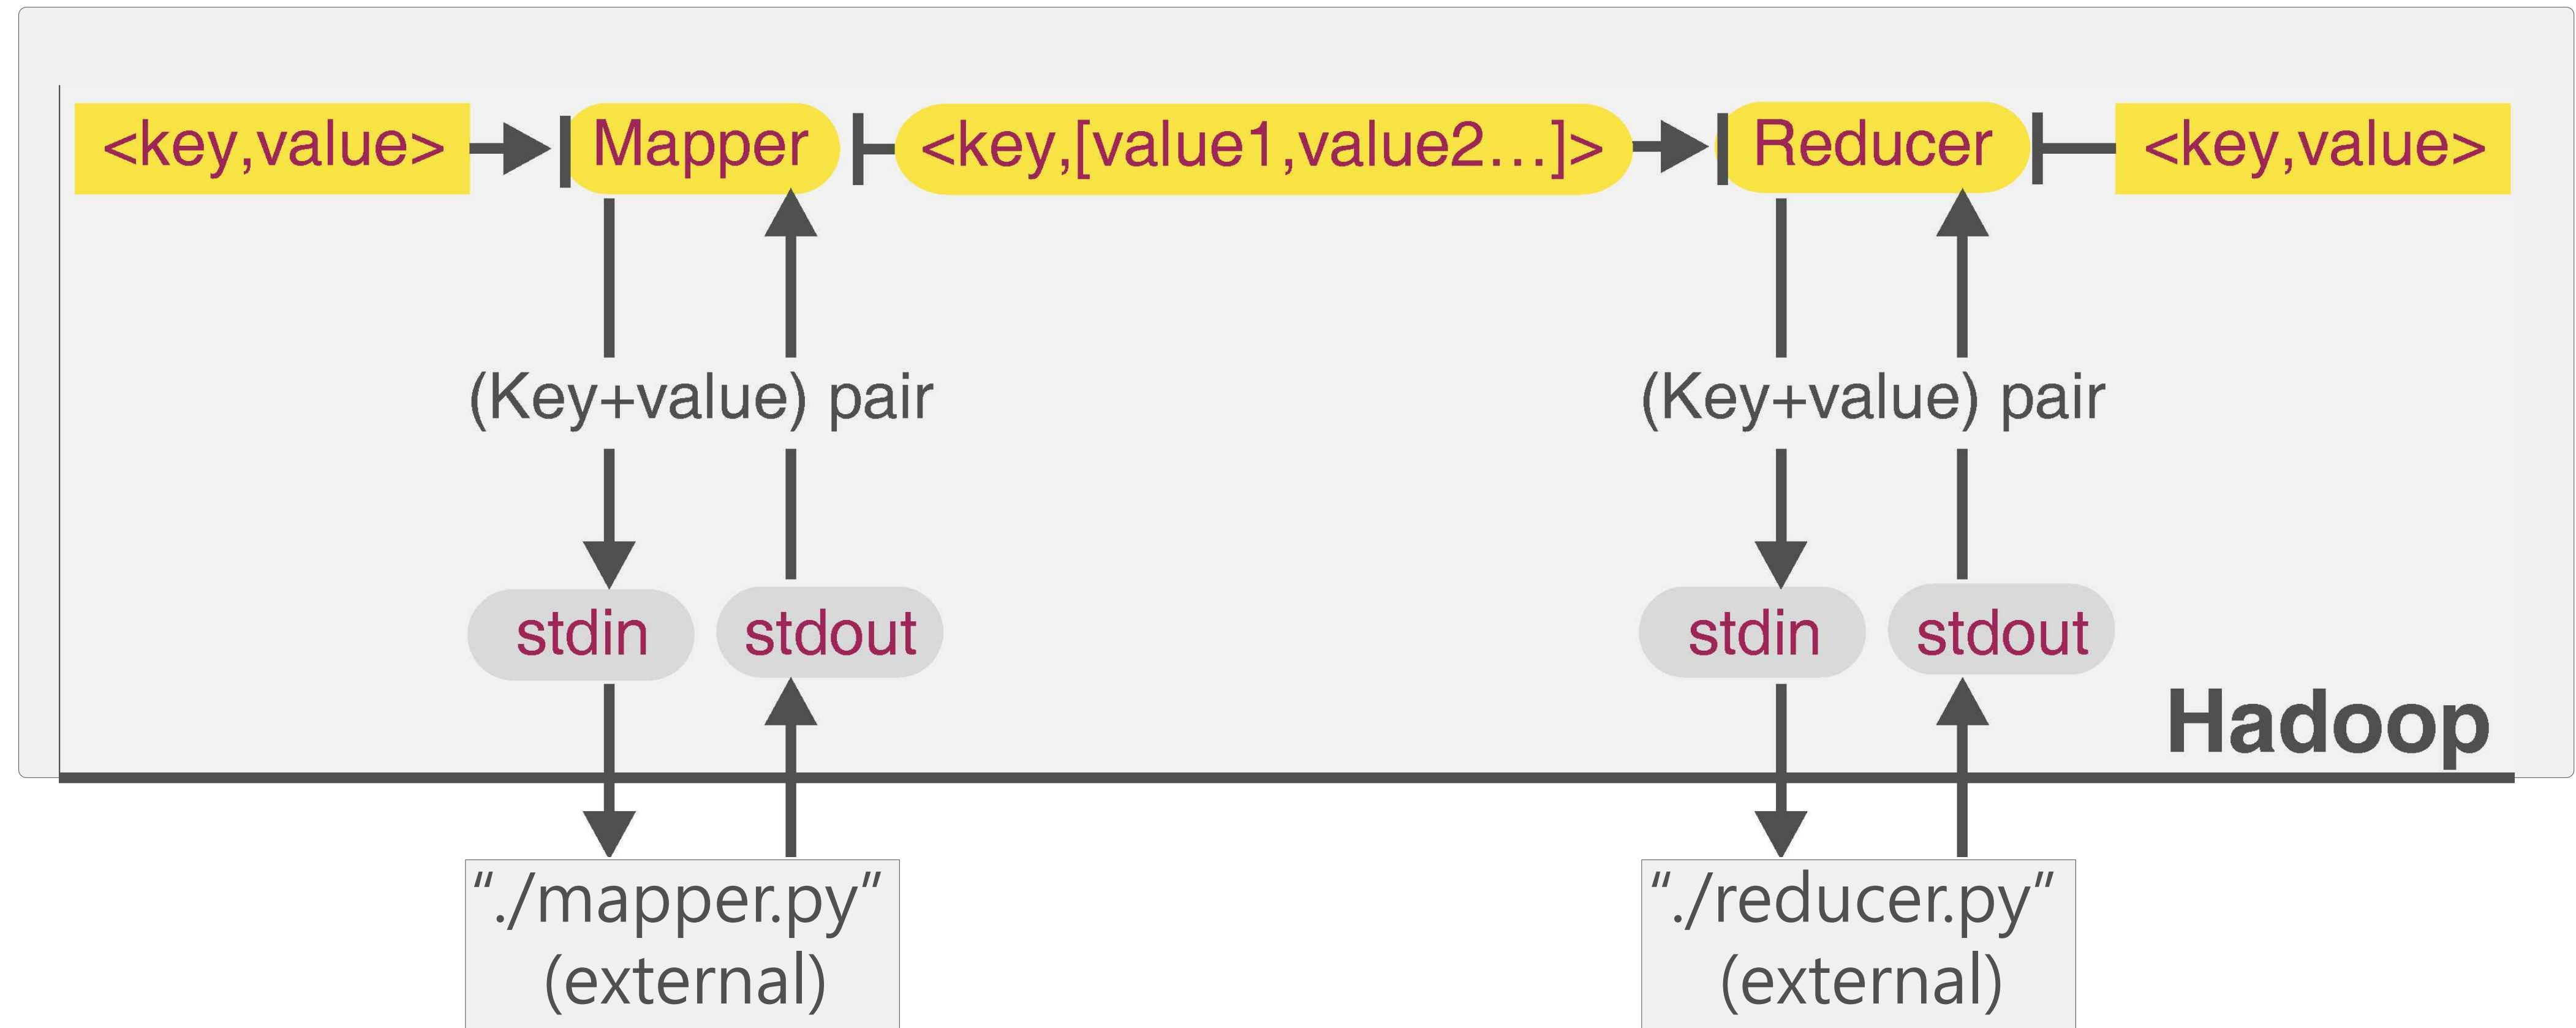



MapReduce

WordCount in Python

Input 파일에서 tab이
key,value를 구분하는 기준

<article id> <tab> <article content>
key value





<article id> <tab> <article content>
key value

```
from __future__ import print_function
import sys
```

```
for line in sys.stdin:
```

```
    article_id, content = line.split("\t", 1)
```

```
        words = content.split()
```

```
            for word in words:
```

```
                print(word, 1, sep="\t")
```

key, value만 분리하기 위한 1

단어별로 나눠서 1을 출력
e.g.) bag 1
game 1

```
yarn jar $HADOOP_STREAMING_JAR \  
    -files mapper.py \  
    -mapper 'python mapper.py' \  
    -numReduceTasks 0 \  
    -input /data/wiki/en_articles \  
    -output word_count
```

```
$ hdfs dfs -text /data/wiki/en_articles/* | head -c 80  
12 <tab> Anarchism           Anarchism is often defined as a political  
philosophy which ...
```

```
$ hdfs dfs -ls -h word_count  
Found 3 items  
-rw-r--r--  3 adral adral 0 2017-03-22 11:40 word_count/_SUCCESS  
-rw-r--r--  3 adral adral 47.8 M 2017-03-22 11:40 word_count/part-00000  
-rw-r--r--  3 adral adral 47.9 M 2017-03-22 11:40 word_count/part-00001
```

```
yarn jar $HADOOP_STREAMING_JAR \  
    -files mapper.py \  
    -mapper 'python mapper.py' \  
    -numReduceTasks 0 \  
    -input /data/wiki/en_articles \  
    -output word_count
```

```
$ hdfs dfs -text /data/wiki/en_articles/*  
| head -c 80
```

~~12~~ <tab> Anarchism Anarchism is
often defined as a political philosophy
which ...

```
$ hdfs dfs -text  
word_count/part-... | head -5
```

| part-00000 | part-00001 |
|------------|-------------|
| Basel 1 | Anarchism 1 |
| Basel 1 | Anarchism 1 |
| (1 | is 1 |
|) 1 | often 1 |
| or 1 | defined 1 |
| ... | ... |

```
yarn jar $HADOOP_STREAMING_JAR \  
    -files mapper.py \  
    -mapper 'python mapper.py' \  
    -numReduceTasks 1 \  
    -input /data/wiki/en_articles \  
    -output word_count
```

```
$ hdfs dfs -text word_count/part-00000 | head
```

```
! 1
```

```
! 1
```

```
! 1
```

```
! 1
```

```
! 1
```

```
...
```

mapper만 실행해도 numReduceTasks를 1이상으로 선언하면
Shuffle&sort가 수행된다.

key별로 정렬이 된 것을 확인가능

mapper.py
수정버전

<article id> <tab> <article content>
key value

```
from __future__ import print_function
import re
import sys
```

```
for line in sys.stdin:
    article_id, content = line.split("\t", 1)
    words = re.split("\W+", content) 특수문자를 지우기 위한 정규식 처리
    for word in words:
        if word:
            print(word, 1, sep="\t")
```

₩W : word가 아닌 것들(!, 공백 등)

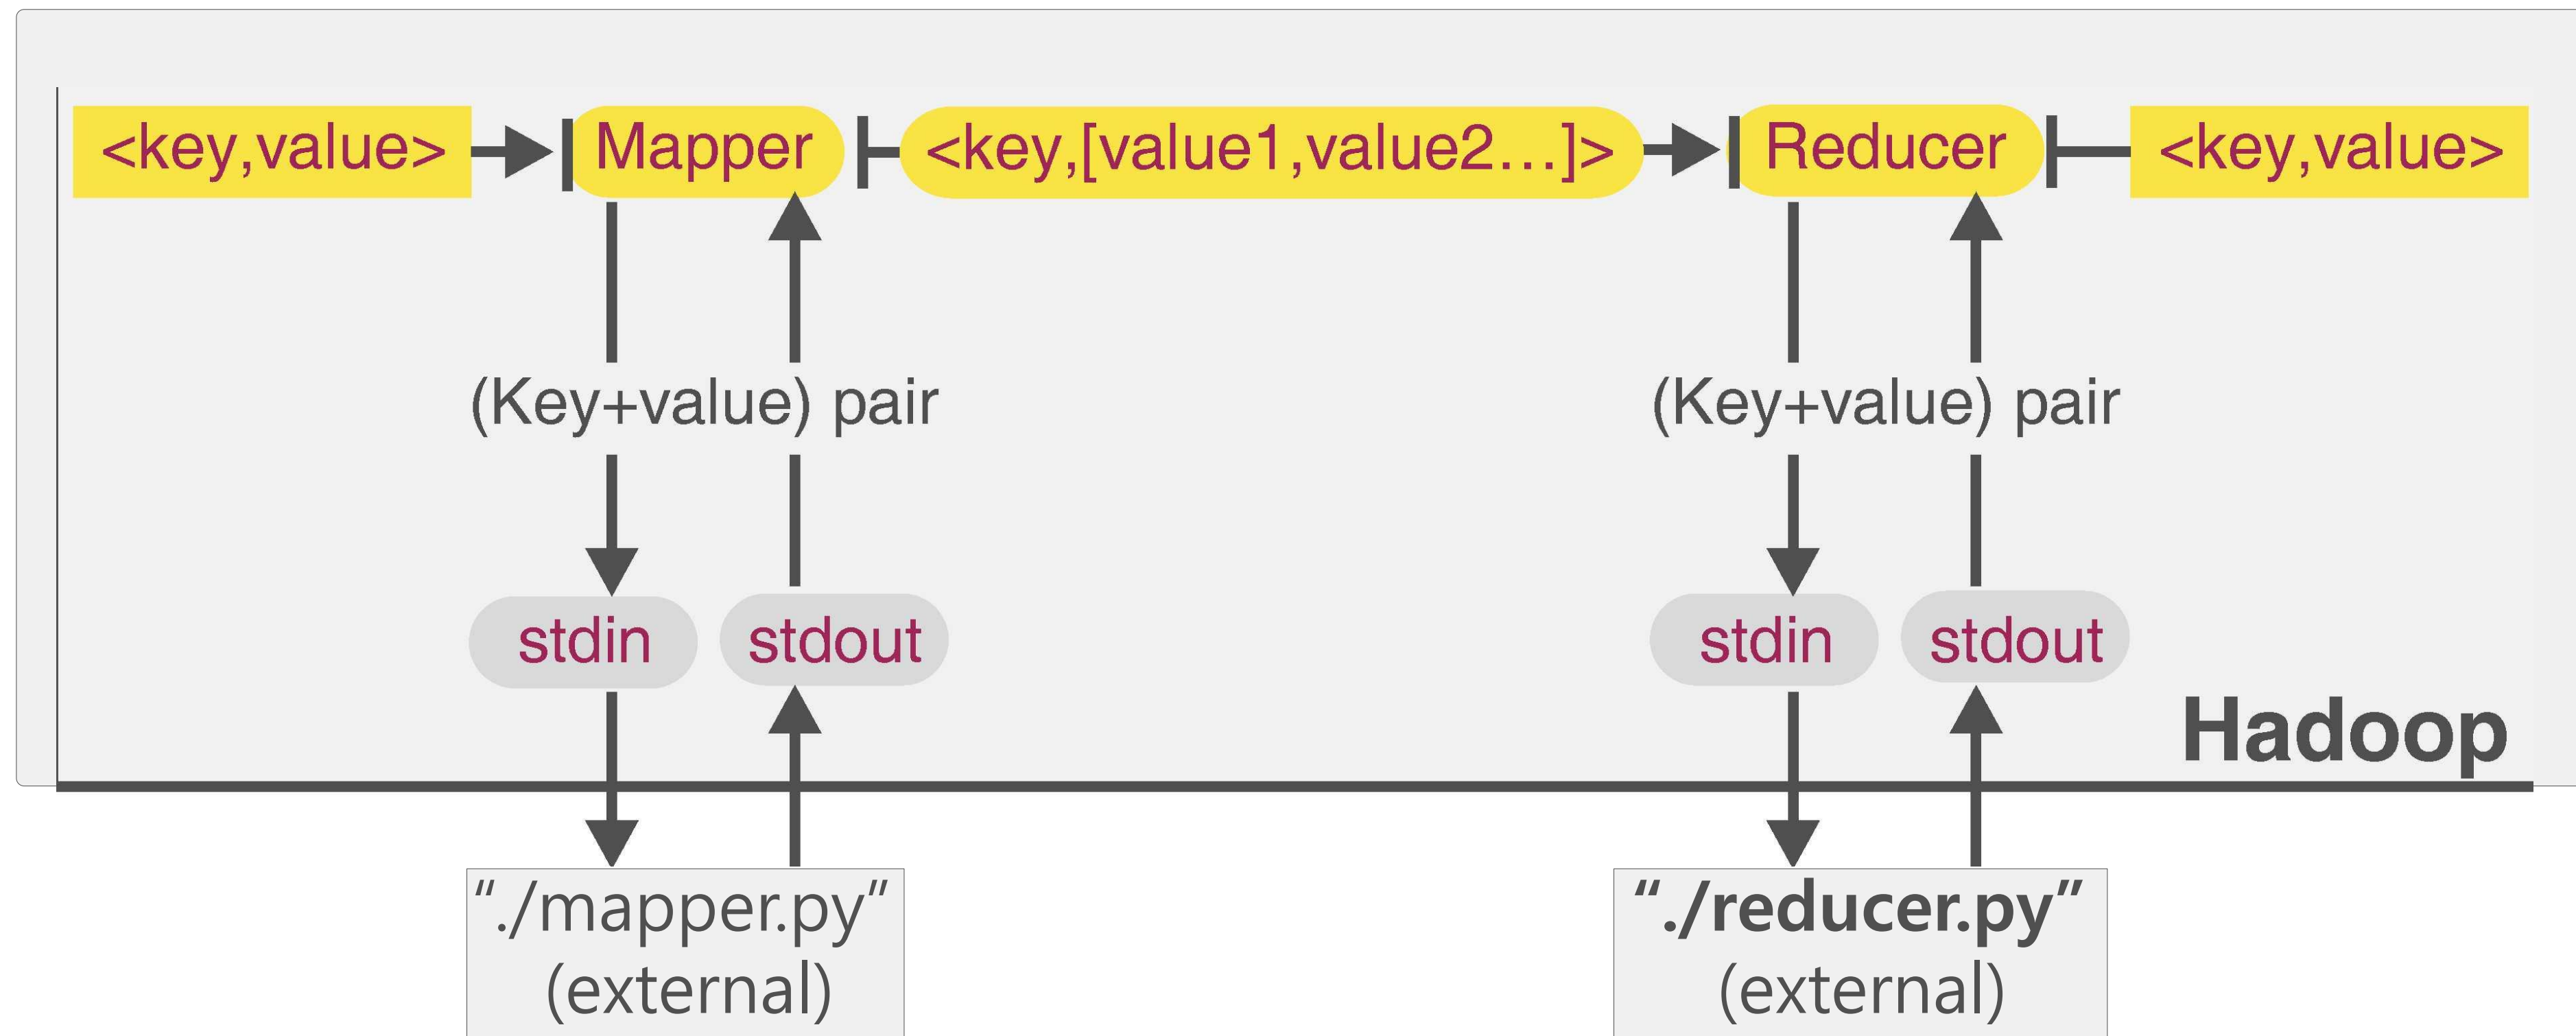

```
yarn jar $HADOOP_STREAMING_JAR \  
        -files mapper.py \  
        -mapper 'python mapper.py' \  
        -numReduceTasks 1 \  
        -input /data/wiki/en_articles \  
        -output word_count
```

특수문자가 제거됨

```
$ hdfs dfs -text word_count/part-00000 | head -4  
0 1  
0 1  
0 1  
0 1
```

```
$ hdfs dfs -tail word_count/part-00000 | tail -4  
zyu1 1  
zyu1 1  
zz 1  
zz 1
```

WordCount




define input format

aggregate sorted data by key

process data

define output format

shuffle & sort된 결과값을
처리해야 함



Reducer (Python): reducer.py

stdout

reducer.py

```
from __future__ import print_function
import sys
```

```
current_word = None 현재 Key를 체크하기 위함
word_count = 0
```


```
for line in sys.stdin:
    word, counts = line.split("\t", 1)
    counts = int(counts)
    if word == current_word:
        word_count += counts
    else:
        if current_word:
            print(current_word, word_count, sep="\t")
        current_word = word
        word_count = counts
```

```
if current_word:
    print(current_word, word_count, sep="\t")
```

마지막 input의 결과를 처리하기 위함

...

| | | |
|---|---------|---|
| | zymosan | 1 |
| | zyu1 | 1 |
| → | zyu1 | 1 |
| | zyu1 | 1 |
| | zyu1 | 1 |
| | zz | 1 |
| | ... | |



```
yarn jar $HADOOP_STREAMING_JAR \  
    -files mapper.py,reducer.py \  
    -mapper 'python mapper.py' \  
    -reducer 'python reducer.py' \  
    -numReduceTasks 1 \  
    -input /data/wiki/en_articles \  
    -output word_count
```

```
$ hdfs dfs -ls -h word_count
```

```
Found 2 items
```

| | | | | | | | |
|------------|---|-------|-------|-------|------------|-------|-----------------------|
| -rw-r--r-- | 3 | adral | adral | 0 | 2017-03-22 | 13:05 | word_count/_SUCCESS |
| -rw-r--r-- | 3 | adral | adral | 3.2 M | 2017-03-22 | 13:05 | word_count/part-00000 |

```
$ hdfs dfs -ls -h word_count
```

```
Found 2 items
```

```
-rw-r--r--      3 adral adral          0 2017-03-22 13:05 word_count/_SUCCESS
-rw-r--r--      3 adral adral    3.2 M 2017-03-22 13:05 word_count/part-00000
```

```
$ hdfs dfs -text word_count/part-00000
```

```
0 14905
```

```
00 844
```

```
000 8186
```

```
...
```

```
zymodemes      1
```

```
zymogen        2
```

```
zymosan        1
```

```
zyu1           4
```

```
zz 2
```

```
...
```

```
zymosan 1
```

```
zyu1 1
```

```
zyu1 1
```

```
zyu1 1
```

```
zyu1 1
```

```
zz 1
```

```
...
```



```
yarn jar $HADOOP_STREAMING_JAR \  
    -files mapper.py,reducer.py \  
    -mapper 'python mapper.py' \  
    -reducer 'python reducer.py' \  
    -input /data/wiki/en_articles \  
    -output word_count
```

numReduceTasks를 지정하지 않으면
자동으로 여러 개를 생성

```
$ hdfs dfs -ls -h word_count
```

```
Found 11 items
```

| | | | | | | | | |
|------------|---|-------|-------|-------|---|------------|-------|-----------------------|
| -rw-r--r-- | 3 | adral | adral | 0 | | 2017-03-22 | 13:19 | word_count/_SUCCESS |
| -rw-r--r-- | 3 | adral | adral | 331.0 | K | 2017-03-22 | 13:18 | word_count/part-00000 |
| -rw-r--r-- | 3 | adral | adral | 332.1 | K | 2017-03-22 | 13:18 | word_count/part-00001 |
| -rw-r--r-- | 3 | adral | adral | 331.7 | K | 2017-03-22 | 13:18 | word_count/part-00002 |
| -rw-r--r-- | 3 | adral | adral | 329.8 | K | 2017-03-22 | 13:18 | word_count/part-00003 |
| -rw-r--r-- | 3 | adral | adral | 326.1 | K | 2017-03-22 | 13:18 | word_count/part-00004 |
| -rw-r--r-- | 3 | adral | adral | 332.2 | K | 2017-03-22 | 13:18 | word_count/part-00005 |
| -rw-r--r-- | 3 | adral | adral | 332.3 | K | 2017-03-22 | 13:18 | word_count/part-00006 |

```
$ hdfs dfs -tail word_count/part-... | tail -5
```

| part-00000 | part-00005 |
|------------|------------|
| ... | ... |
| | zsu 1 |
| zuang 1 | |
| zucchini 5 | |
| | zuchetto 1 |
| zuerst 1 | |
| | zure 1 |
| ... | ... |

각각의 reducer에서는 결과가 정렬되지만 globally sort는 되지 않는다.

TotalOrderPartitioner라는 옵션이 따로 존재

see: TotalOrderPartitioner

Summary

You know **what** MapReduce Streaming **is** and how it works

You know **how to write** MapReduce Bash and Python
Streaming applications

You should be able **to solve** WordCount or similar problems
in MapReduce in Python **by yourself**

BigDATAteam