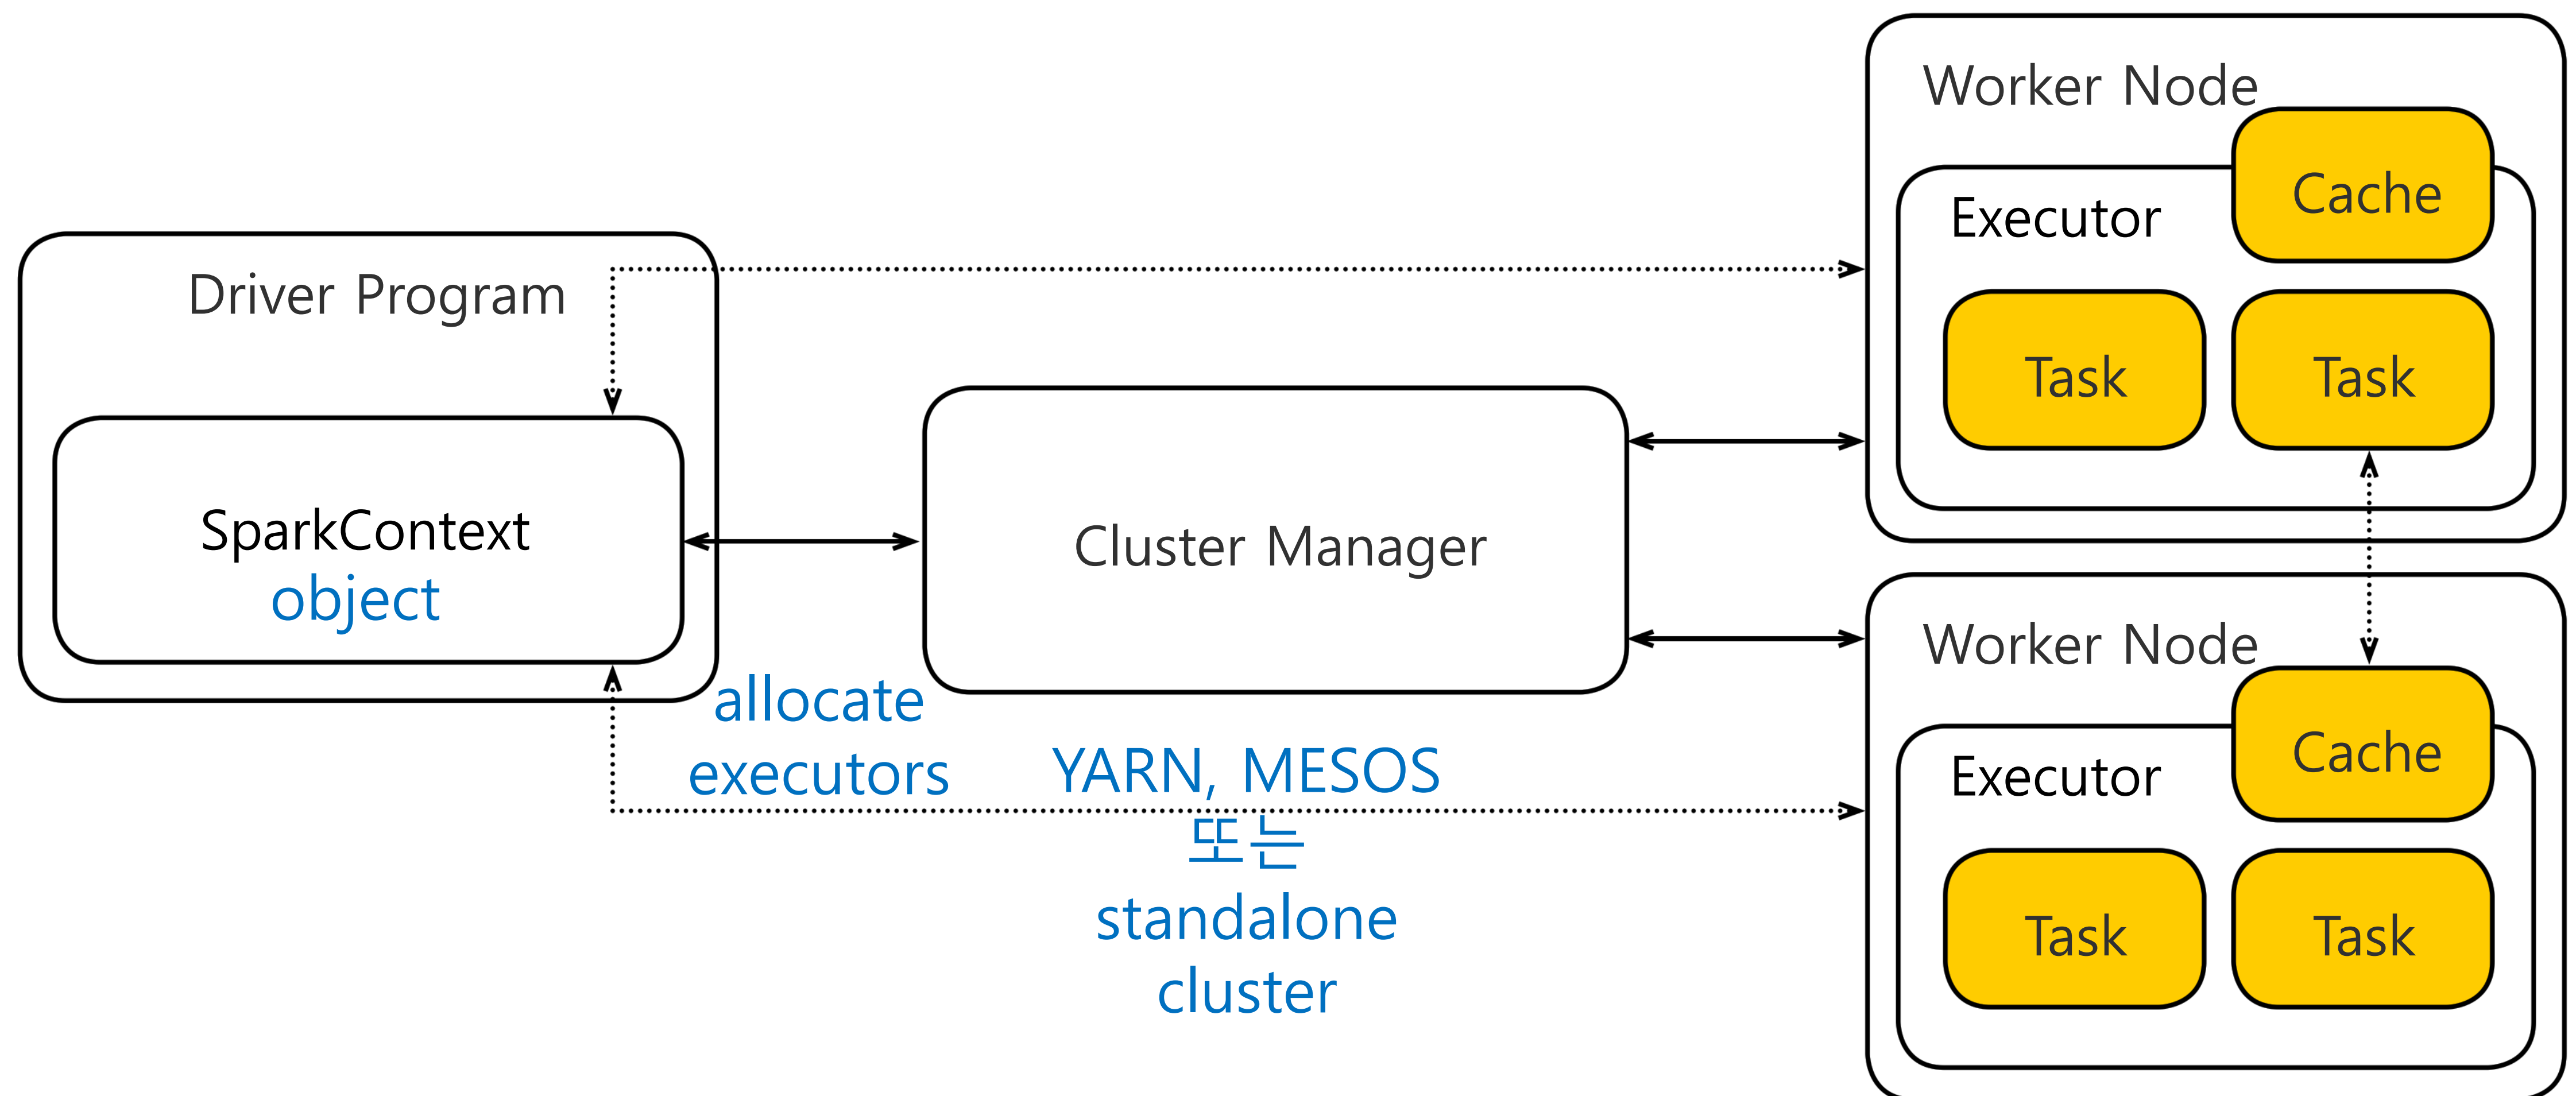


Yandex

Execution & scheduling

SparkContext

- » Tells your application how to access a cluster
- » Coordinates processes on the cluster to run your application



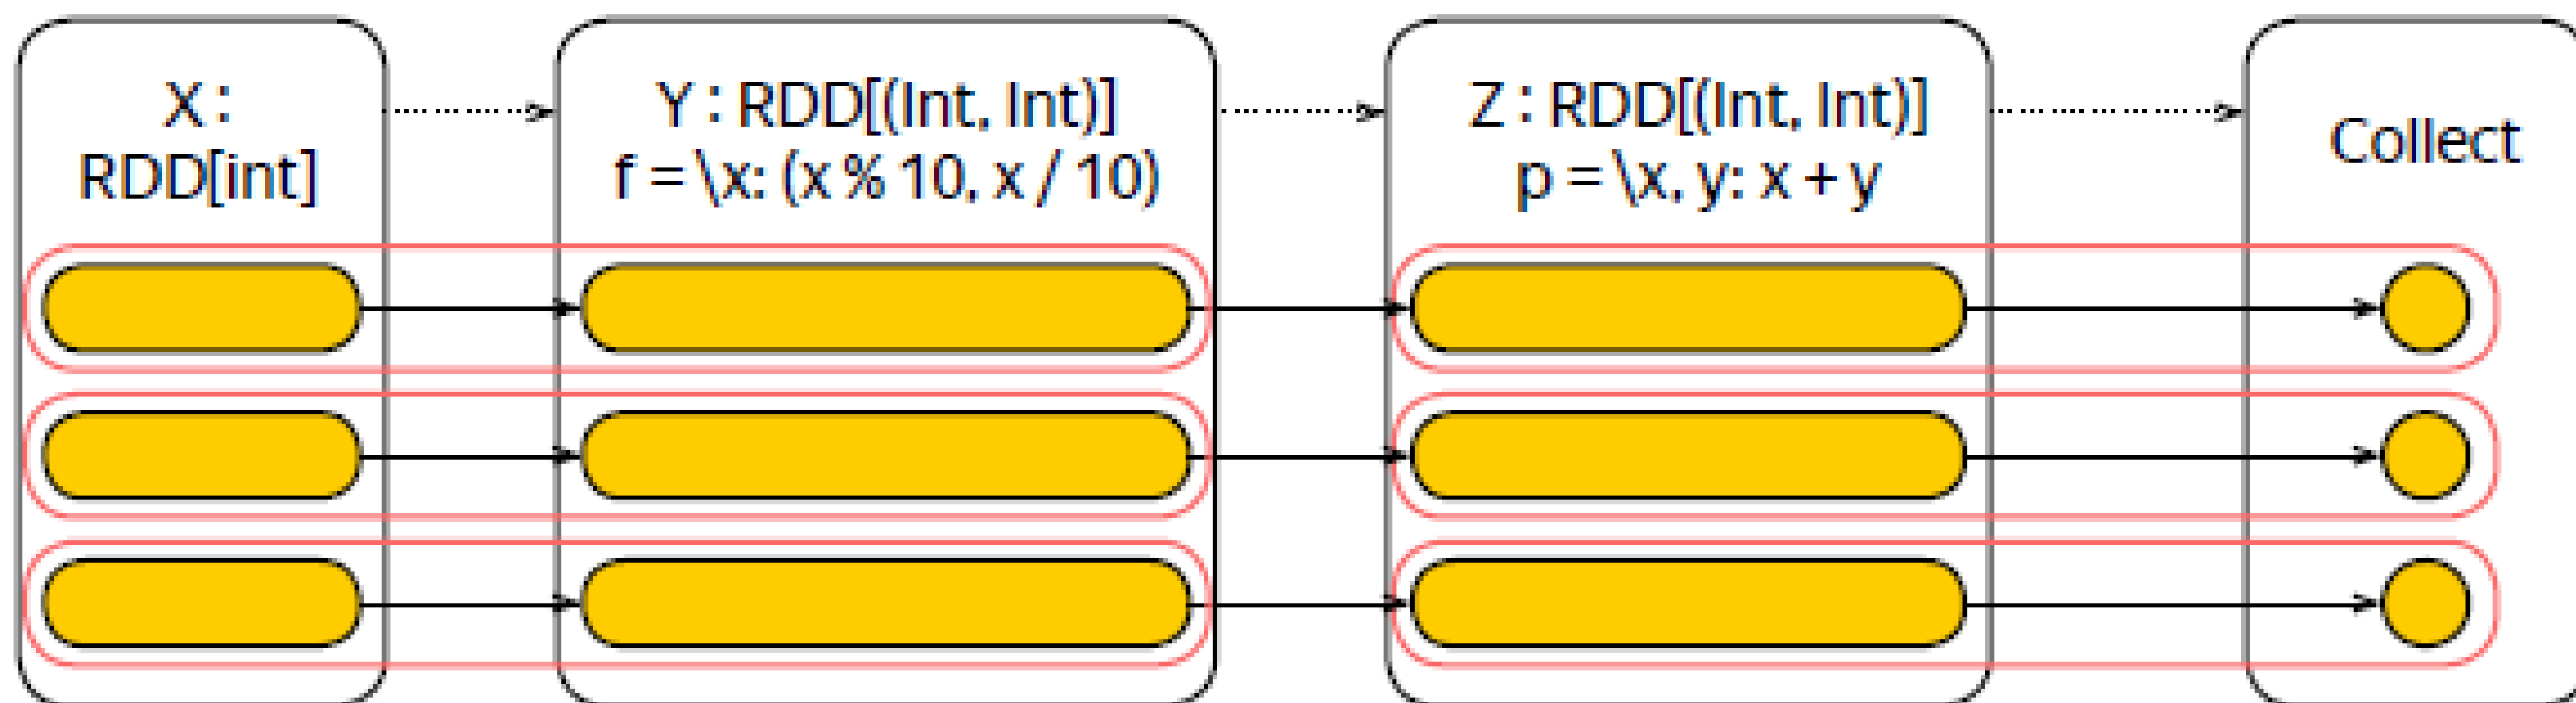
Jobs, stages, tasks

- › **Task** is a unit of work to be done
 - ›› Tasks are created by a **job scheduler** for every job stage
- › **Job** is spawned in response to a Spark action
 - ›› Job is divided in smaller sets of tasks called **stages**

Jobs, stages, tasks (example)

```
>> Z = X  
    .map(lambda x: (x % 10, x / 10))  
    .reduceByKey(lambda x, y: x + y)  
    .collect()
```

1. Invoking an action... `collect()`
2. ...spawns the job...
3. ... that gets divided into the stages by the job scheduler...
4. ...and tasks are created for every job stage.



Jobs, stages, tasks

Stage는 파이프라인을 위한 구조

- › Job **stage** is a pipelined computation spanning between materialization boundaries
 - ››not immediately executable **RDD Level**
- › Task is a job stage bound to particular partitions
 - ››immediately executable **Partition Level**
- ›› Materialization happens when reading, shuffling or passing data to an action **Materialization == building**
 - ››narrow dependencies allow pipelining
 - ››wide dependencies forbid it

SparkContext의 역할

- ›› Tracks liveness of the executors
 - ›› required to provide fault-tolerance
- ›› Schedules multiple concurrent jobs
 - ›› to control the resource allocation **within** the application
- ›› Performs dynamic resource allocation
 - ›› to control the resource allocation **between** different applications

Summary

- » The SparkContext is the core of your application
- » The driver communicates directly with the executors
- » Execution goes as follows:
Action → Job → Job Stages → Tasks
- » Transformations with narrow dependencies allow pipelining

Caching & persistence

Intermediate Data

Quick reminder

- » RDDs are partitioned
- » Execution is build around the partitions
- › **Block** is a unit of input and output in Spark

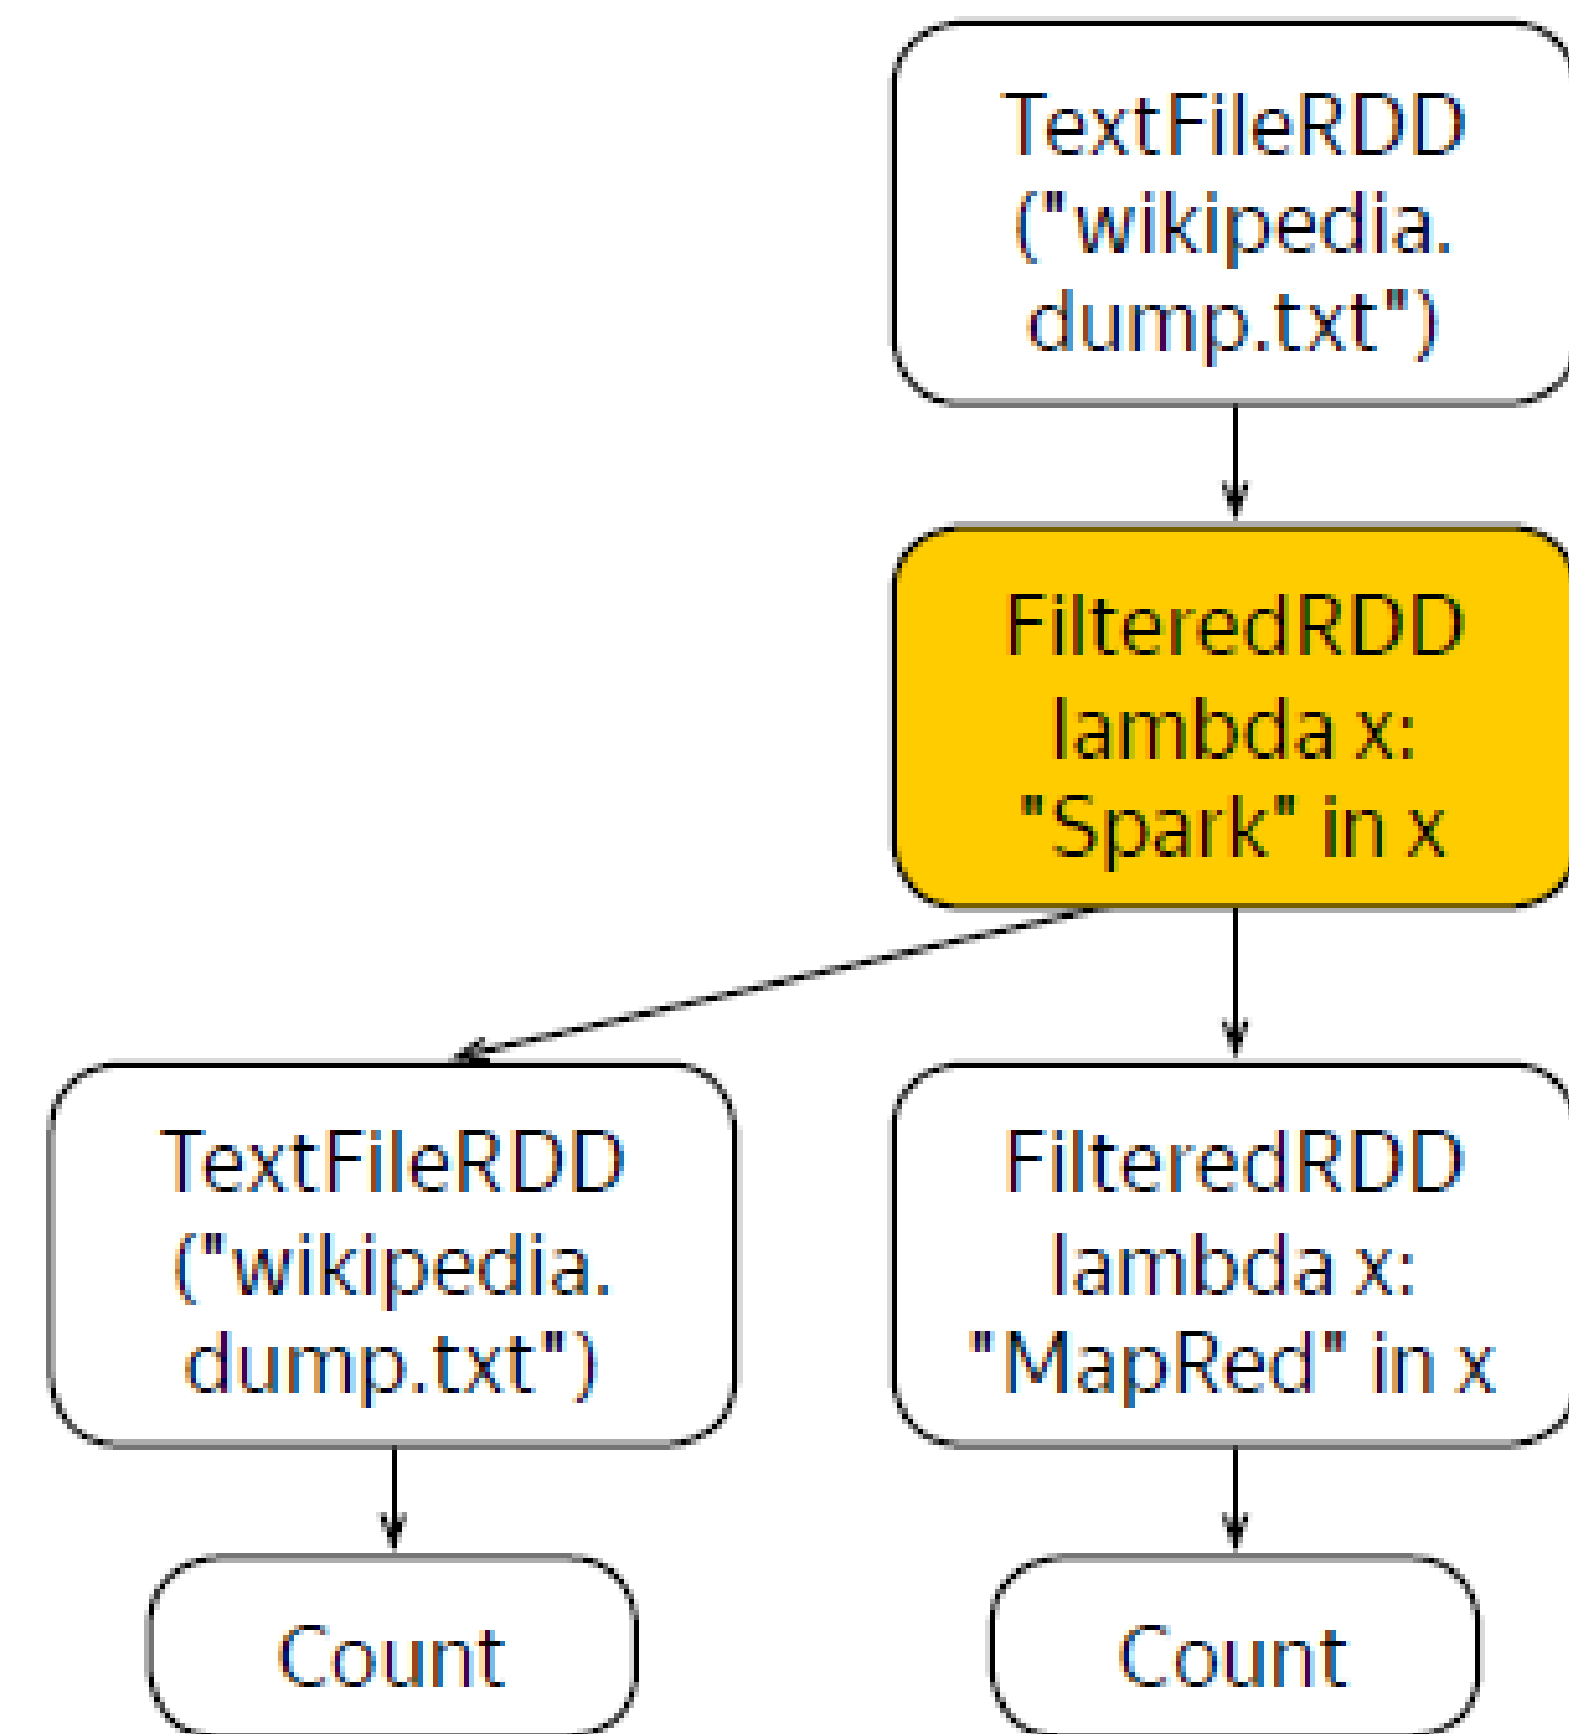
Motivating example

```
sc = SparkContext(...)
wiki = sc.textFile("wikipedia.dump.txt")

spark_articles = wiki.filter(
    lambda x: "Spark" in x)

spark_articles.cache()

hadoop_articles = spark_articles.filter(
    lambda x: "Hadoop" in x)
mapreduce_articles = spark_articles.filter(
    lambda x: "MapRed" in x)
print(hadoop_articles.count())
print(mapreduce_articles.count())
```



메모리에 RDD를 cache로 저장하면 불필요한 작업을 줄일 수 있다.
(같은 작업 반복 X)

Controlling persistence level

>> rdd.persist(storageLevel)

>>sets RDD's storage to persist across operations after it is computed for the first time

>>storageLevel is a set of flags controlling the persistence, typical values are

DISK_ONLY

– save the data to the disk,

MEMORY_ONLY

– keep the data in the memory

MEMORY_AND_DISK

– keep the data in the memory; when out of memory – save it to the disk

DISK_ONLY_2, MEMORY_ONLY_2, MEMORY_AND_DISK_2

– same as about, but make two replicas ← improves failure recovery times!

>> rdd.cache() = rdd.persist(MEMORY_ONLY)

cache함수는 메모리에 저장한다는 함수의 shortcut이다.

Best practices

일반적인 데이터 persist 방법

- » For interactive sessions
 - » cache preprocessed data
- » For batch computations
 - » cache dictionaries
 - » cache other datasets that are accessed multiple times
- » For iterative computations
 - » cache static data
- » And do benchmarks!

Summary

- » Performance may be improved by persisting data across operations
 - » in interactive sessions, iterative computations and hot datasets
- » You can control the persistence of a dataset
 - » whether to store in the memory or on the disk
 - » how many replicas to create

Broadcast variables

Shared Data

Broadcast variable

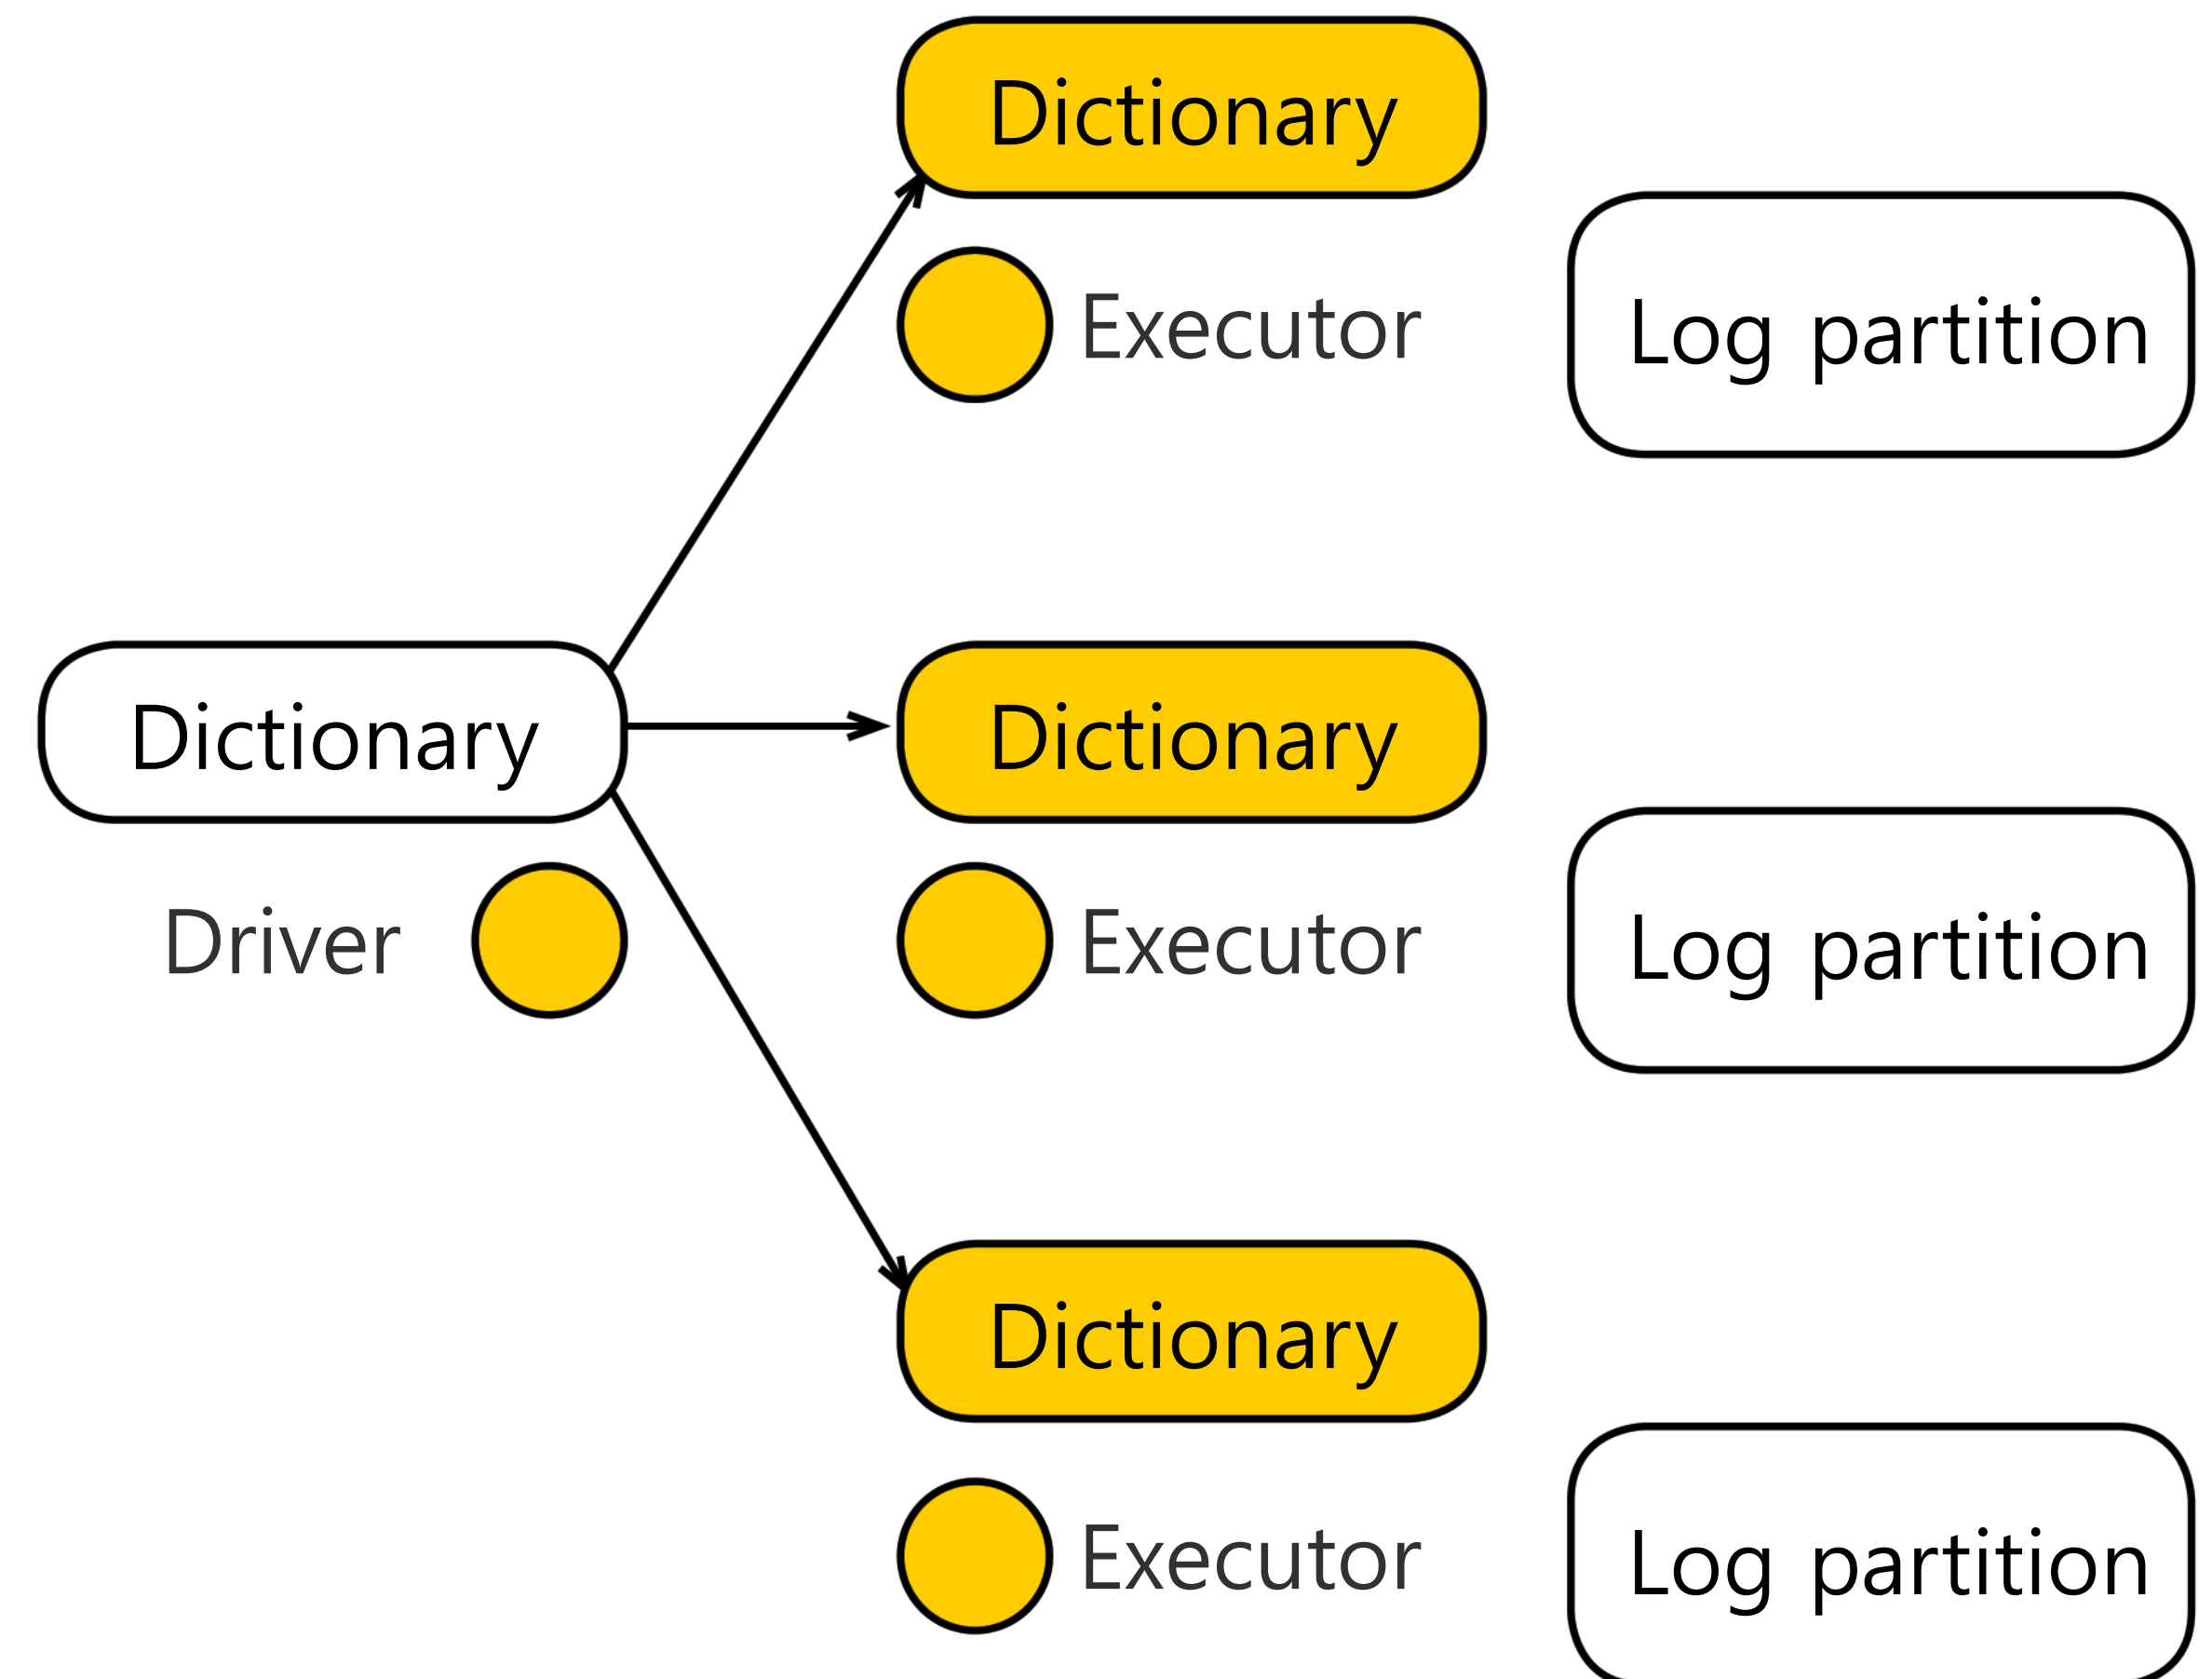
- » **Broadcast variable** is a read-only variable that is efficiently shared among tasks
- » Distribution is done by a torrent-like protocol (**extremely fast!**)
- » Distributed efficiently compared to captured variables

일반 variable을 closure에 넣으면 1 to many protocol
(한 곳에서 많은 executor로 전달해야 함)

Broadcast variable은 many to many protocol (토렌트와 같은 방식)

Motivating example

- › Input:
1TB partitioned log, 1GB IP dictionary
- › Task:
resolve IP addresses
- › Idea:
distribute the dictionary
query it locally



Motivating example

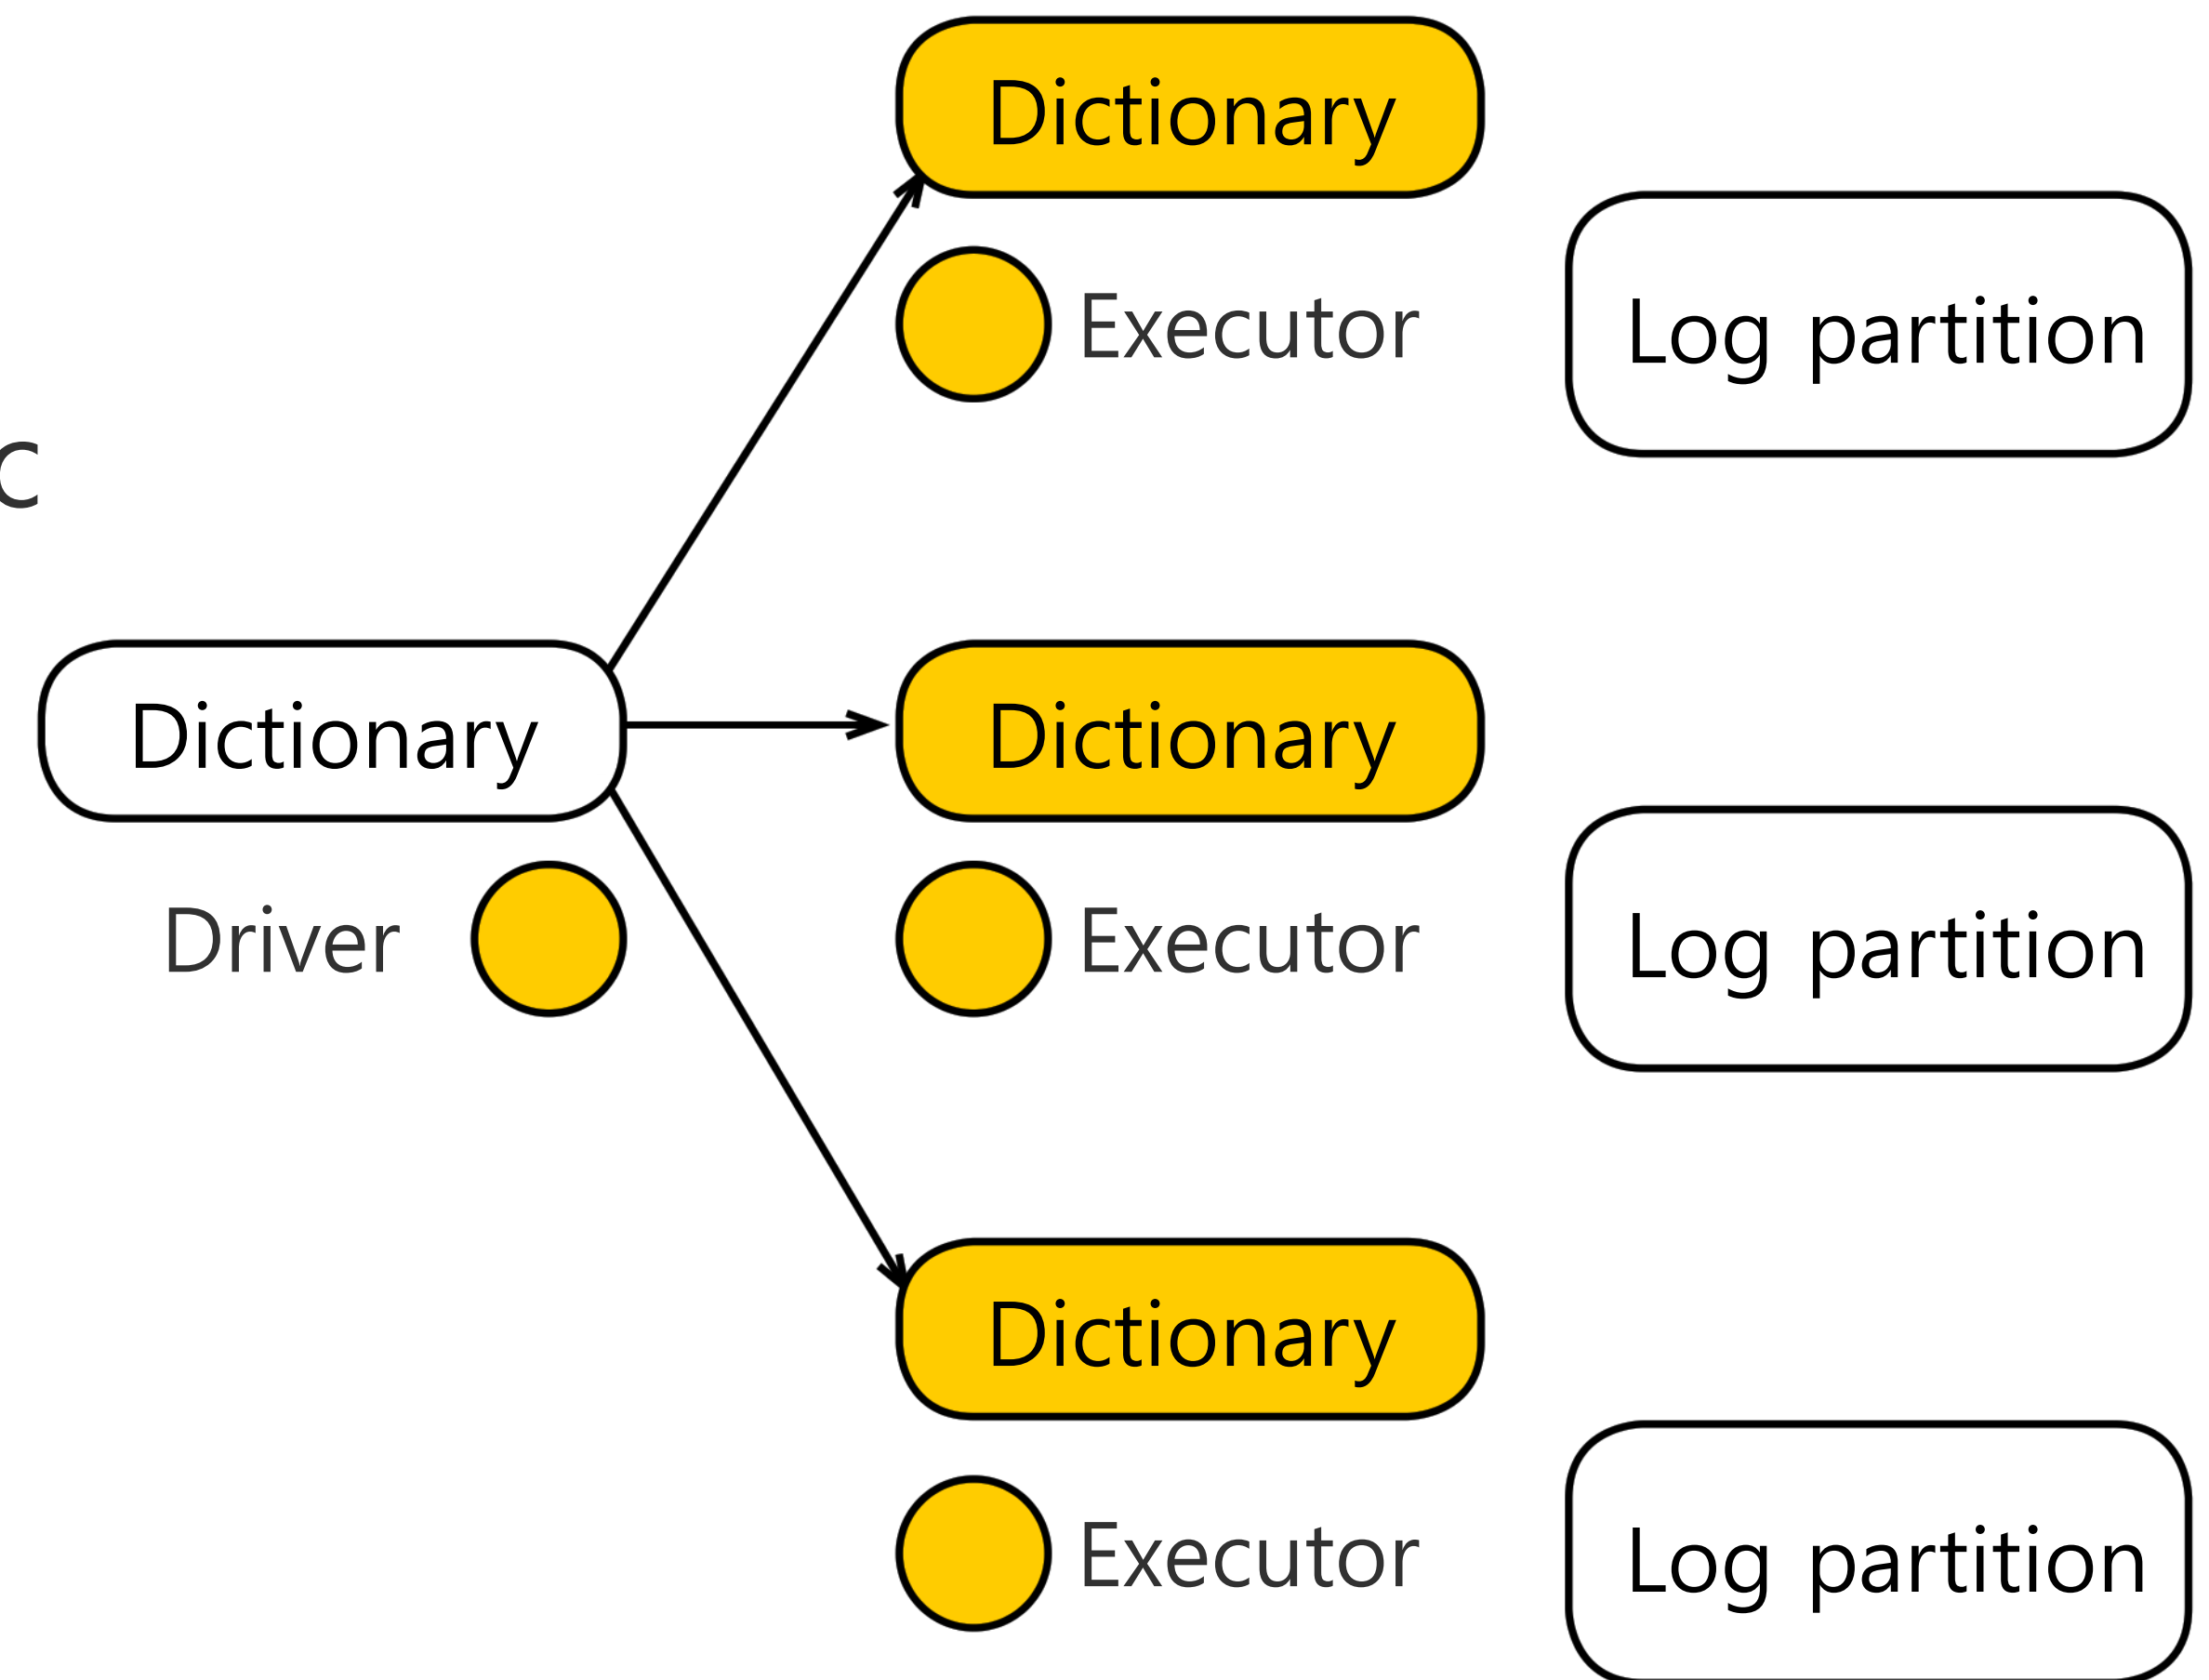
Serial distribution via the closure
(from the driver to every executor)
 $\sim 1000 \text{ (tasks)} * 1\text{GB} = 1\text{TB}$ of traffic

1GB 데이터를 driver가 모두 전달

Parallel distribution via
the broadcast variable
(torrent-like)

$\sim 1\text{-}2 \text{ GB}$ of traffic

Faster!



Motivating example 2

```
sc = SparkContext(conf=...)
```

```
# compute the dictionary
```

```
my_dict_rdd = sc.textFile(...).map(...).filter(...)
```

```
my_dict_data = my_dict_rdd.collect()
```

```
# distributed the dictionary via the broadcast variable
```

```
broadcast_var = sc.broadcast(my_dict_data)
```

```
# use the broadcast variable within the task
```

```
my_data_rdd = sc.textFile(...).filter(
```

```
lambda x: x in broadcast_var.value)
```

Summary

- » Broadcast variables are read-only shared variables with effective sharing mechanism
단, memory에 맞는 양의 데이터를 활용가능하다.
- » Useful to share dictionaries, models

Accumulator variables

Useful for the control flow, monitoring, profiling & debugging

Accumulator variable

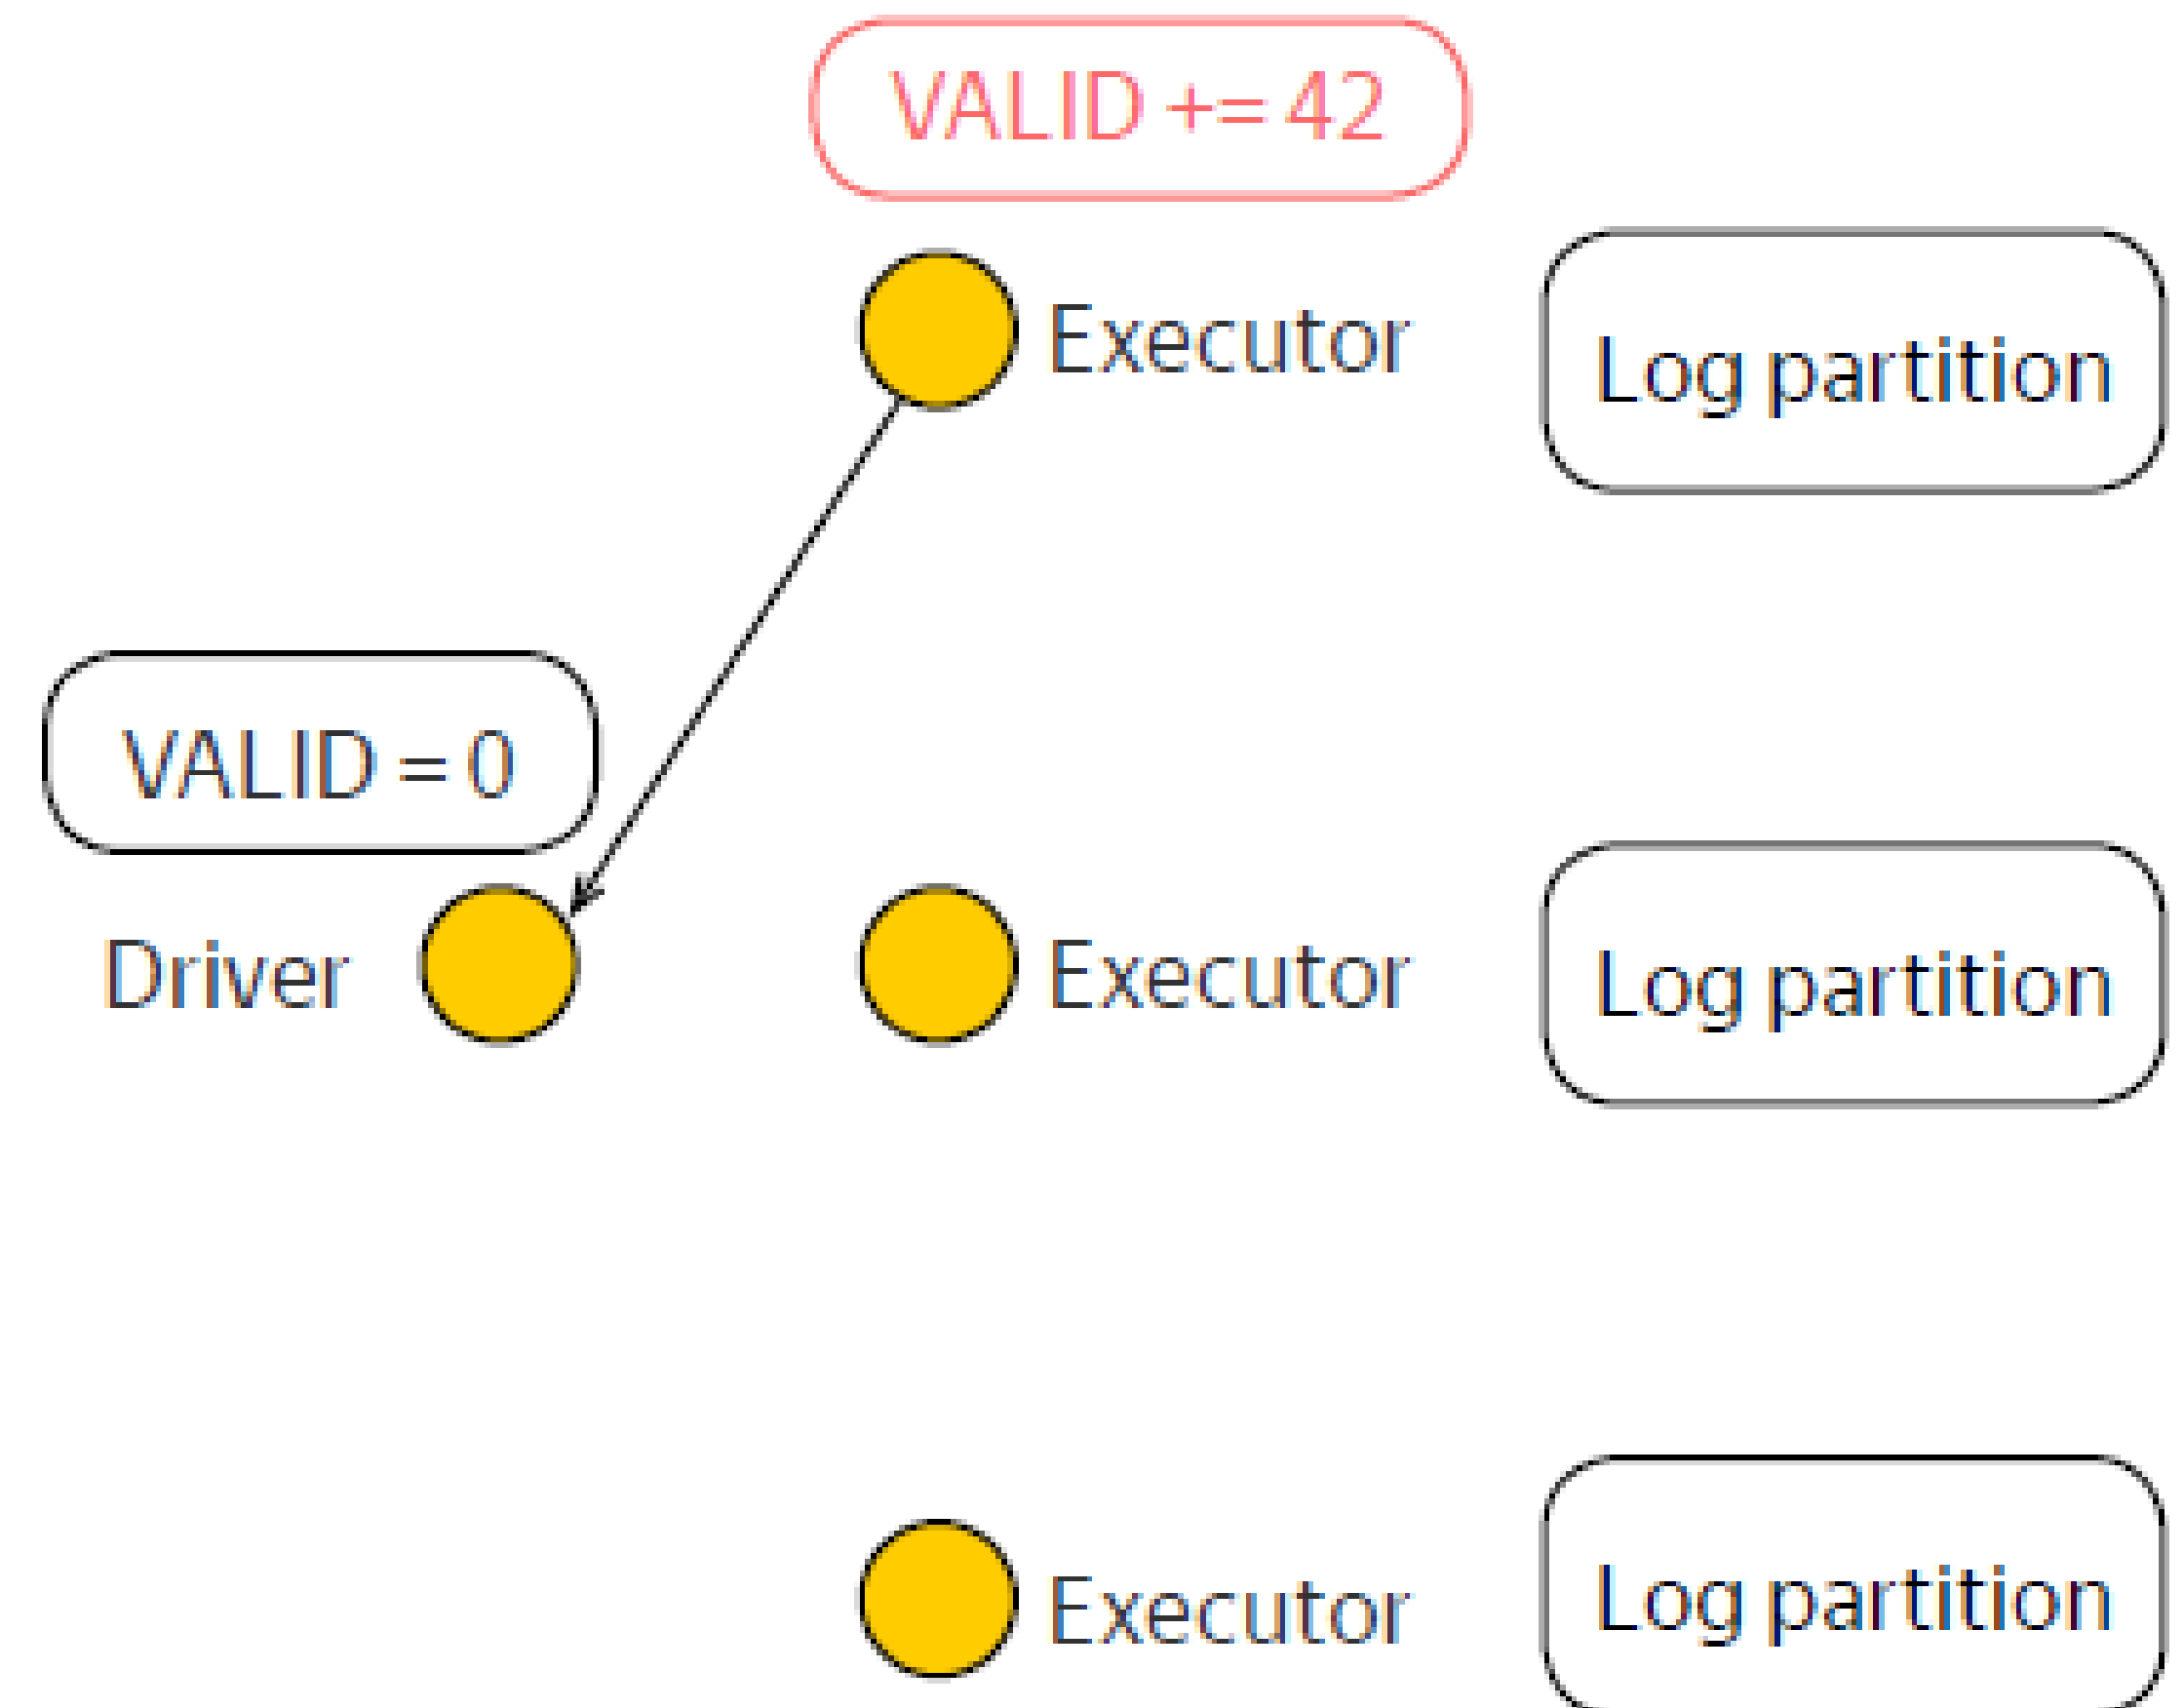
- › Accumulator variable is a read-write variable that is shared among tasks
- ›› Writes are restricted to increments! *synchronization 문제를 피하기 위해*
 - ››i. e.: `var += delta`
 - ››addition may be replaced by any associate, commutative operation
- ›› Reads are allowed *only by the driver program!* *task에서는 읽을 수 없다.*

Guarantees on the updates

- ›› In actions updates are applied exactly once
 - action에서는 accumulator에 한 번만 적용된다.*
- ›› In transformations there are no guarantees as the transformation code may be re-executed
 - transformation은 재실행될 수 있기 때문에 보장할 수 없다.*

Example (similar to the previous video)

- › Input:
1TB partitioned log
- › Task:
resolve IP addresses
AND
collect metrics:
of valid records



Example (similar to the previous video)

› Input:

1TB partitioned log


› Task:

resolve IP addresses

AND

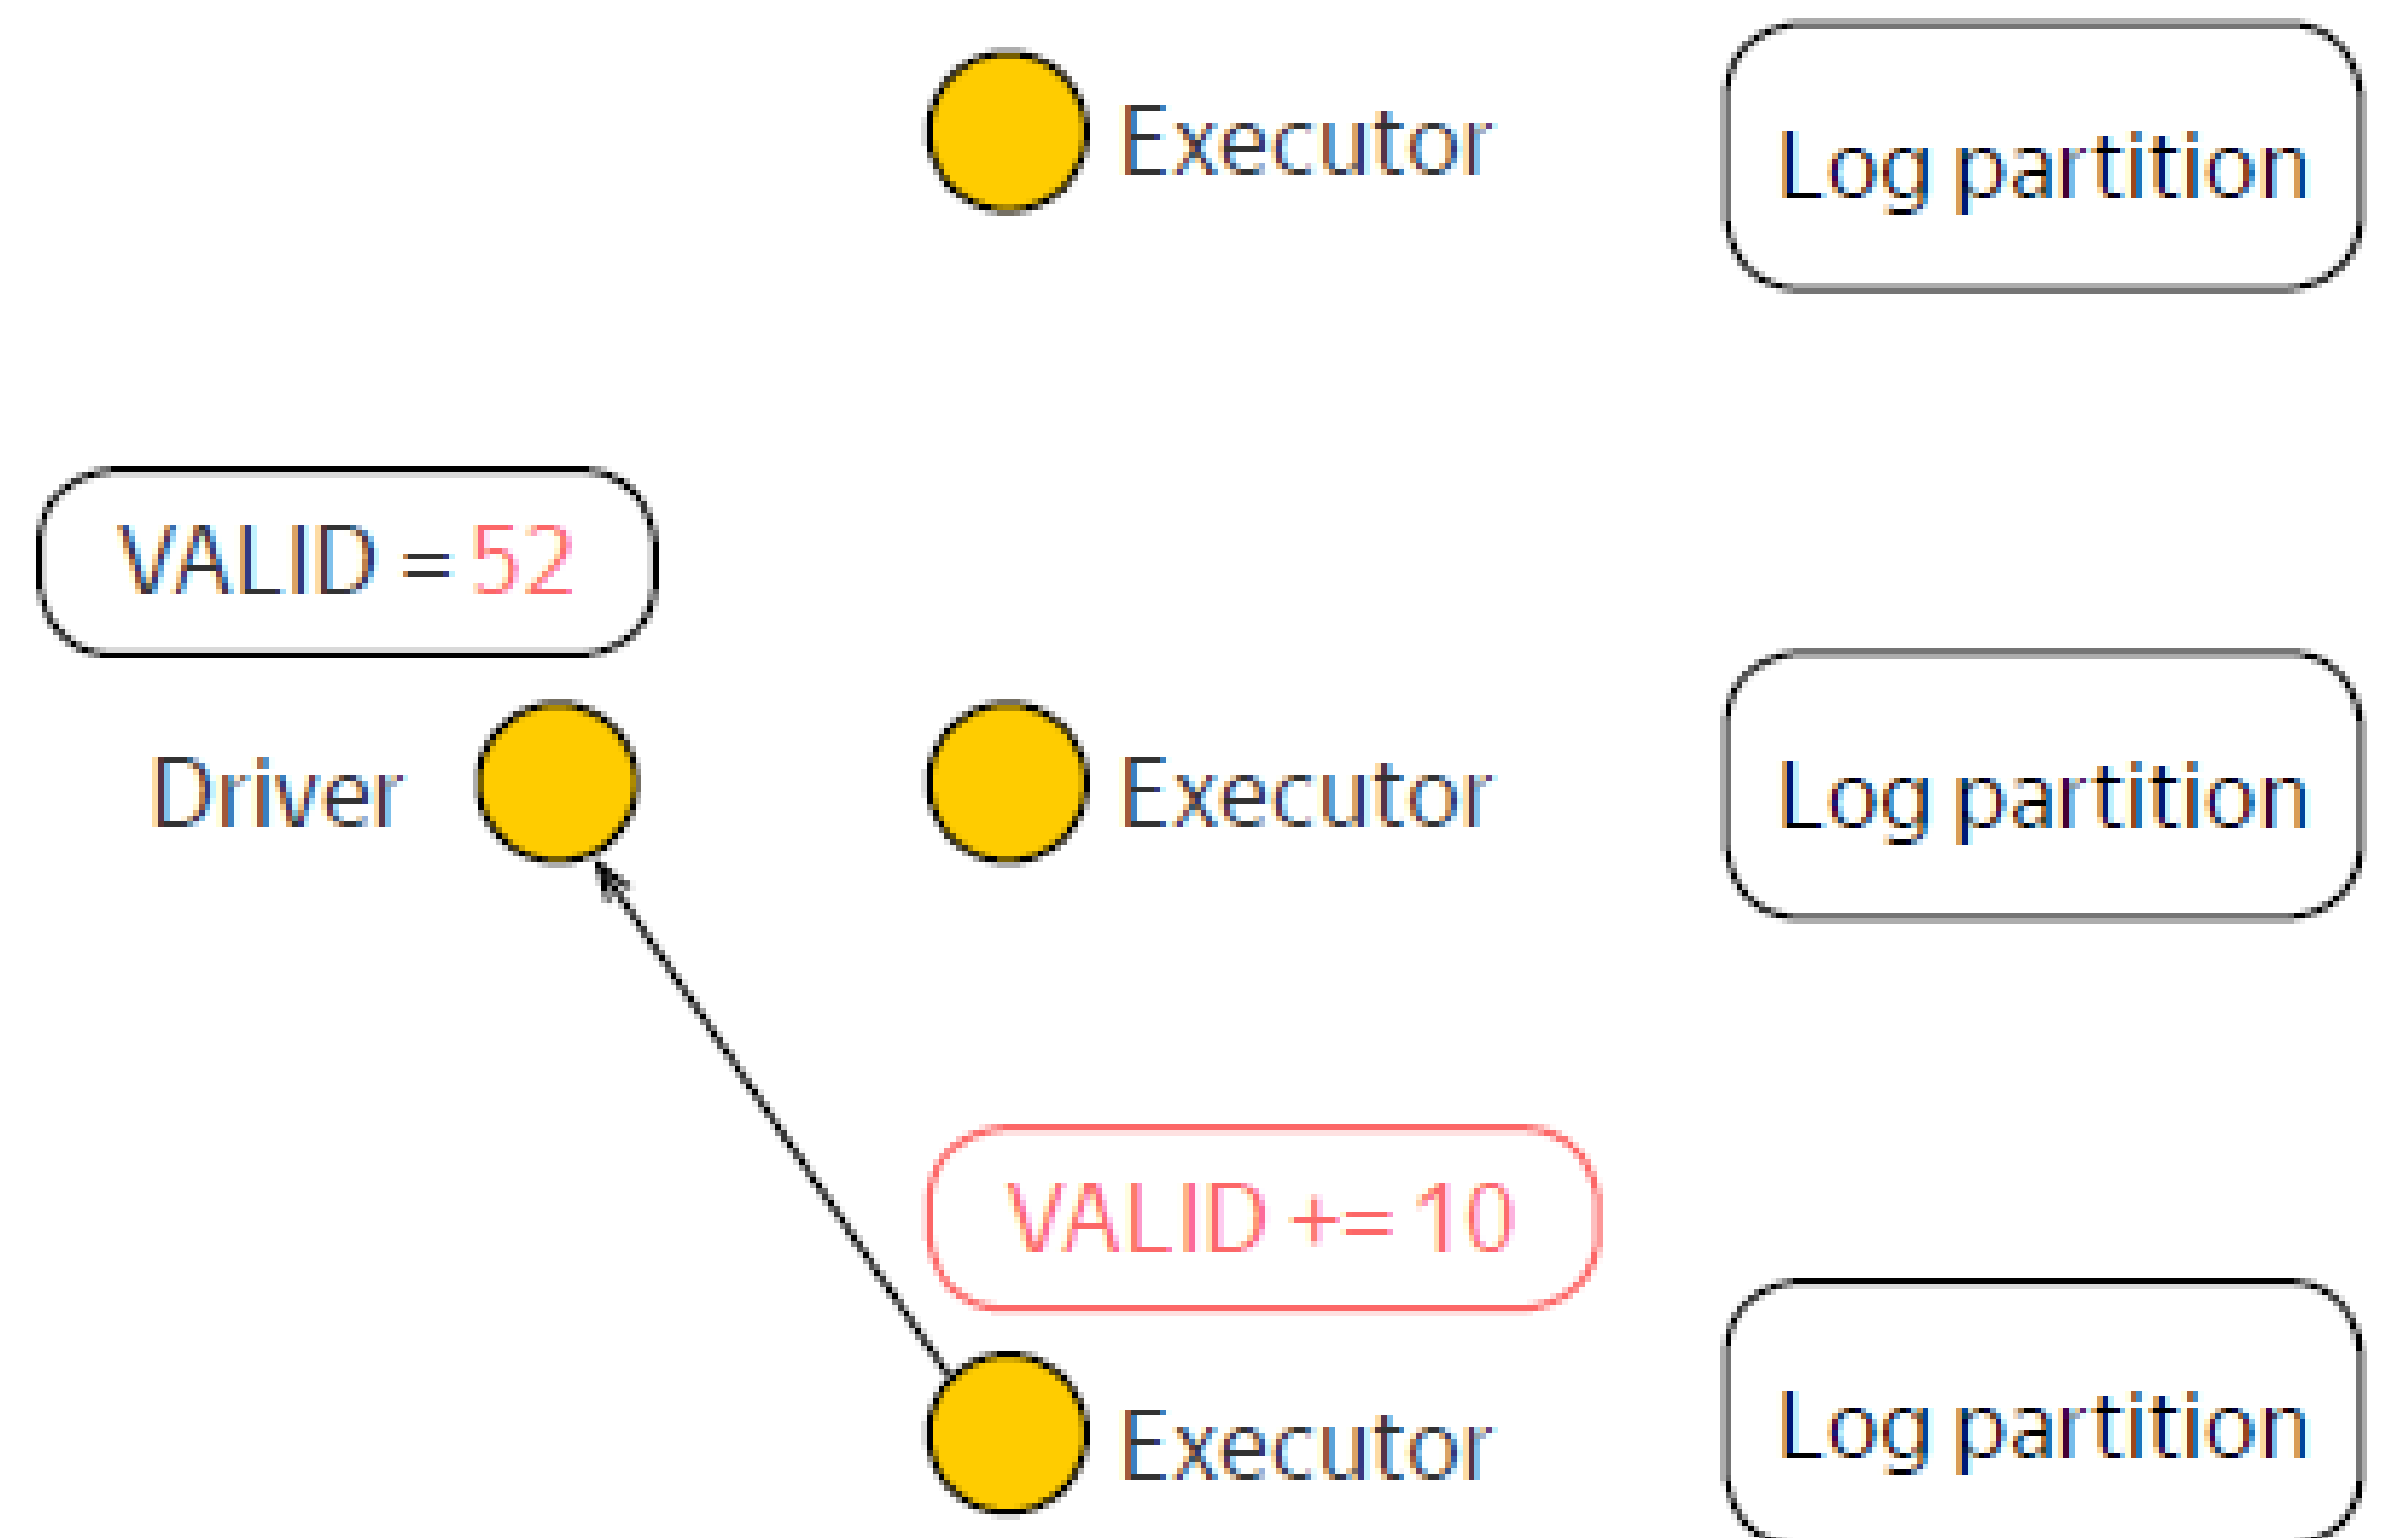
collect metrics:

of valid records



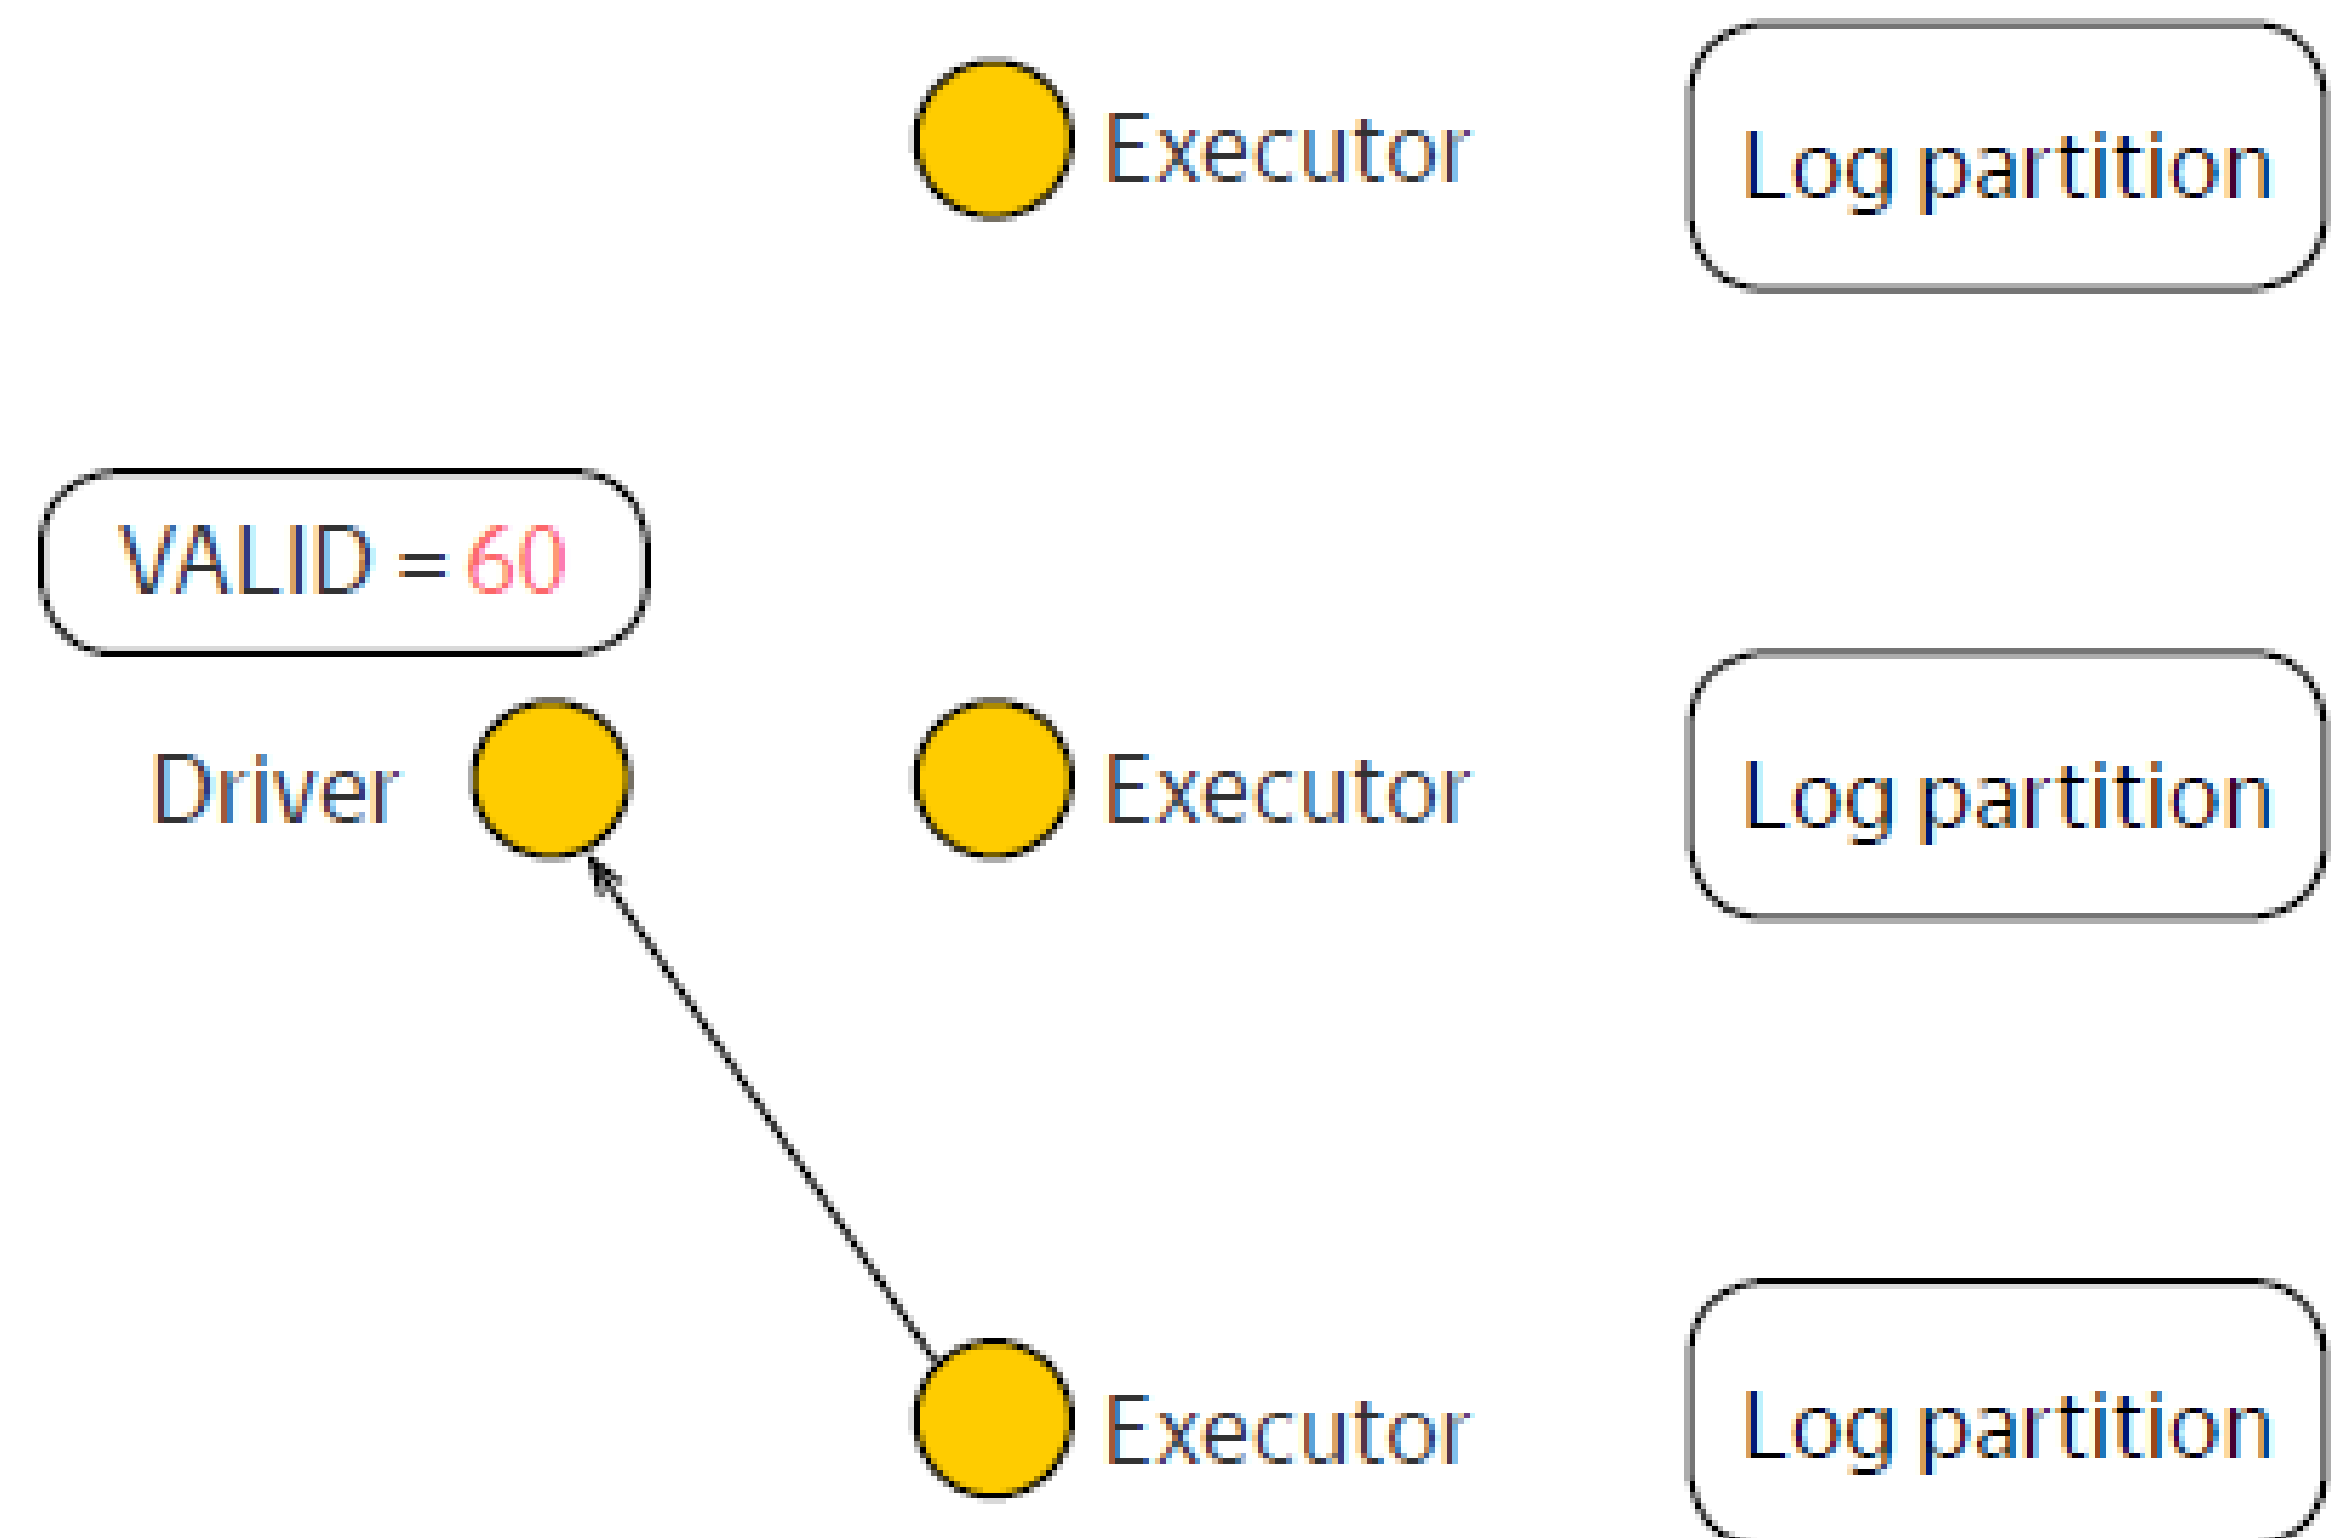
Example (similar to the previous video)

- › Input:
1TB partitioned log
- › Task:
resolve IP addresses
AND
collect metrics:
of valid records



Example (similar to the previous video)

- › Input:
1TB partitioned log
- › Task:
resolve IP addresses
AND
collect metrics:
of valid records



Use cases

- » Performance counters

- »# of processed records, total elapsed time, total error and so on and so forth

- » Simple control flow

- »conditionals: stop on reaching a threshold for corrupted records

- »loops: decide whether to run the next iteration of an algorithm or not

- » Monitoring

- »export values to the monitoring system

- » Profiling & debugging

Summary

- » Accumulators are read-write shared variables with restricted updates
 - » increments only
 - » can use custom associative, commutative operation for the updates
 - » can read the total value only in the driver
- » Useful for the control flow, monitoring, profiling & debugging

BigDATAteam