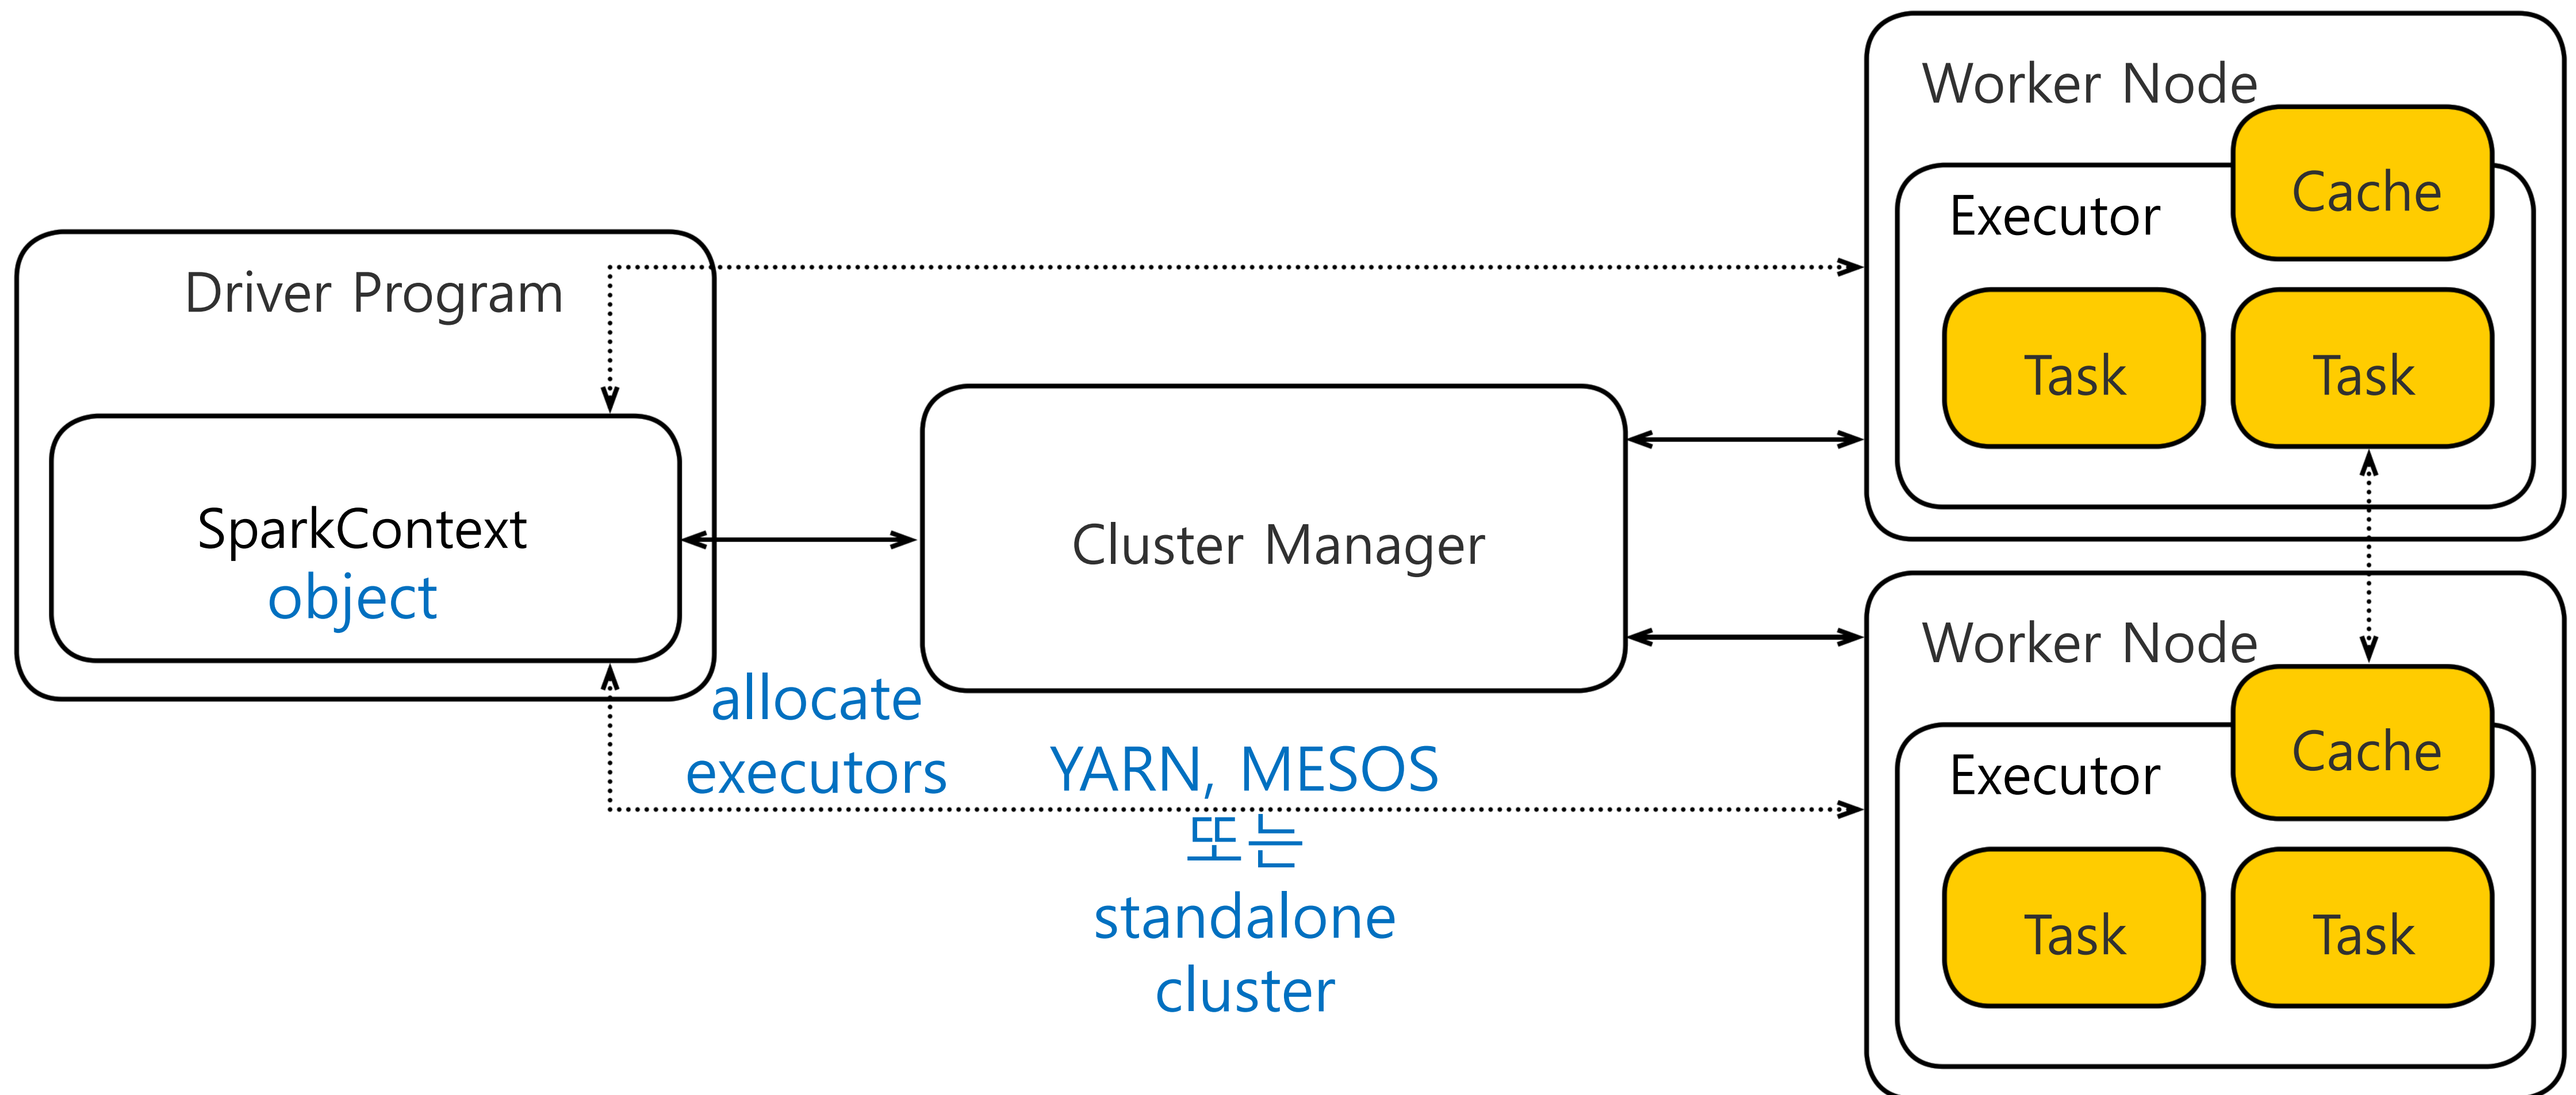


**Y**andex

# Execution & scheduling

# SparkContext

- » Tells your application how to access a cluster
- » Coordinates processes on the cluster to run your application



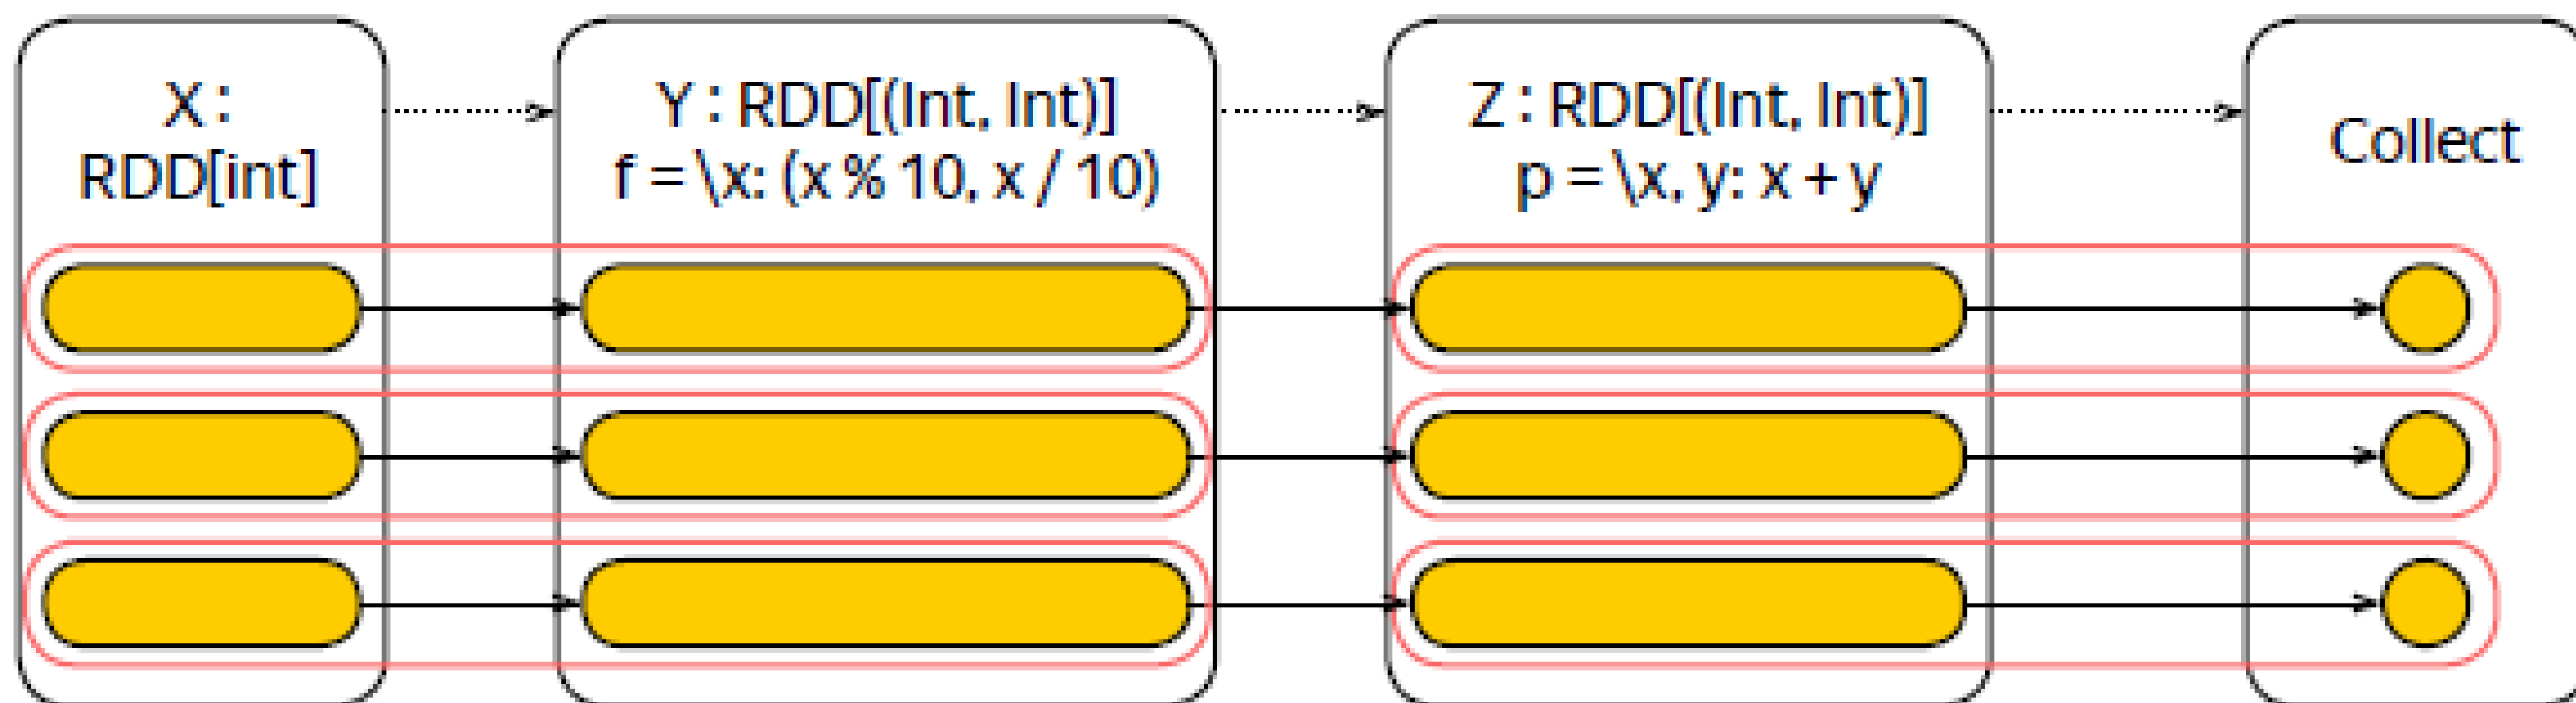
# Jobs, stages, tasks

- › **Task** is a unit of work to be done
  - ›› Tasks are created by a **job scheduler** for every job stage
- › **Job** is spawned in response to a Spark action
  - ›› Job is divided in smaller sets of tasks called **stages**

# Jobs, stages, tasks (example)

```
>> Z = X  
    .map(lambda x: (x % 10, x / 10))  
    .reduceByKey(lambda x, y: x + y)  
    .collect()
```

1. Invoking an action... `collect()`
2. ...spawns the job...
3. ... that gets divided into the stages by the job scheduler...
4. ...and tasks are created for every job stage.



# Jobs, stages, tasks

Stage는 파이프라인을 위한 구조

- › Job **stage** is a pipelined computation spanning between materialization boundaries
  - ››not immediately executable **RDD Level**
- › Task is a job stage bound to particular partitions
  - ››immediately executable **Partition Level**
- ›› Materialization happens when reading, shuffling or passing data to an action **Materialization == building**
  - ››narrow dependencies allow pipelining
  - ››wide dependencies forbid it

# SparkContext의 역할

- ›› Tracks liveness of the executors
  - ›› required to provide fault-tolerance
- ›› Schedules multiple concurrent jobs
  - ›› to control the resource allocation **within** the application
- ›› Performs dynamic resource allocation
  - ›› to control the resource allocation **between** different applications

# Summary

- » The SparkContext is the core of your application
- » The driver communicates directly with the executors
- » Execution goes as follows:  
Action → Job → Job Stages → Tasks
- » Transformations with narrow dependencies allow pipelining



# Caching & persistence

Intermediate Data

# Quick reminder

- » RDDs are partitioned
- » Execution is build around the partitions
- › **Block** is a unit of input and output in Spark

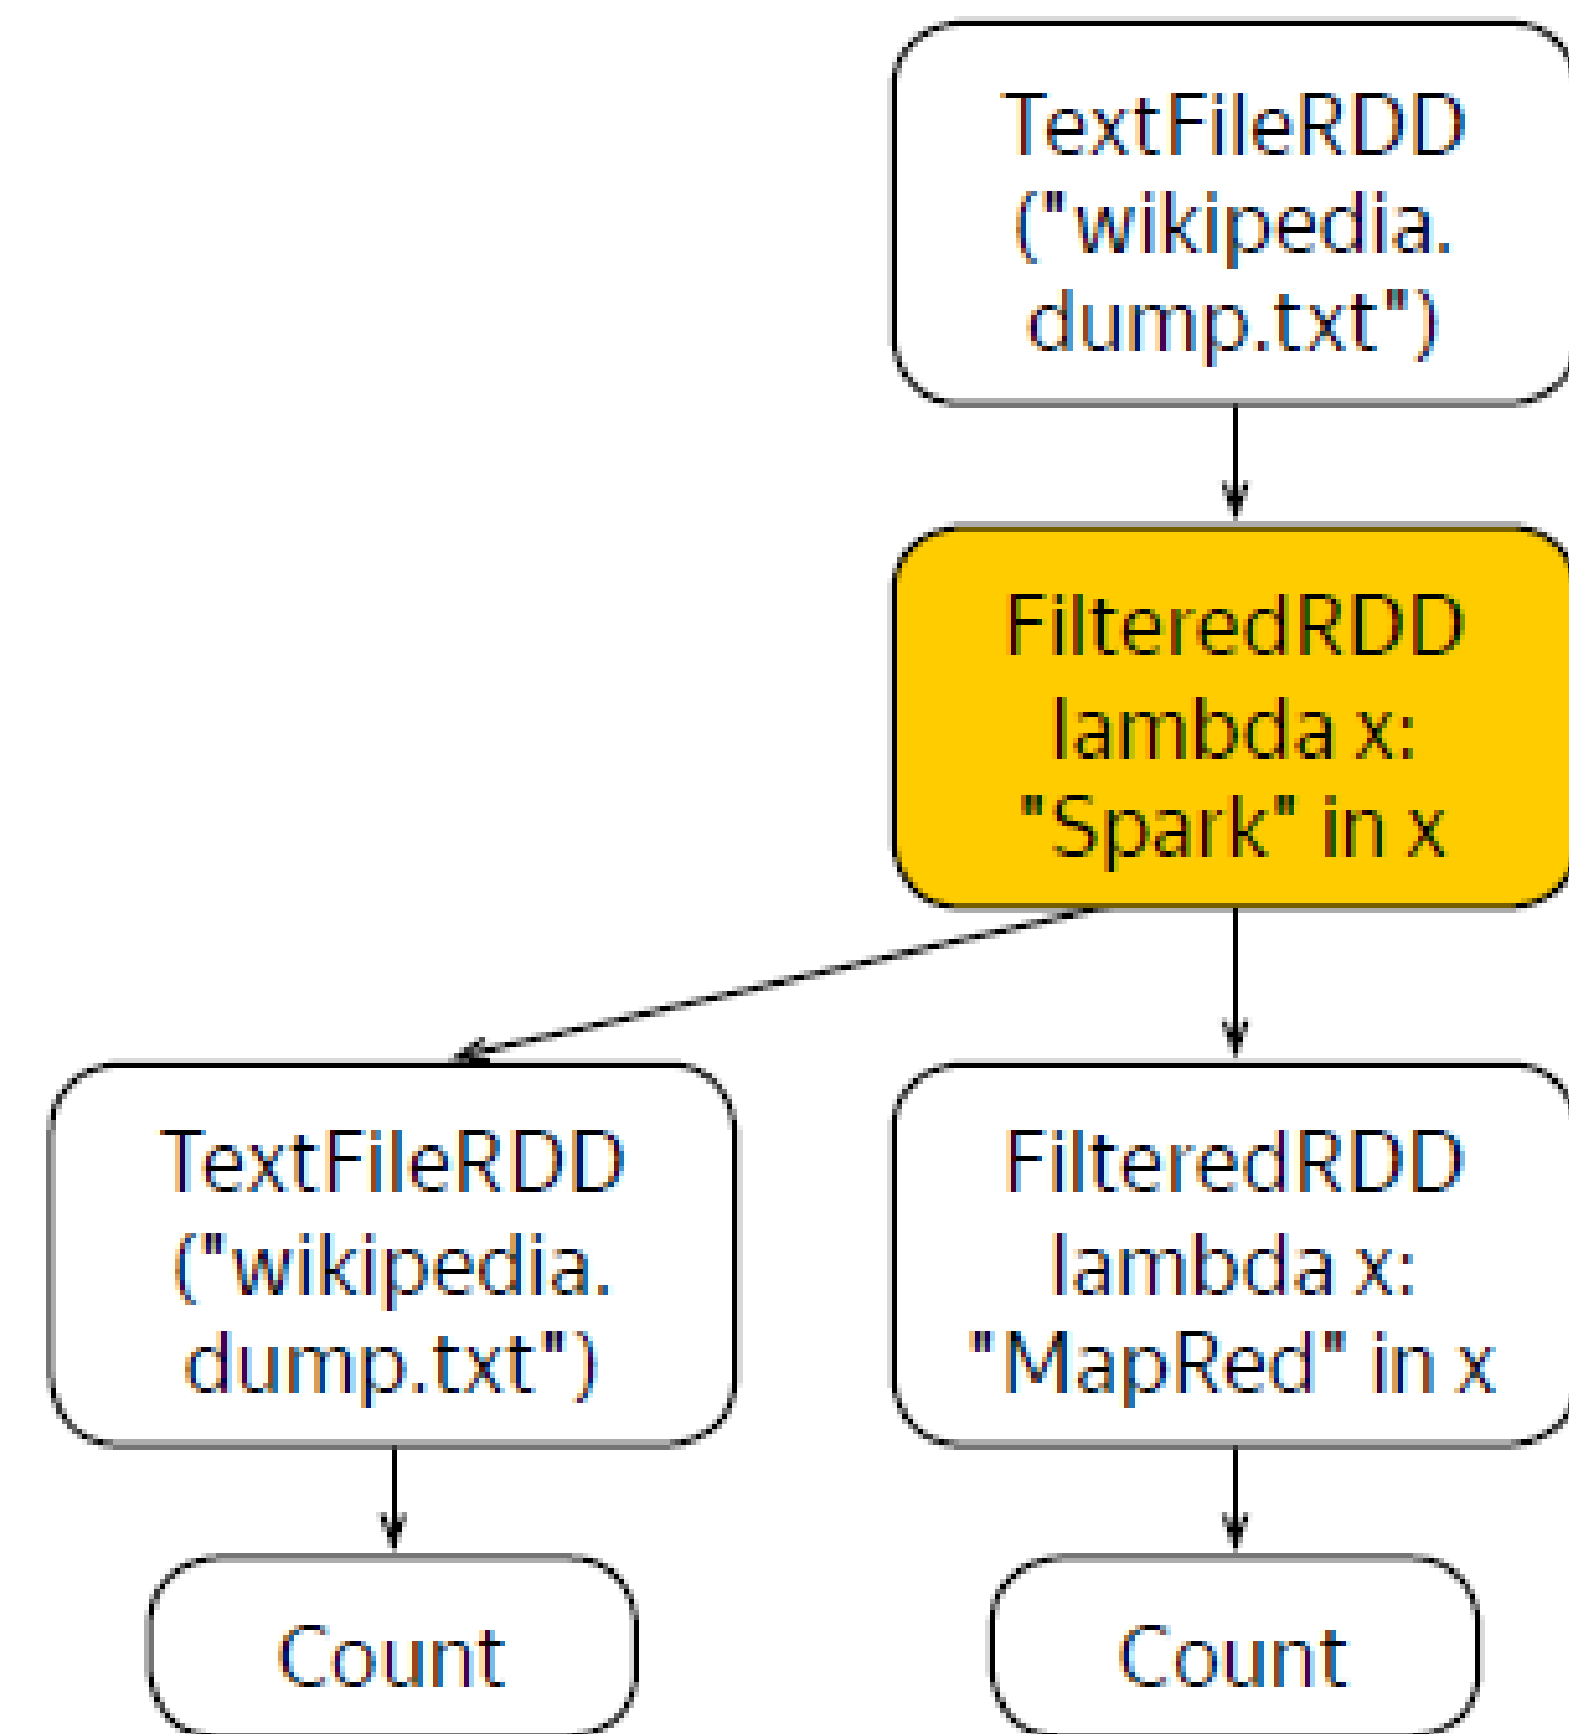
# Motivating example

```
sc = SparkContext(...)
wiki = sc.textFile("wikipedia.dump.txt")

spark_articles = wiki.filter(
    lambda x: "Spark" in x)

spark_articles.cache()

hadoop_articles = spark_articles.filter(
    lambda x: "Hadoop" in x)
mapreduce_articles = spark_articles.filter(
    lambda x: "MapRed" in x)
print(hadoop_articles.count())
print(mapreduce_articles.count())
```



메모리에 RDD를 cache로 저장하면 불필요한 작업을 줄일 수 있다.  
(같은 작업 반복 X)

# Controlling persistence level

>> rdd.persist(storageLevel)

>>sets RDD's storage to persist across operations after it is computed for the first time

>>storageLevel is a set of flags controlling the persistence, typical values are

DISK\_ONLY

– save the data to the disk,

MEMORY\_ONLY

– keep the data in the memory

MEMORY\_AND\_DISK

– keep the data in the memory; when out of memory – save it to the disk

DISK\_ONLY\_2, MEMORY\_ONLY\_2, MEMORY\_AND\_DISK\_2

– same as about, but make two replicas ← improves failure recovery times!

>> rdd.cache() = rdd.persist(MEMORY\_ONLY)

cache함수는 메모리에 저장한다는 함수의 shortcut이다.

# Best practices

## 일반적인 데이터 persist 방법

- » For interactive sessions
  - » cache preprocessed data
- » For batch computations
  - » cache dictionaries
  - » cache other datasets that are accessed multiple times
- » For iterative computations
  - » cache static data
- » And do benchmarks!

# Summary

- » Performance may be improved by persisting data across operations
  - » in interactive sessions, iterative computations and hot datasets
- » You can control the persistence of a dataset
  - » whether to store in the memory or on the disk
  - » how many replicas to create

# Broadcast variables

Shared Data

# Broadcast variable

- » **Broadcast variable** is a read-only variable that is efficiently shared among tasks
- » Distribution is done by a torrent-like protocol (**extremely fast!**)
- » Distributed efficiently compared to captured variables

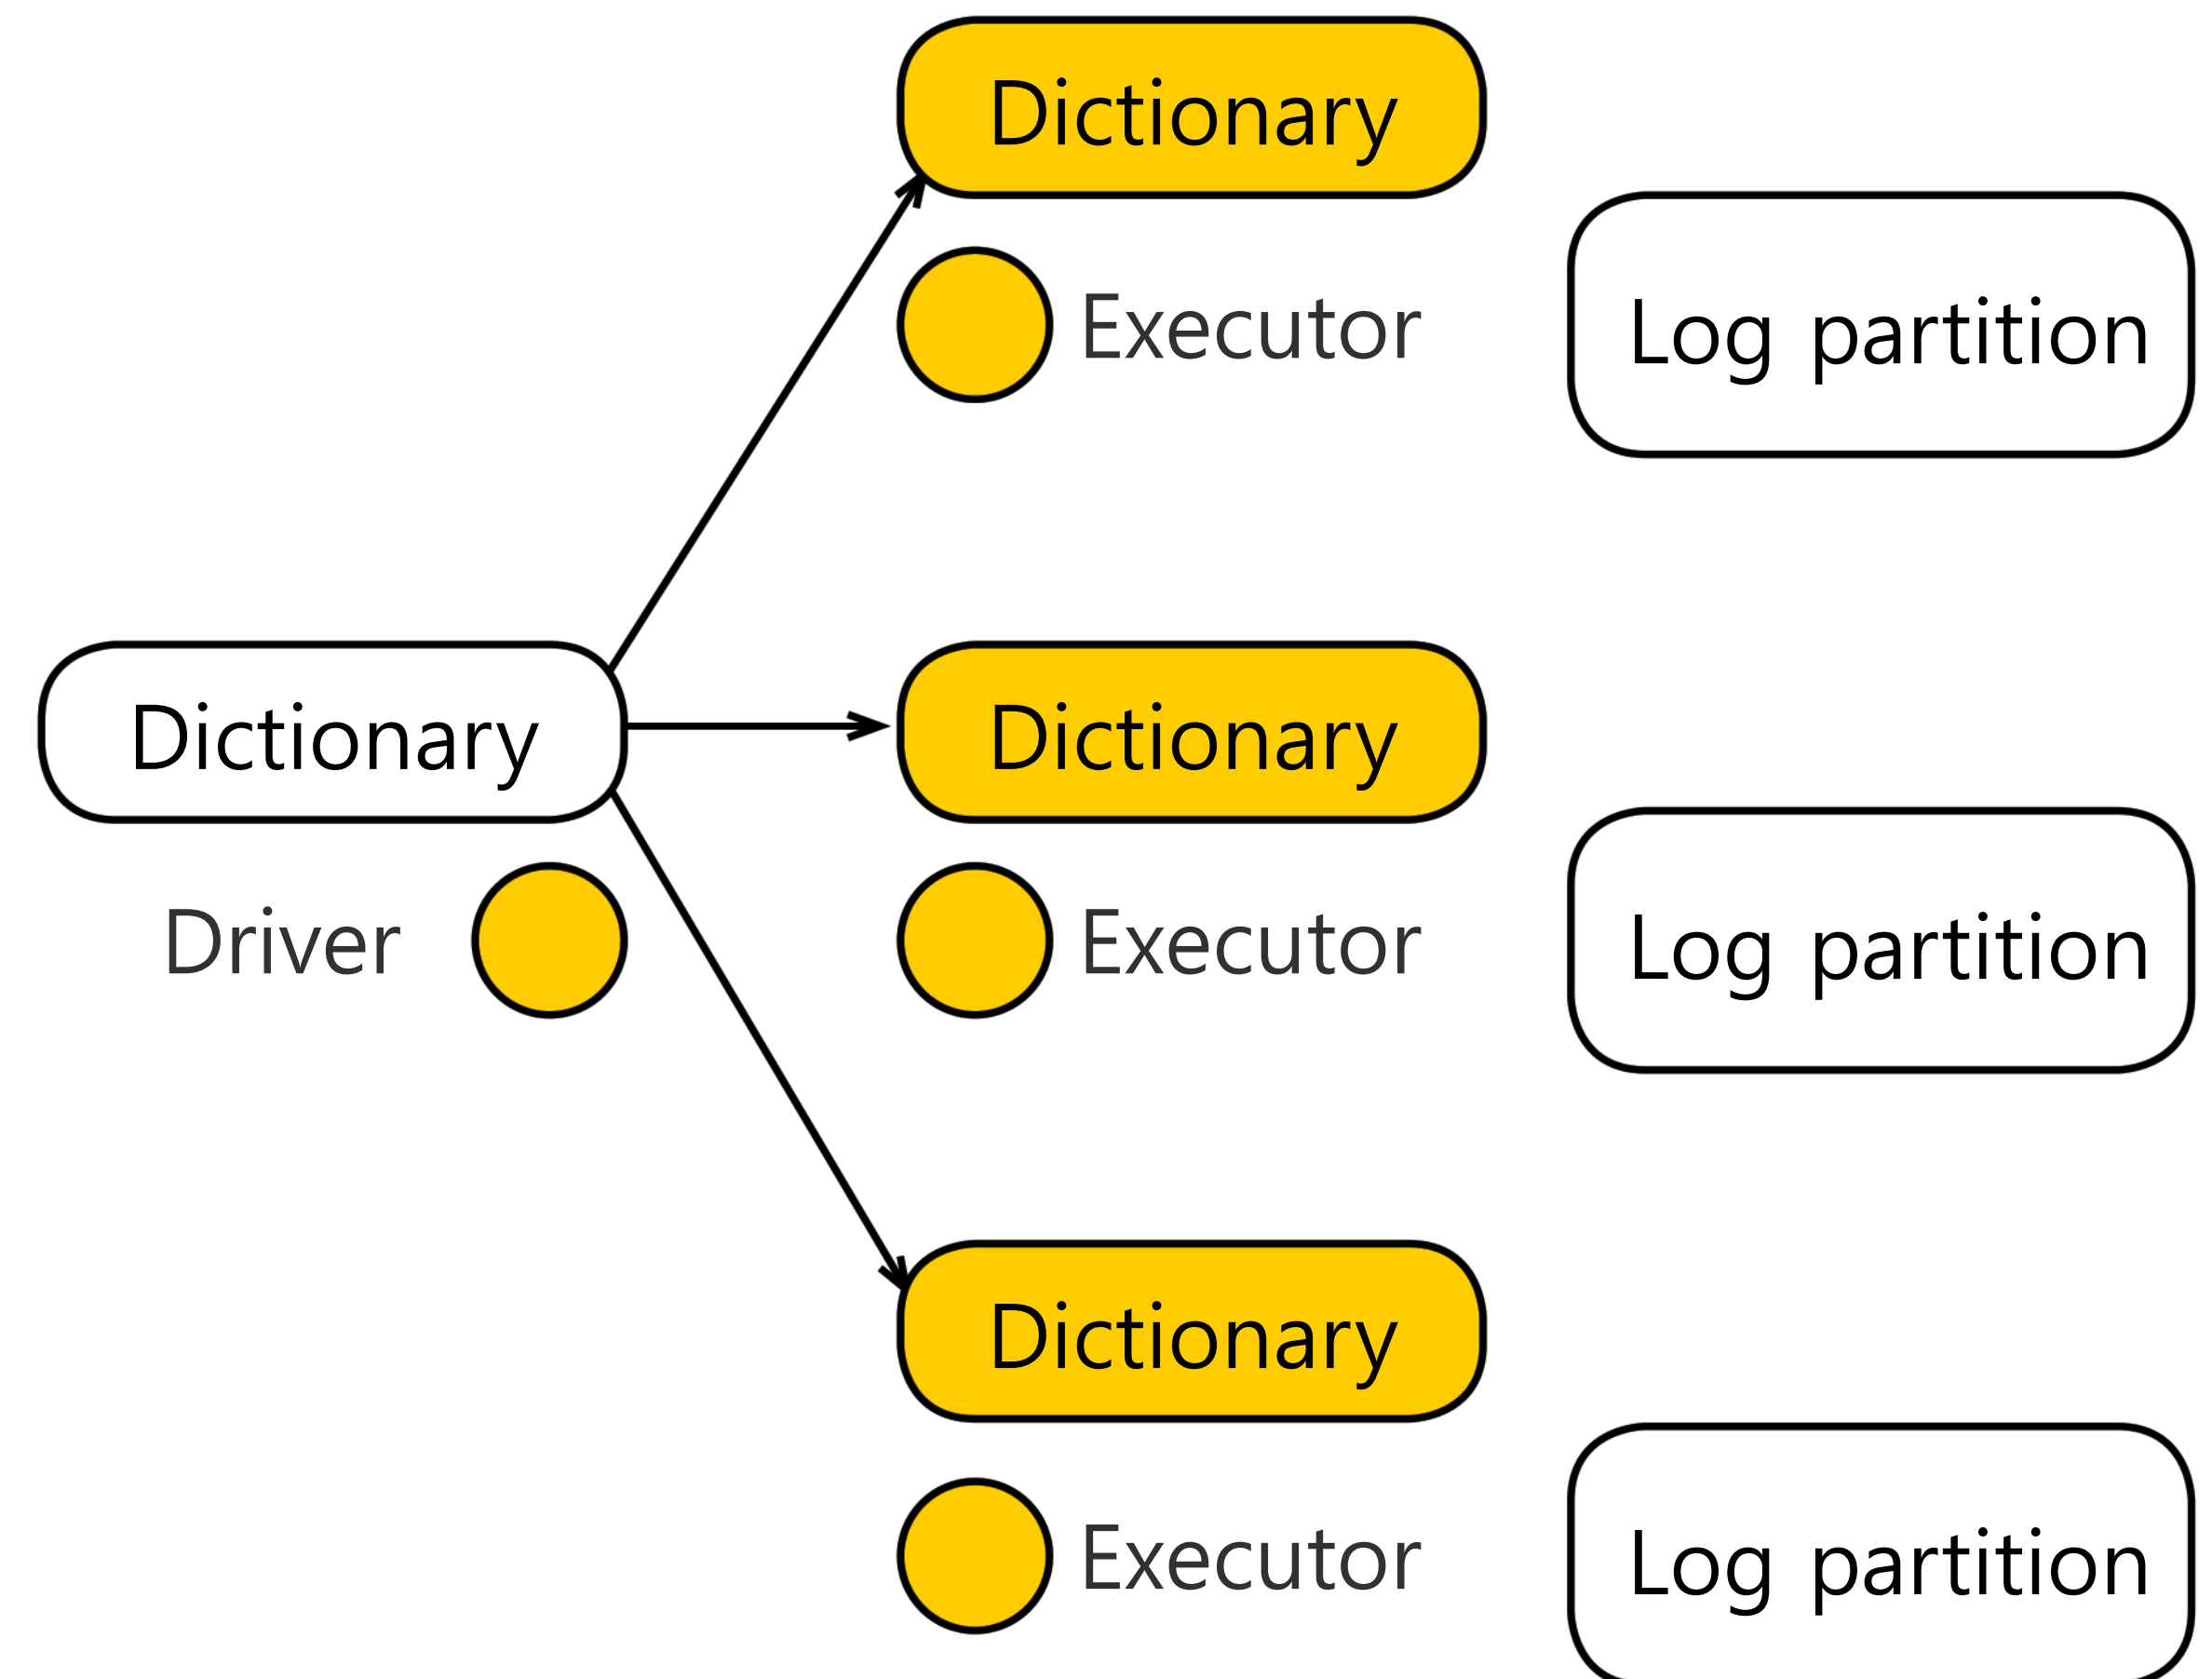
일반 variable을 closure에 넣으면 1 to many protocol  
( 한 곳에서 많은 executor로 전달해야 함 )

Broadcast variable은 many to many protocol (토렌트와 같은 방식)



# Motivating example

- › Input:  
1TB partitioned log, 1GB IP dictionary
- › Task:  
resolve IP addresses
- › Idea:  
distribute the dictionary  
query it locally



# Motivating example

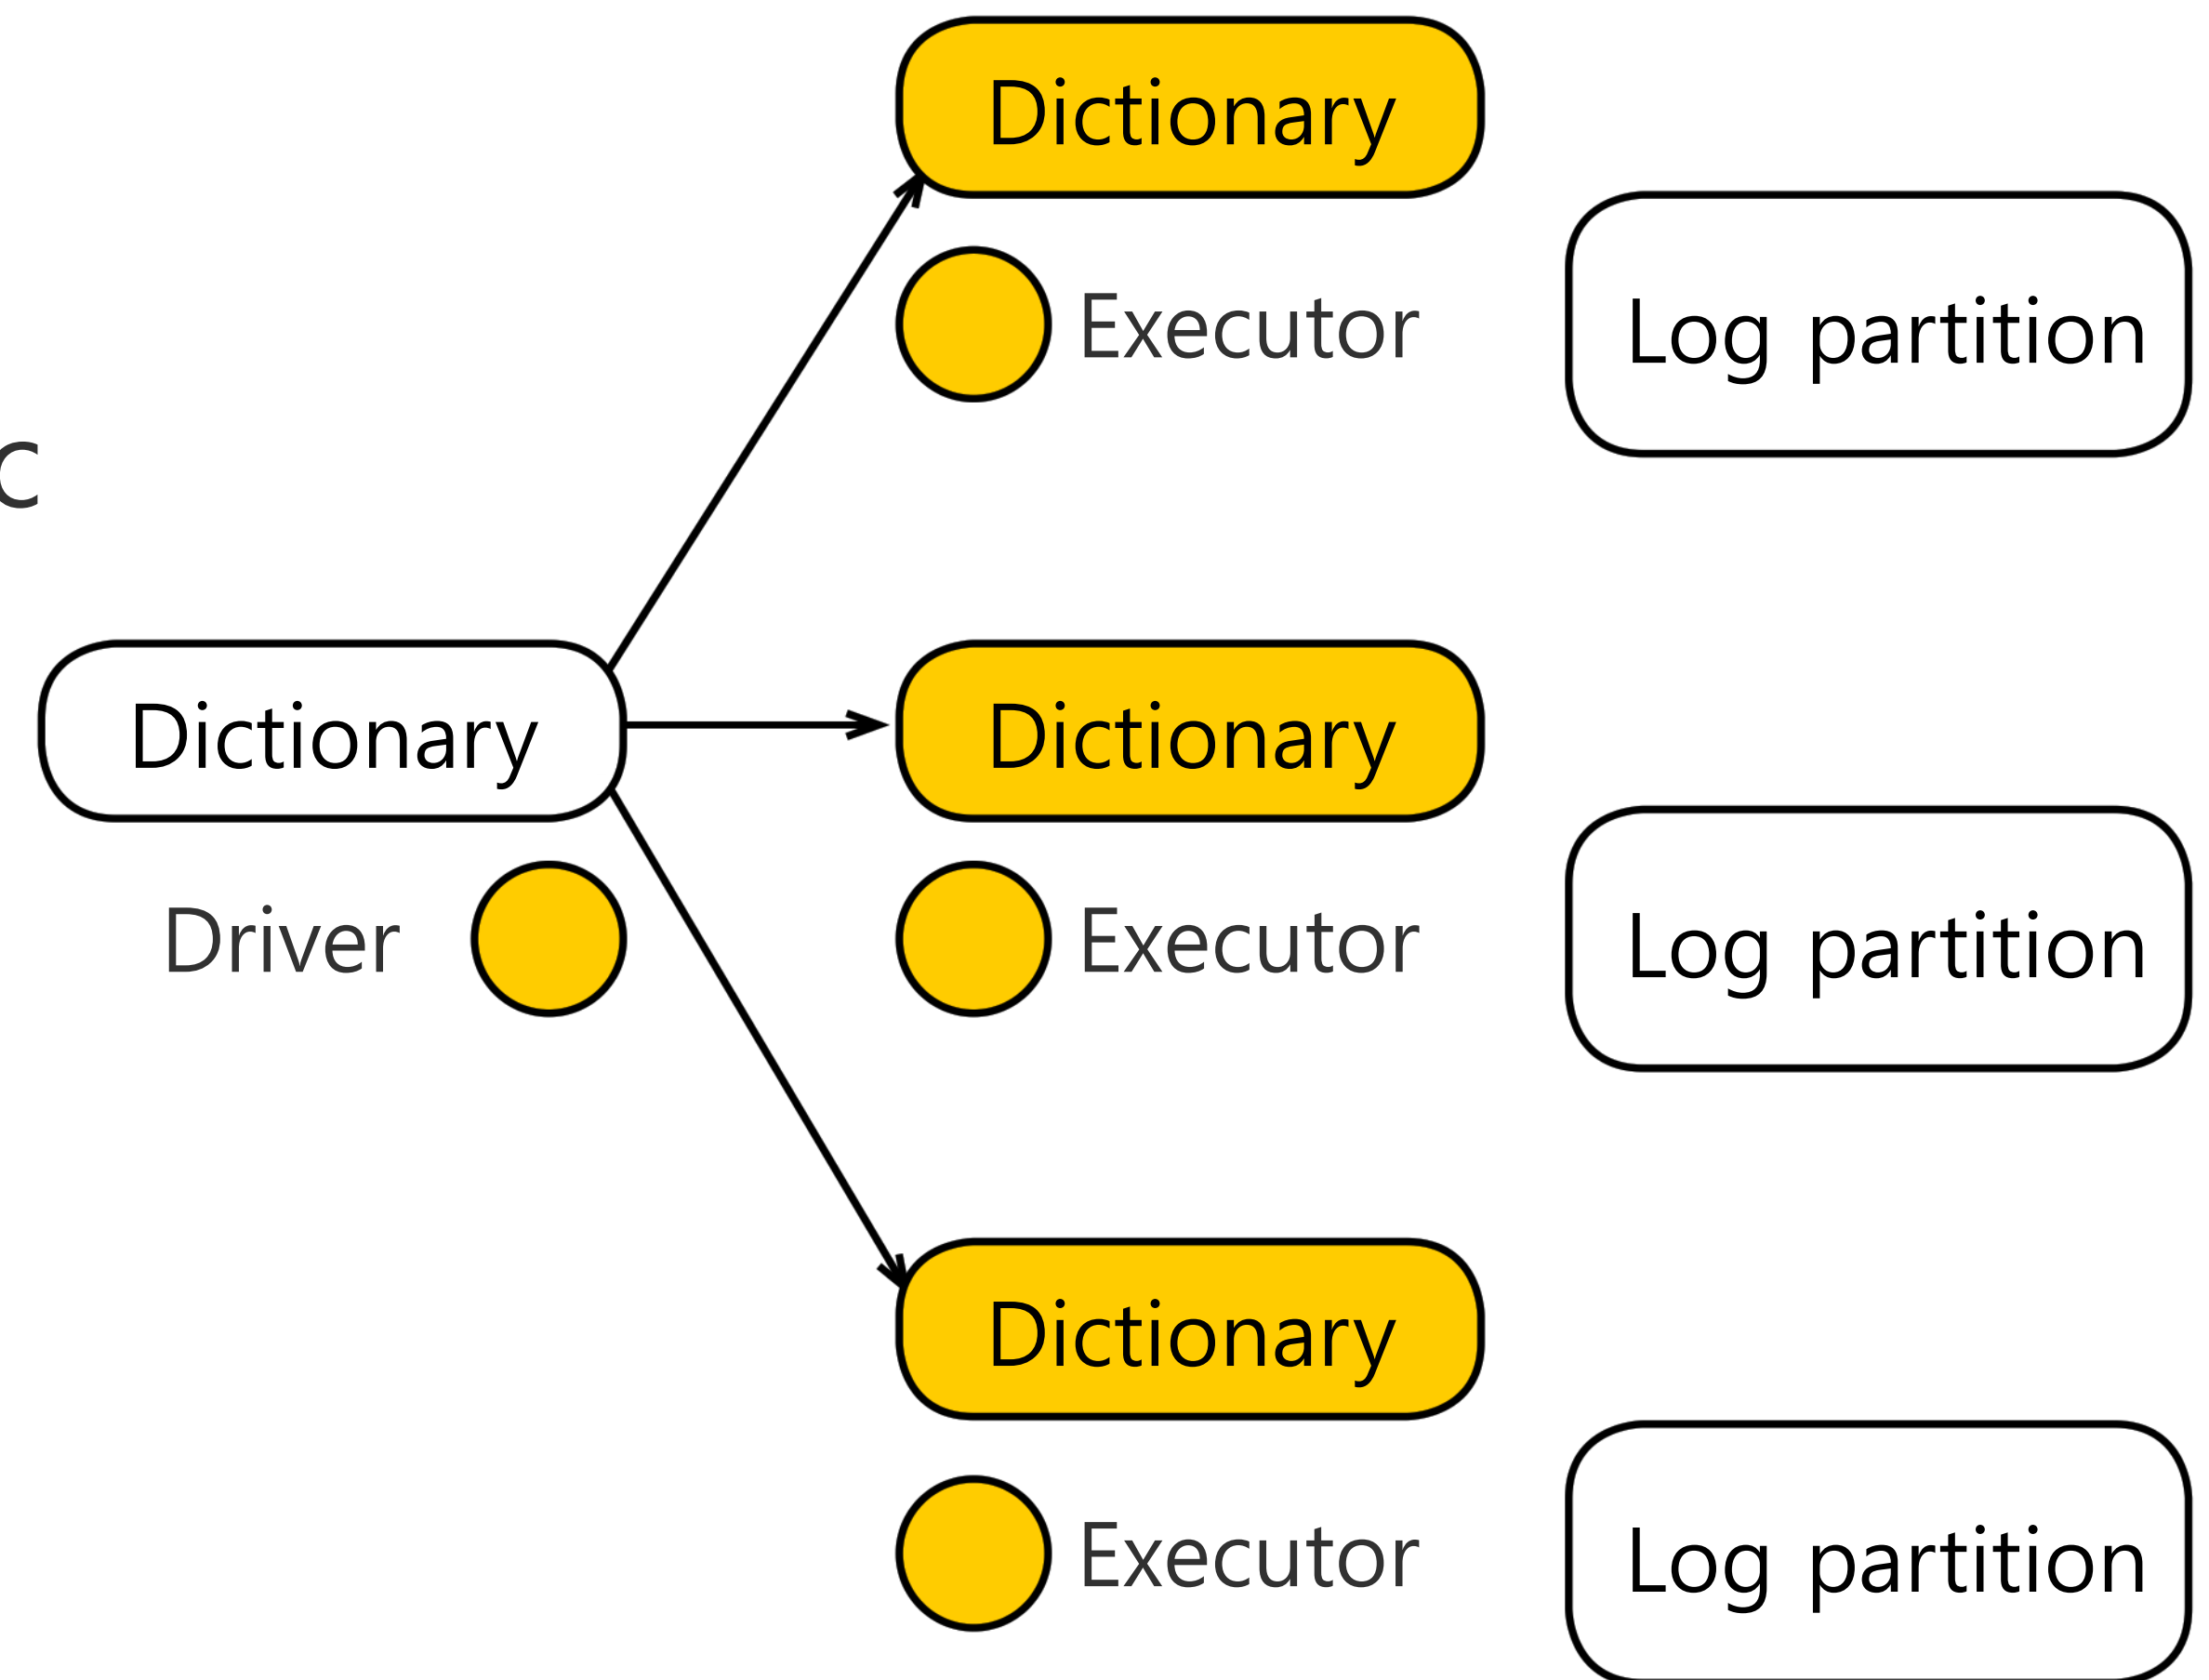
Serial distribution via the closure  
(from the driver to every executor)  
 $\sim 1000 \text{ (tasks)} * 1\text{GB} = 1\text{TB}$  of traffic

1GB 데이터를 driver가 모두 전달

Parallel distribution via  
the broadcast variable  
(torrent-like)

$\sim 1\text{-}2 \text{ GB}$  of traffic

Faster!



# Motivating example 2

```
sc = SparkContext(conf=...)
```

```
# compute the dictionary
```

```
my_dict_rdd = sc.textFile(...).map(...).filter(...)
```

```
my_dict_data = my_dict_rdd.collect()
```

```
# distributed the dictionary via the broadcast variable
```

```
broadcast_var = sc.broadcast(my_dict_data)
```

```
# use the broadcast variable within the task
```

```
my_data_rdd = sc.textFile(...).filter(
```

```
lambda x: x in broadcast_var.value)
```

# Summary

- » Broadcast variables are read-only shared variables with effective sharing mechanism  
단, memory에 맞는 양의 데이터를 활용가능하다.
- » Useful to share dictionaries, models

# Accumulator variables

Useful for the control flow, monitoring, profiling & debugging

# Accumulator variable

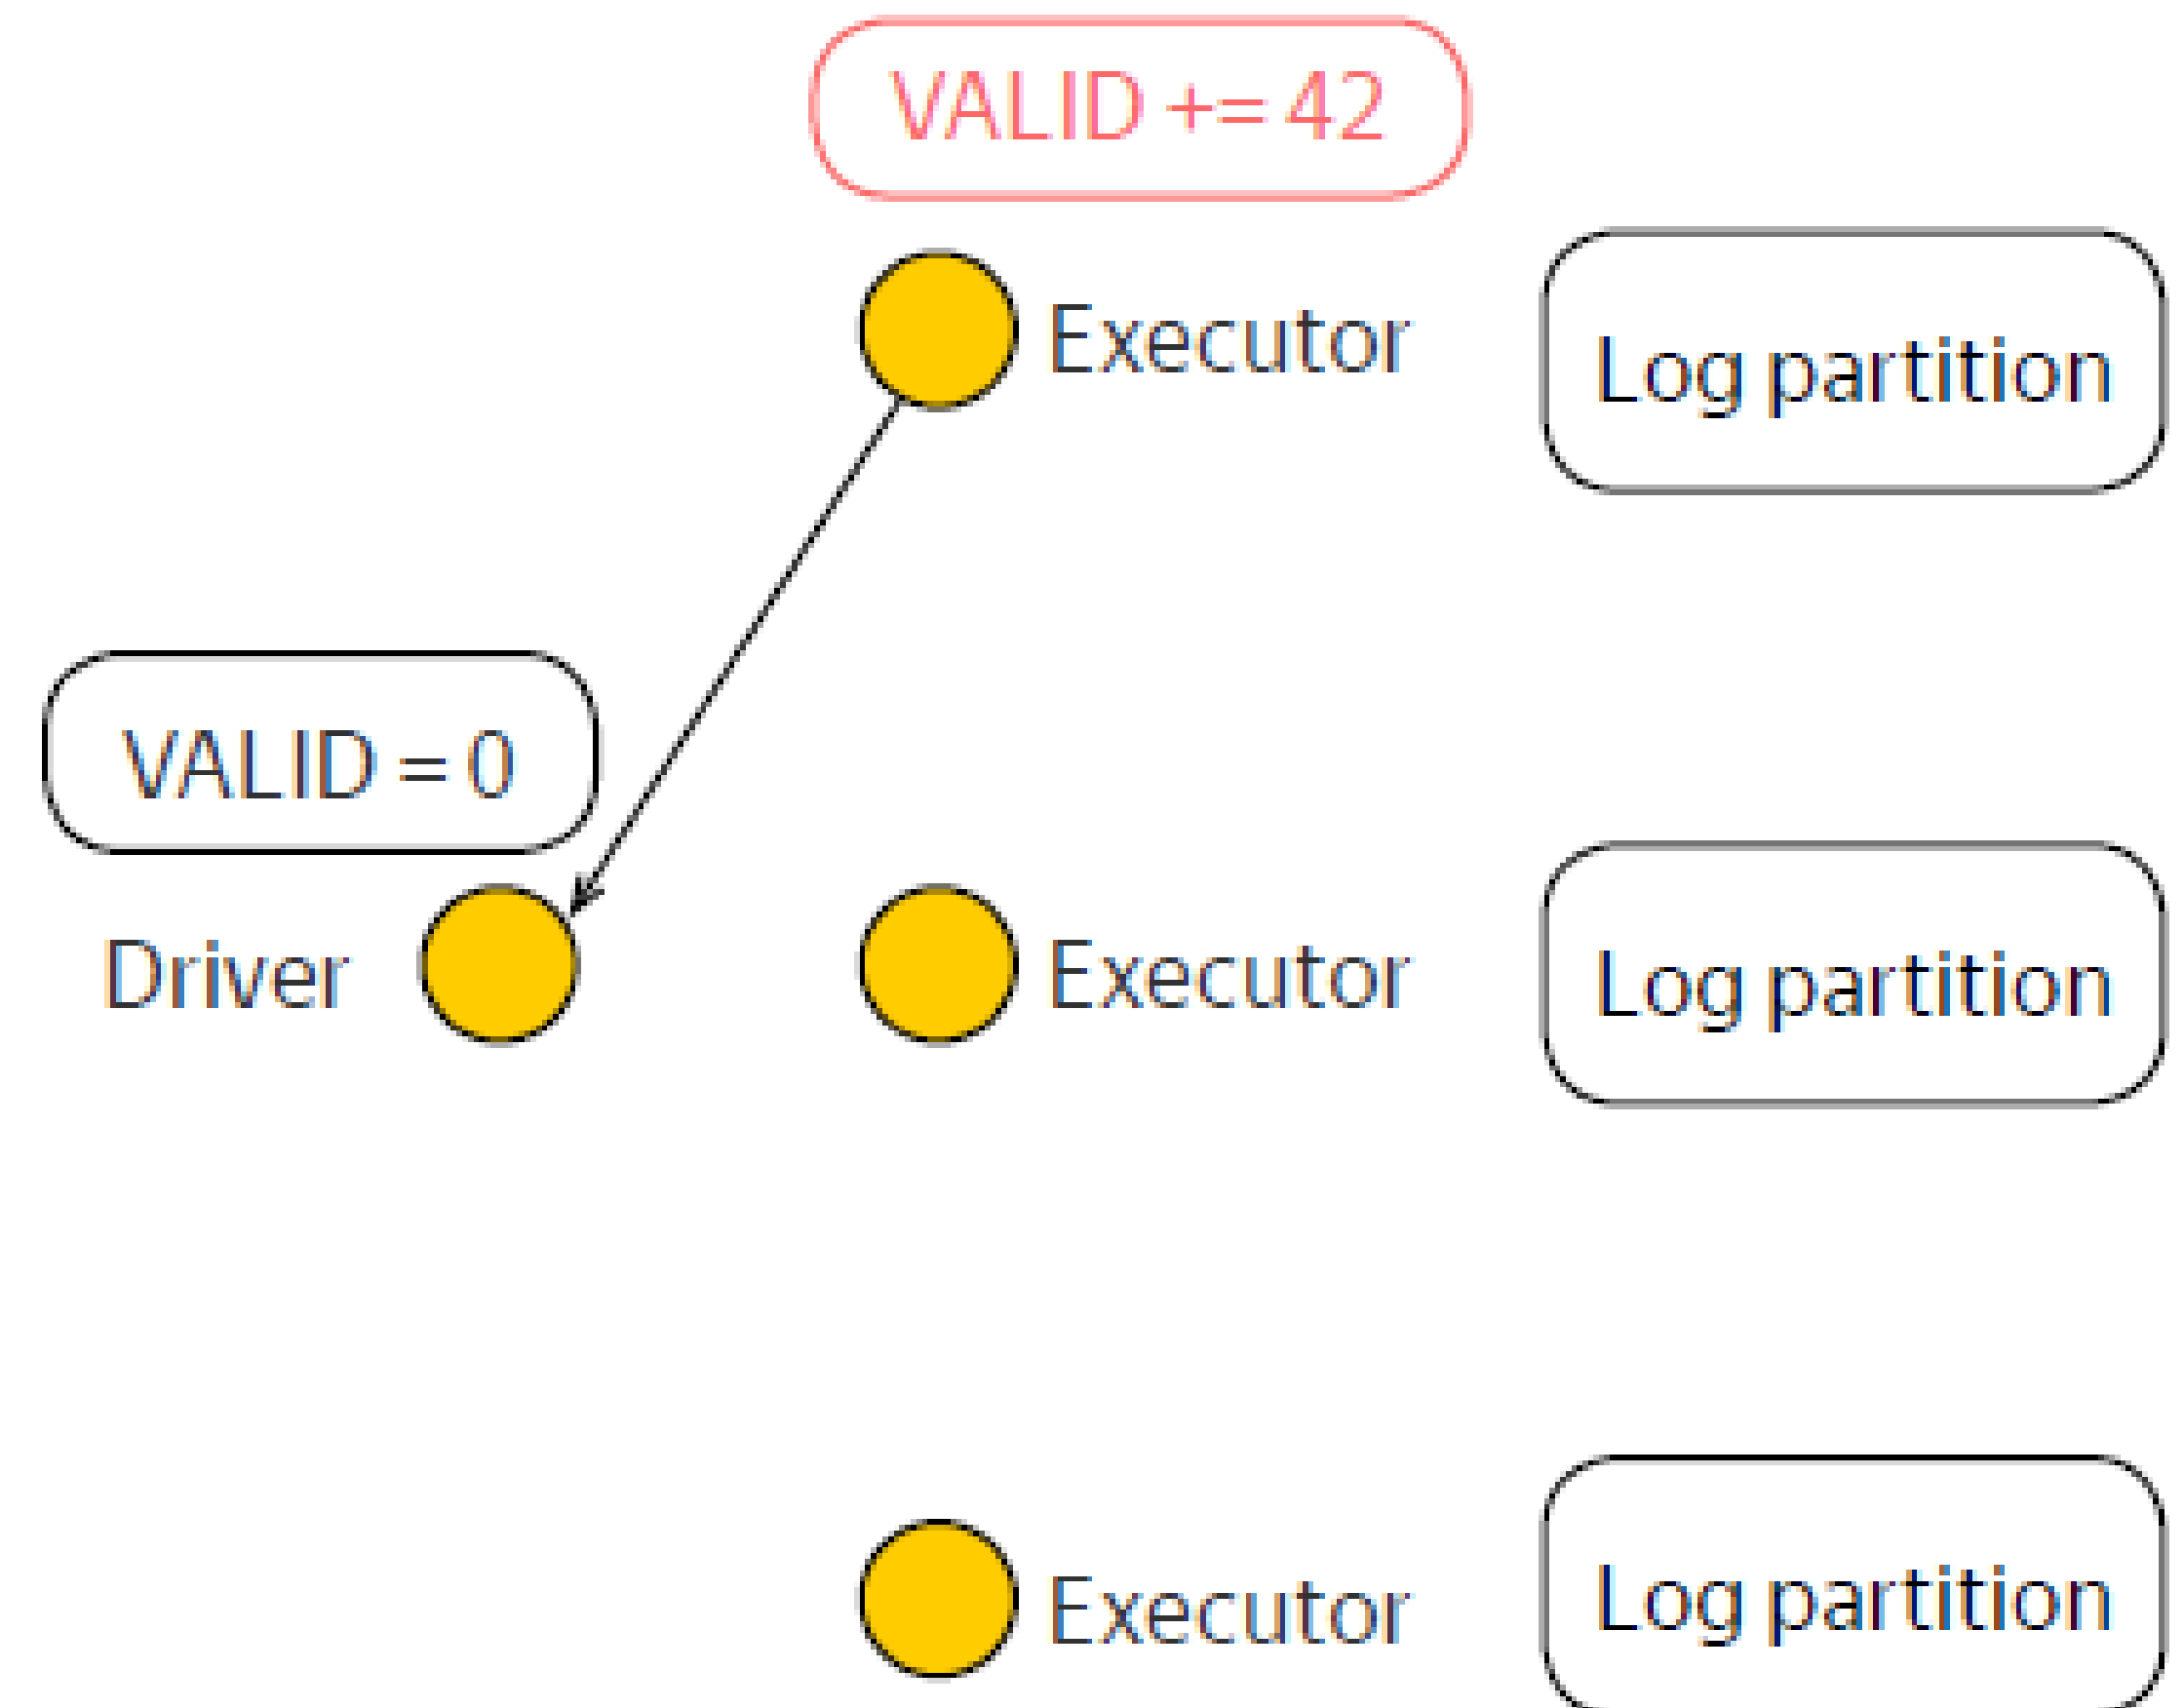
- › Accumulator variable is a read-write variable that is shared among tasks
- ›› Writes are restricted to increments! *synchronization 문제를 피하기 위해*
  - ››i. e.: `var += delta`
  - ››addition may be replaced by any associate, commutative operation
- ›› Reads are allowed *only by the driver program!* *task에서는 읽을 수 없다.*

## Guarantees on the updates

- ›› In actions updates are applied exactly once
  - action에서는 accumulator에 한 번만 적용된다.*
- ›› In transformations there are no guarantees as the transformation code may be re-executed
  - transformation은 재실행될 수 있기 때문에 보장할 수 없다.*

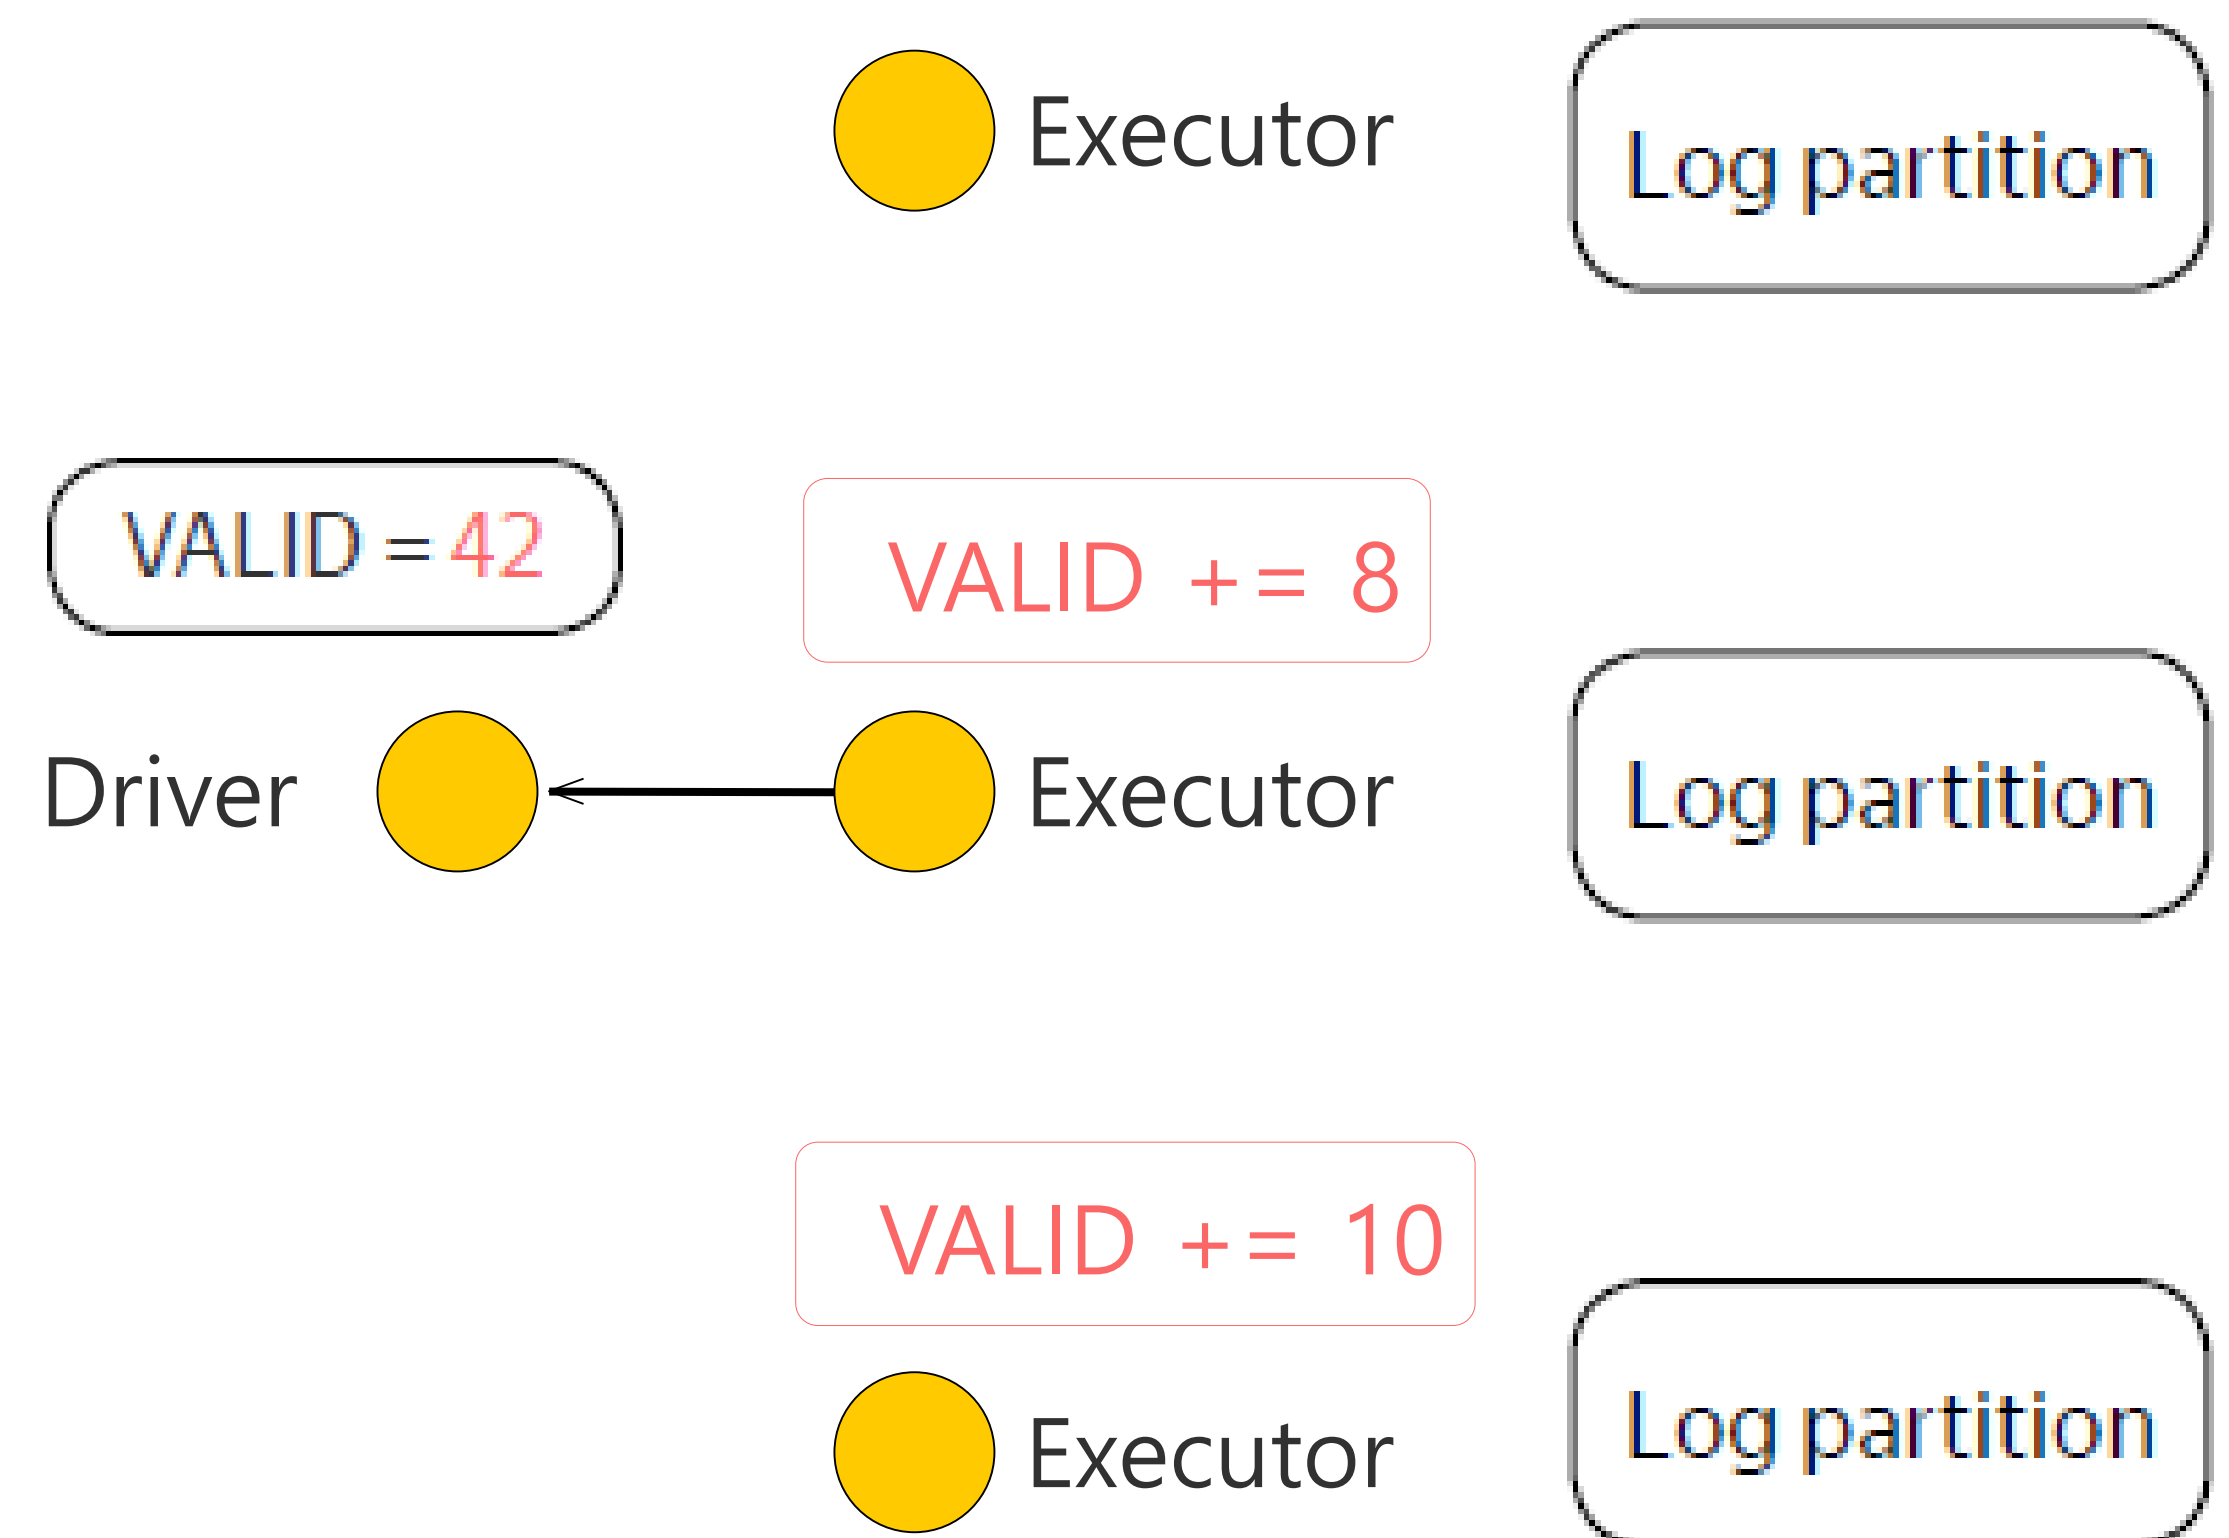
# Example (similar to the previous video)

- › Input:  
1TB partitioned log
- › Task:  
resolve IP addresses  
**AND**  
collect metrics:  
# of valid records



# Example (similar to the previous video)

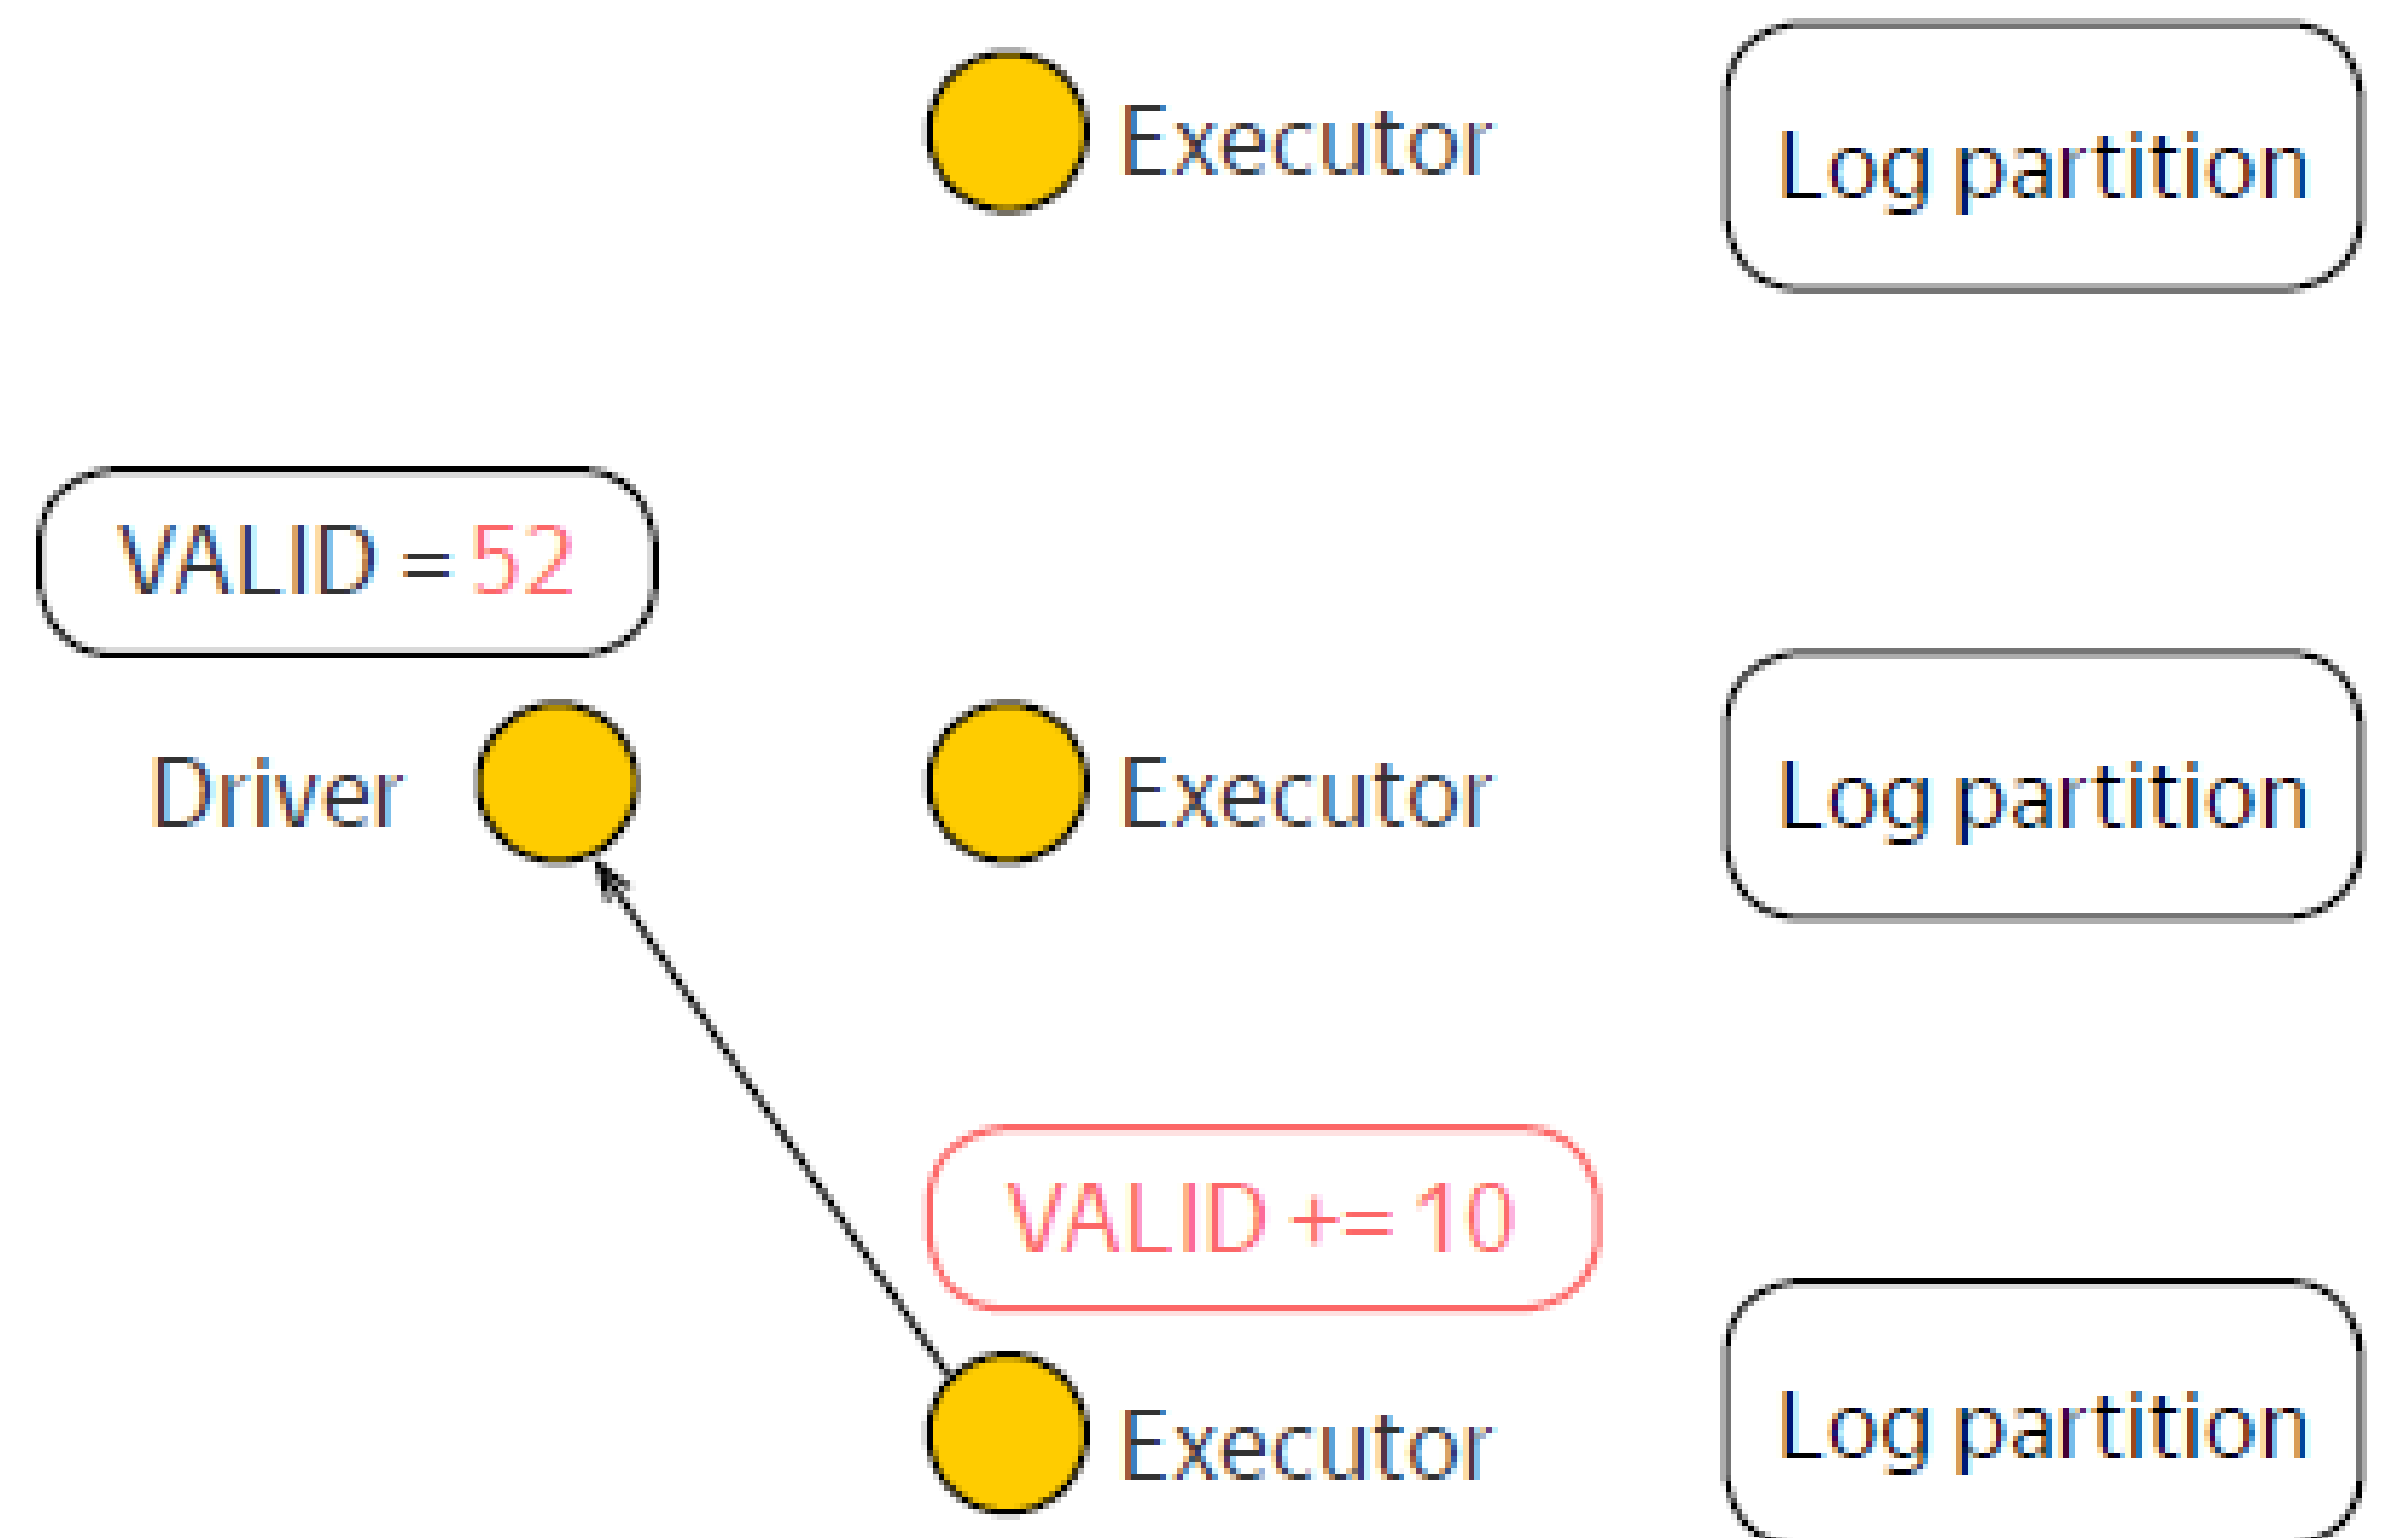
- › Input:  
1TB partitioned log
- › Task:  
resolve IP addresses  
**AND**  
collect metrics:  
# of valid records





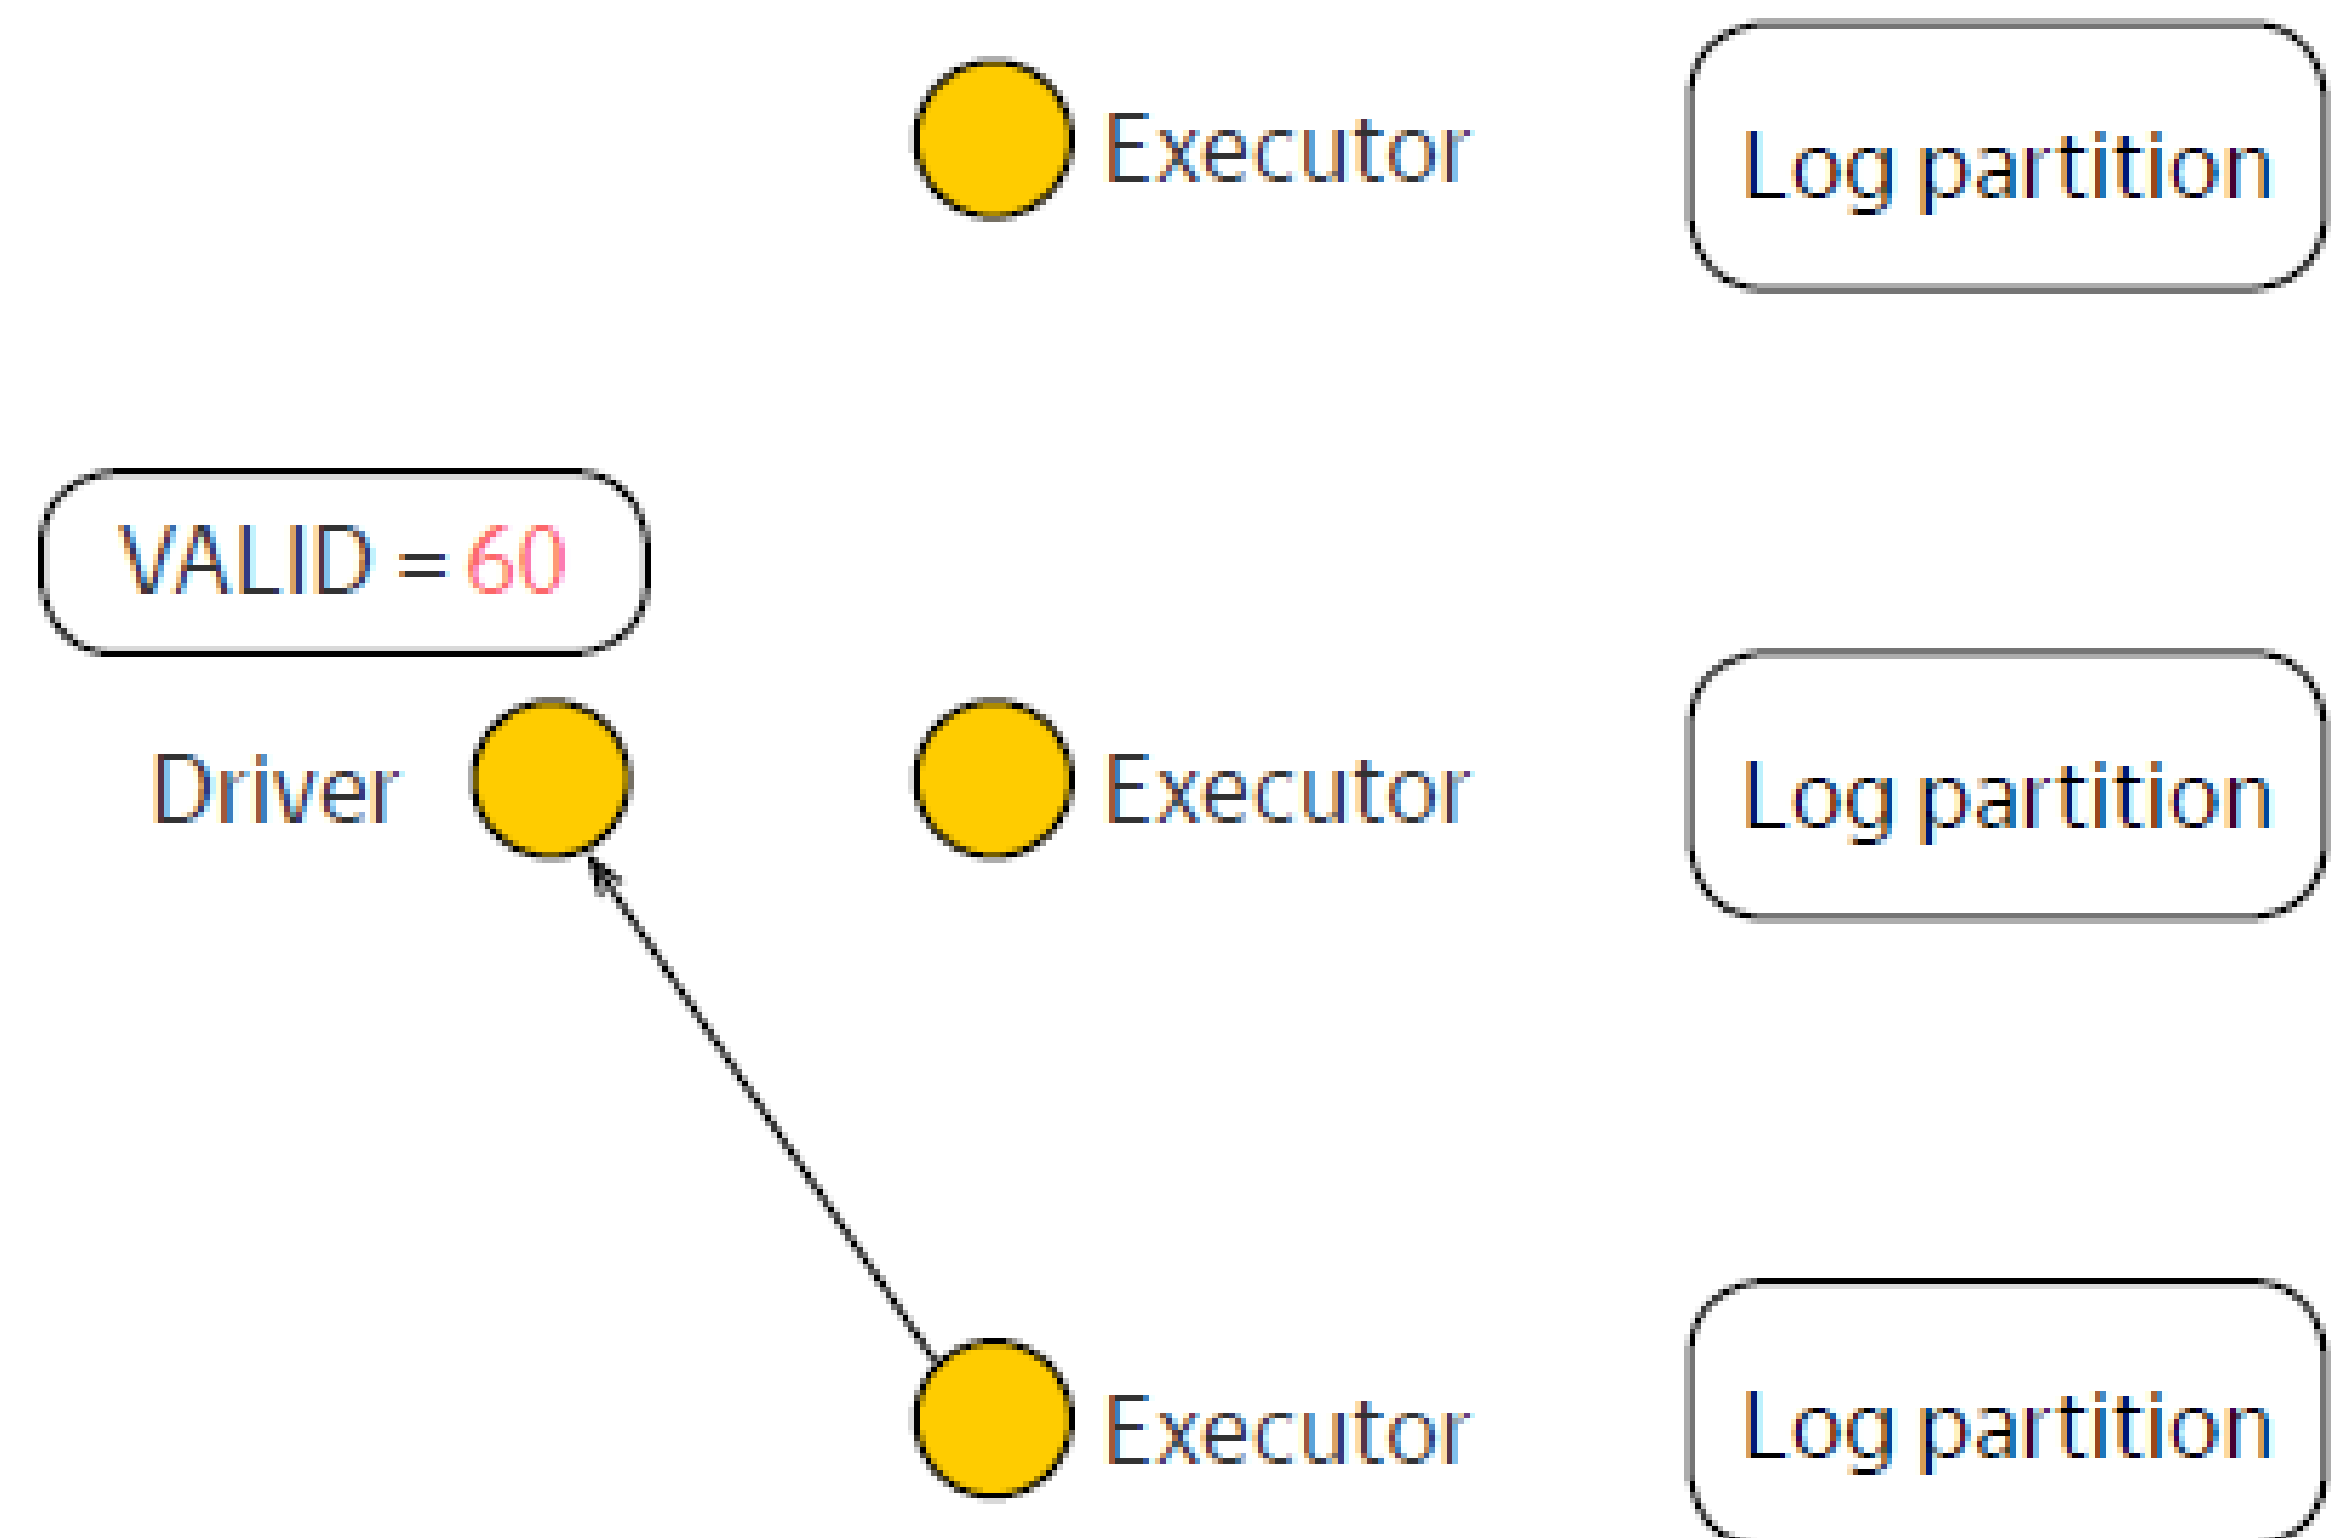
# Example (similar to the previous video)

- › Input:  
1TB partitioned log
- › Task:  
resolve IP addresses  
**AND**  
collect metrics:  
# of valid records



# Example (similar to the previous video)

- › Input:  
1TB partitioned log
- › Task:  
resolve IP addresses  
**AND**  
collect metrics:  
# of valid records



# Use cases

- » Performance counters
  - » # of processed records, total elapsed time, total error and so on and so forth
- » Simple control flow
  - » conditionals: stop on reaching a threshold for corrupted records
  - » loops: decide whether to run the next iteration of an algorithm or not
- » Monitoring
  - » export values to the monitoring system
- » Profiling & debugging

# Summary

- » Accumulators are read-write shared variables with restricted updates
  - » increments only
  - » can use custom associative, commutative operation for the updates
  - » can read the total value only in the driver
- » Useful for the control flow, monitoring, profiling & debugging

# Getting started with Spark & Python

# Installing Spark locally

- › Navigate to <http://spark.apache.org/docs/latest/#downloading> and follow the instructions
- › At the time of making this video
  - › download .tar.gz
  - › extract it
  - › run `./bin/pyspark` from the extracted directory
- › If you have IPython installed, you can run `PYSPARK_DRIVER_PYTHON=ipython pyspark`

## ./bin/pyspark

Python 2.7.10 (default, Feb 6 2017, 23:53:20)

[GCC 4.2.1 Compatible Apple LLVM 8.0.0 (clang-800.0.34)] on darwin

Type "help", "copyright", "credits" or "license" for more information.

Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties

Setting default log level to "WARN".

To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).

Welcome to

```
 _ _ _ _ _  
/  _  _  _  \\  
\\V_\\_/_/_/_/  
/_/. _\\_/_/_/_/_\\ version 2.2.0  
/_/_
```

Using Python version 2.7.10 (default, Feb 6 2017 23:53:20)

SparkSession available as 'spark'.

```
>>> sc
```

```
<SparkContext master=local[*] appName=PySparkShell>
```

spark shell에서는 sc가 자동으로 생성된다.

```
>>>
```

## SparkContext 생성 방법

```
from pyspark import SparkConf, SparkContext
```

```
sc = SparkContext(conf=SparkConf().setAppName("MyApp").setMaster("local"))
```

### Master 설정 방법 ( Cluster Mode )

Possible values for the master URL:

**local[K]** — local mode with K threads

**spark://HOST:PORT** — standalone Spark cluster

**mesos://HOST:PORT** — Mesos cluster

**yarn** — YARN cluster

# Spark UI



## Spark Jobs (?)

**User:** sandello

**Total Uptime:** 10 min

**Scheduling Mode:** FIFO

▶ [Event Timeline](#)

## Spark Jobs [\(?\)](#)

**User:** sandello

**Total Uptime:** 15 min

**Scheduling Mode:** FIFO

**Completed Jobs:** 1

[▶ Event Timeline](#)

### Completed Jobs (1)


Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	<a href="#">top at &lt;ipython-input-15-0b58397eb290&gt;:1</a>	2017/07/23 22:16:13	2 s	2/2	<div><div>8/8</div></div>

Details for Job 0

Status: SUCCEEDED

Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



Completed Stages (2)

Stage Id ▾	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
1	<a href="#">top at &lt;ipython-input-15-0b58397eb290&gt;:1</a>	<a href="#">+details</a>	2017/07/23 22:16:15	0,1 s	4/4			7.3 KB	
0	<a href="#">join at &lt;ipython-input-10-9c257a414347&gt;:1</a>	<a href="#">+details</a>	2017/07/23 22:16:13	2 s	4/4	23.0 KB			7.3 KB

## Details for Stage 0 (Attempt 0)

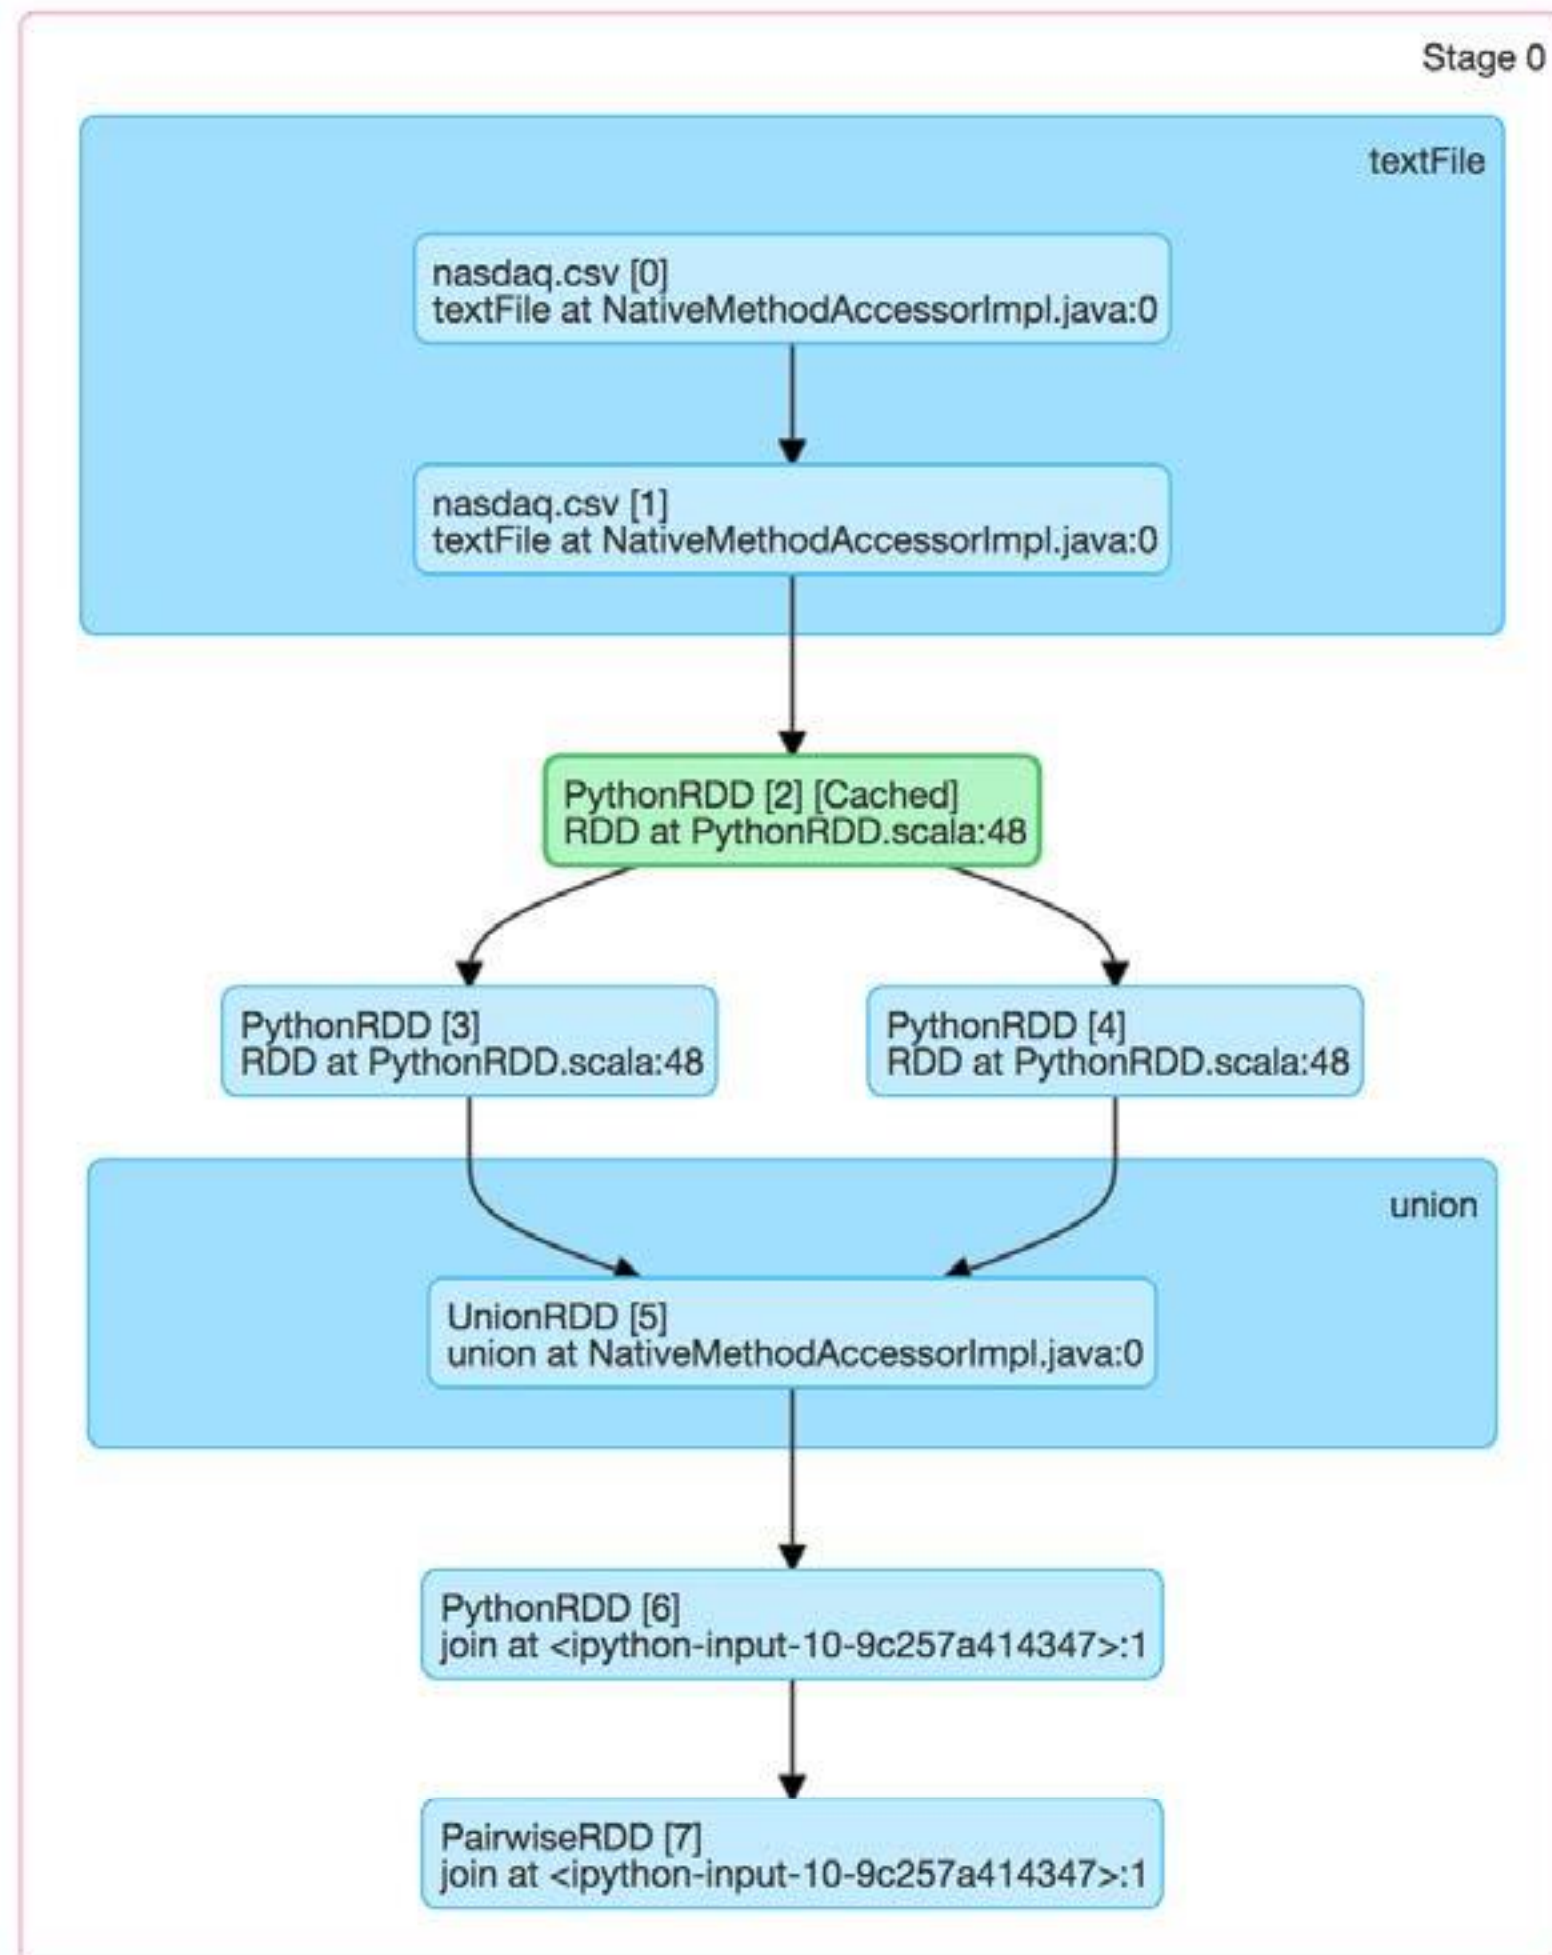
Total Time Across All Tasks: 7 s

Locality Level Summary: Process local: 4

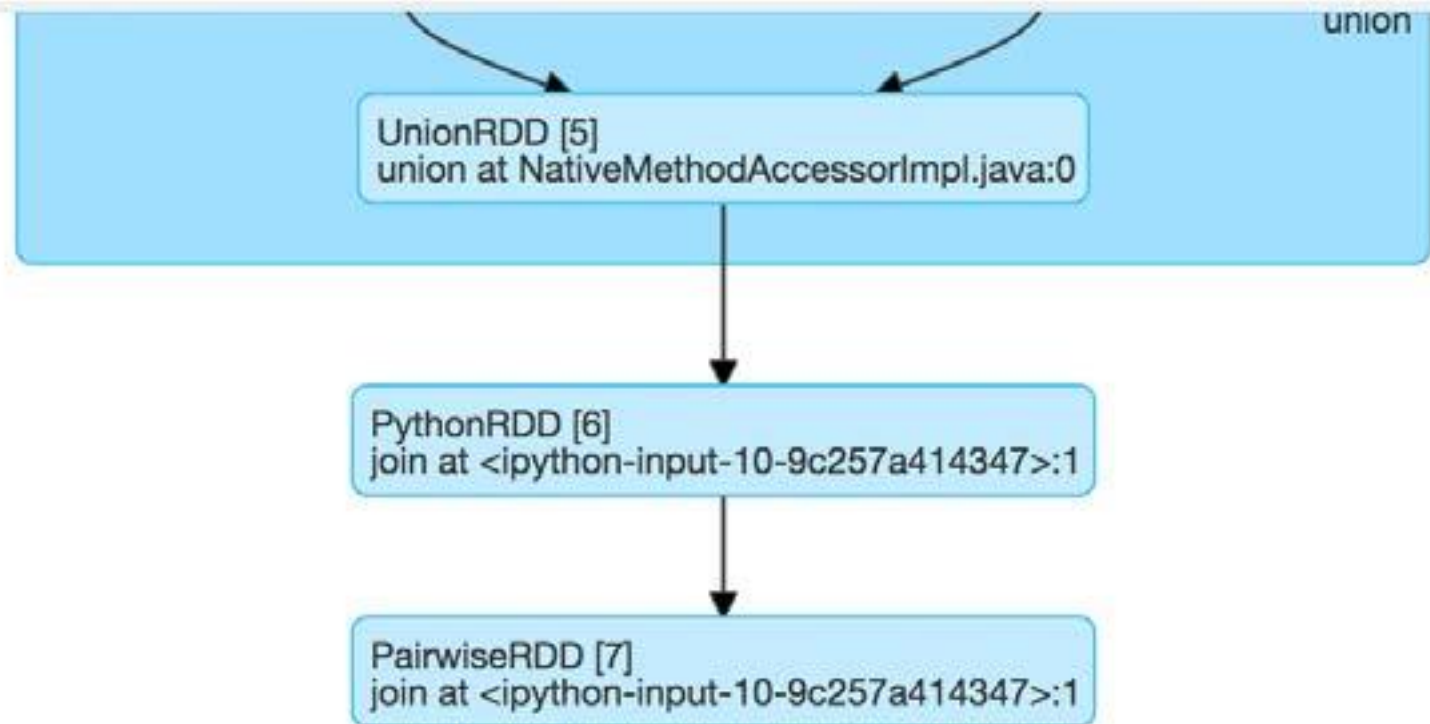
Input Size / Records: 23.0 KB / 153

Shuffle Write: 7.3 KB / 32

▼ DAG Visualization







- ▶ Show Additional Metrics
- ▶ Event Timeline

Summary Metrics for 4 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	2 s	2 s	2 s	2 s	2 s
GC Time	0 ms	0 ms	0 ms	0 ms	0 ms
Input Size / Records	3.1 KB / 7	3.1 KB / 7	5.6 KB / 69	11.1 KB / 70	11.1 KB / 70
Shuffle Write Size / Records	1849.0 B / 8	1856.0 B / 8	1865.0 B / 8	1867.0 B / 8	1867.0 B / 8

▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks	Input Size / Records	Shuffle Write Size / Records	Blacklisted
driver	77.88.19.2:53626	8 s	4	0	0	4	23.0 KB / 153	7.3 KB / 32	0

Tasks (4)

Index ▲	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Input Size / Records	Write Time	Shuffle Write Size / Records	Errors
0	0	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2017/07/23 22:16:13	2 s		11.1 KB / 70	9 ms	1865.0 B / 8	
1	1	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2017/07/23 22:16:13	2 s		3.1 KB / 7	20 ms	1856.0 B / 8	
2	2	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2017/07/23 22:16:13	2 s		3.1 KB / 7	36 ms	1849.0 B / 8	
3	3	0	SUCCESS	PROCESS_LOCAL	driver / localhost	2017/07/23 22:16:13	2 s		5.6 KB / 69	10 ms	1867.0 B / 8	

## Storage

### RDDs

RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size on Disk
<a href="#">PythonRDD</a>	Memory Serialized 1x Replicated	2	100%	6.3 KB	0.0 B

## RDD Storage Info for PythonRDD

**Storage Level:** Memory Serialized 1x Replicated

**Cached Partitions:** 2

**Total Partitions:** 2

**Memory Size:** 6.3 KB

**Disk Size:** 0.0 B

### Data Distribution on 1 Executors

Host	On Heap Memory Usage	Off Heap Memory Usage	Disk Usage
77.88.19.2:53626	6.3 KB (366.3 MB Remaining)	0.0 B (0.0 B Remaining)	0.0 B

### 2 Partitions

Block Name ▲	Storage Level	Size in Memory	Size on Disk	Executors
rdd_2_0	Memory Serialized 1x Replicated	3.1 KB	0.0 B	77.88.19.2:53626
rdd_2_1	Memory Serialized 1x Replicated	3.1 KB	0.0 B	77.88.19.2:53626



## Executors

[Show Additional Metrics](#)

### Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Blacklisted
<b>Active(1)</b>	3	29.9 KB / 384.1 MB	0.0 B	4	0	0	8	8	8 s (0 ms)	23.5 KB	0.0 B	7.4 KB	0
<b>Dead(0)</b>	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	0
<b>Total(1)</b>	3	29.9 KB / 384.1 MB	0.0 B	4	0	0	8	8	8 s (0 ms)	23.5 KB	0.0 B	7.4 KB	0

### Executors

Show  entries

Search: 

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	77.88.19.2:53626	Active	3	29.9 KB / 384.1 MB	0.0 B	4	0	0	8	8	8 s (0 ms)	23.5 KB	0.0 B	7.4 KB	<a href="#">Thread Dump</a>

Showing 1 to 1 of 1 entries

Previous

1

Next



# Summary

- » You have learned how to open and use Spark UI
- » In the next course of the specialization you will learn how to use the interface to optimize your application performance

**BigDATA**team