


Yandex

Telecommunications Analytics

Map and Reduce Side Joins

Telecommunications - SMS, Call, Internet - MI



Milano Grid

- » Square ID
- » Time Interval
- » Country Code
- » SMS-in Activity
- » SMS-out Activity
- » Call-in Activity
- » Call-out Activity
- » Internet Traffic Activity

Schema

해결할 문제 : Square ID를 South와 North 지역으로 구분해 SMS-in activity 합을 구한다.

BIG data

로그 데이터


- ›› Square ID
- ›› Time Interval
- ›› Country Code
- ›› SMS-in Activity
- ›› SMS-out Activity
- ›› Call-in Activity
- ›› Call-out Activity
- ›› Internet Traffic Activity

```
1 1383260400000 0 0.08136262351125882
1 1383260400000 39 0.14186425470242922
0.1567870050390246 0.16093793691701822
0.052274848528573205 11.028366381681026
1 1383261000000 0 0.13658782275823106
0.02730046487718618
1 1383261000000 33
0.026137424264286602
...
```

small data

Square ID에 대한 위치정보

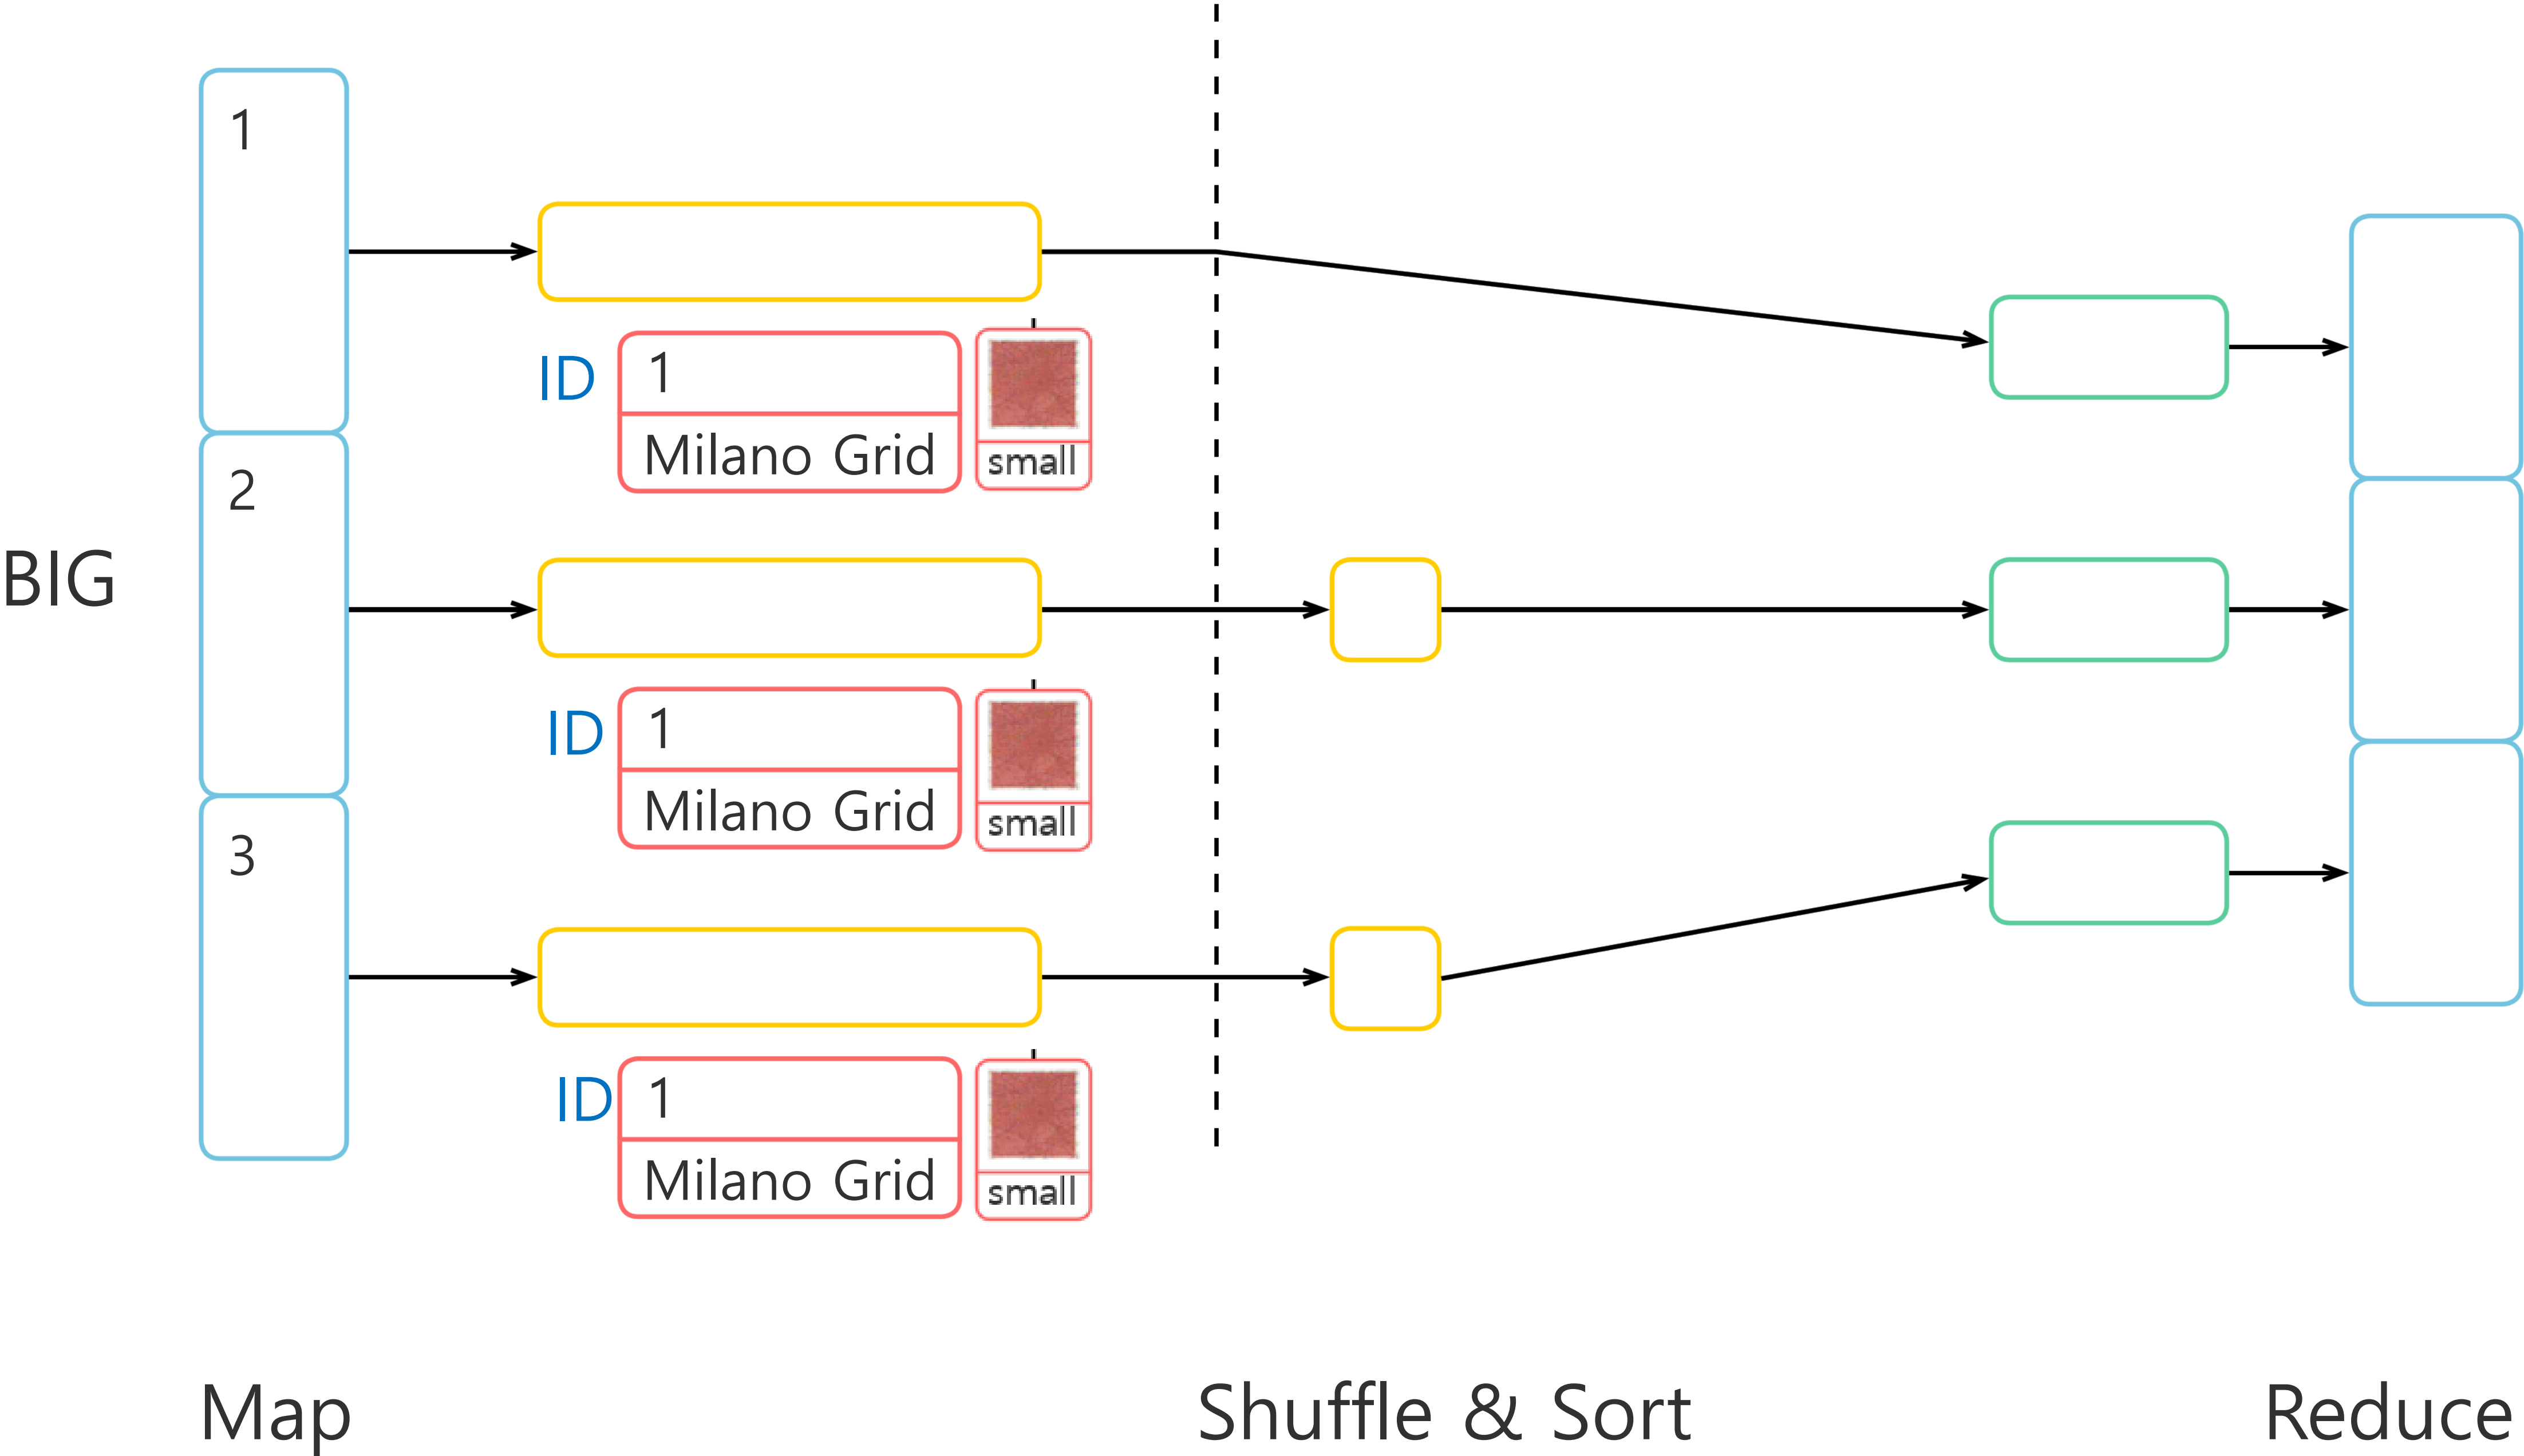
지역을 구분하기 위한 Join이 필요



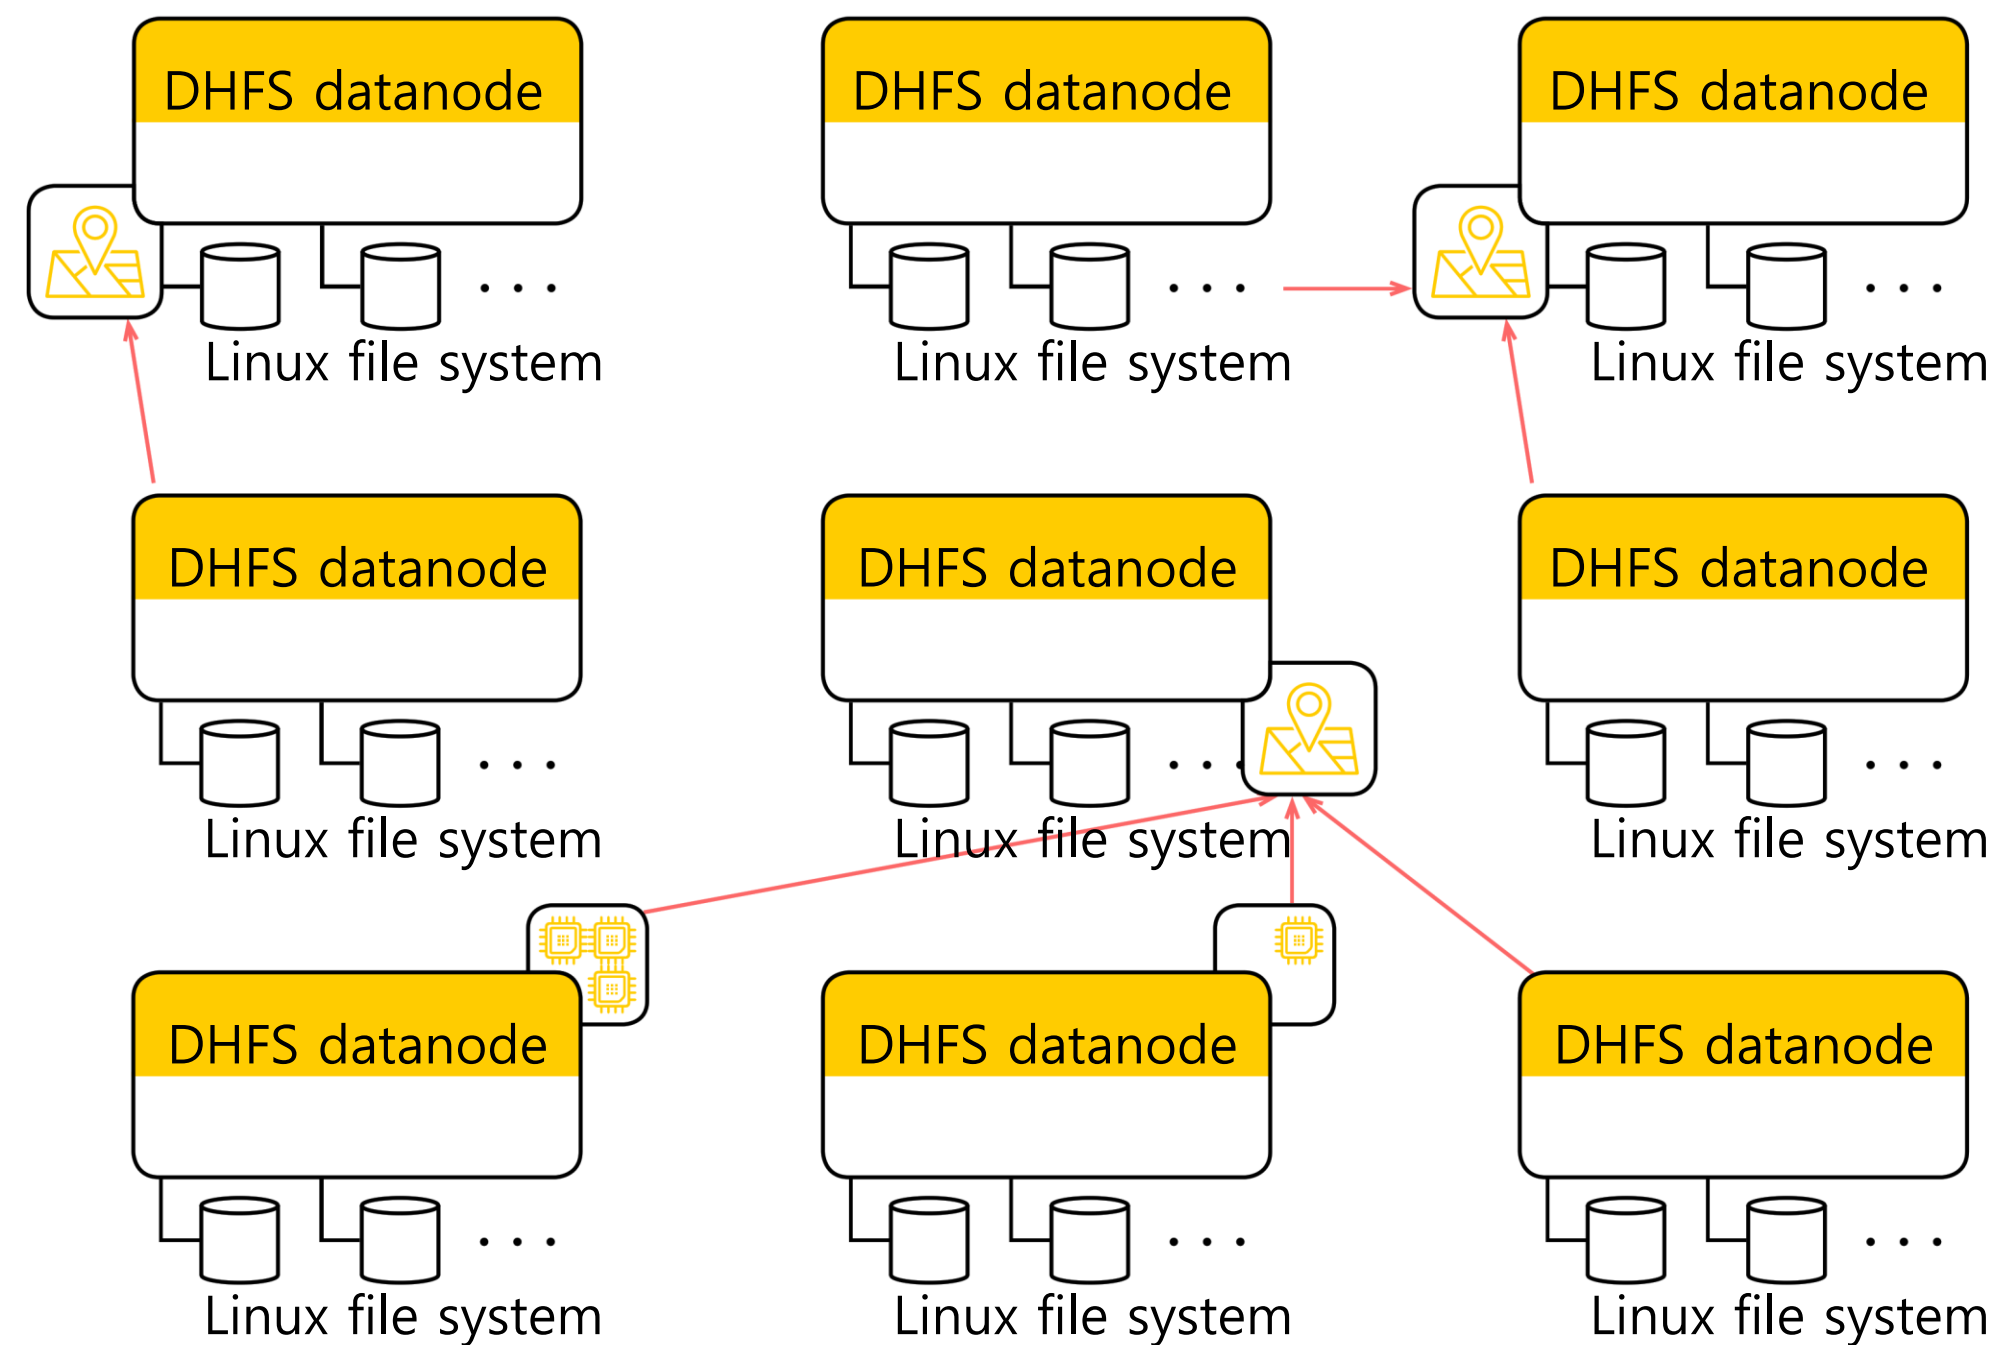
```
{'type': 'Polygon', 'coordinates':
[[[9.0114910478323, 45.35880131440966],
[9.014491488013135, 45.35880097314403],
[9.0144909480813, 45.35668565341486],
[9.011490619692509,
45.356685994655464], [9.0114910478323,
45.35880131440966]]]}
...
```

Map-Side Join

Mapper에서 데이터를 메모리로 읽어 join 하는 경우

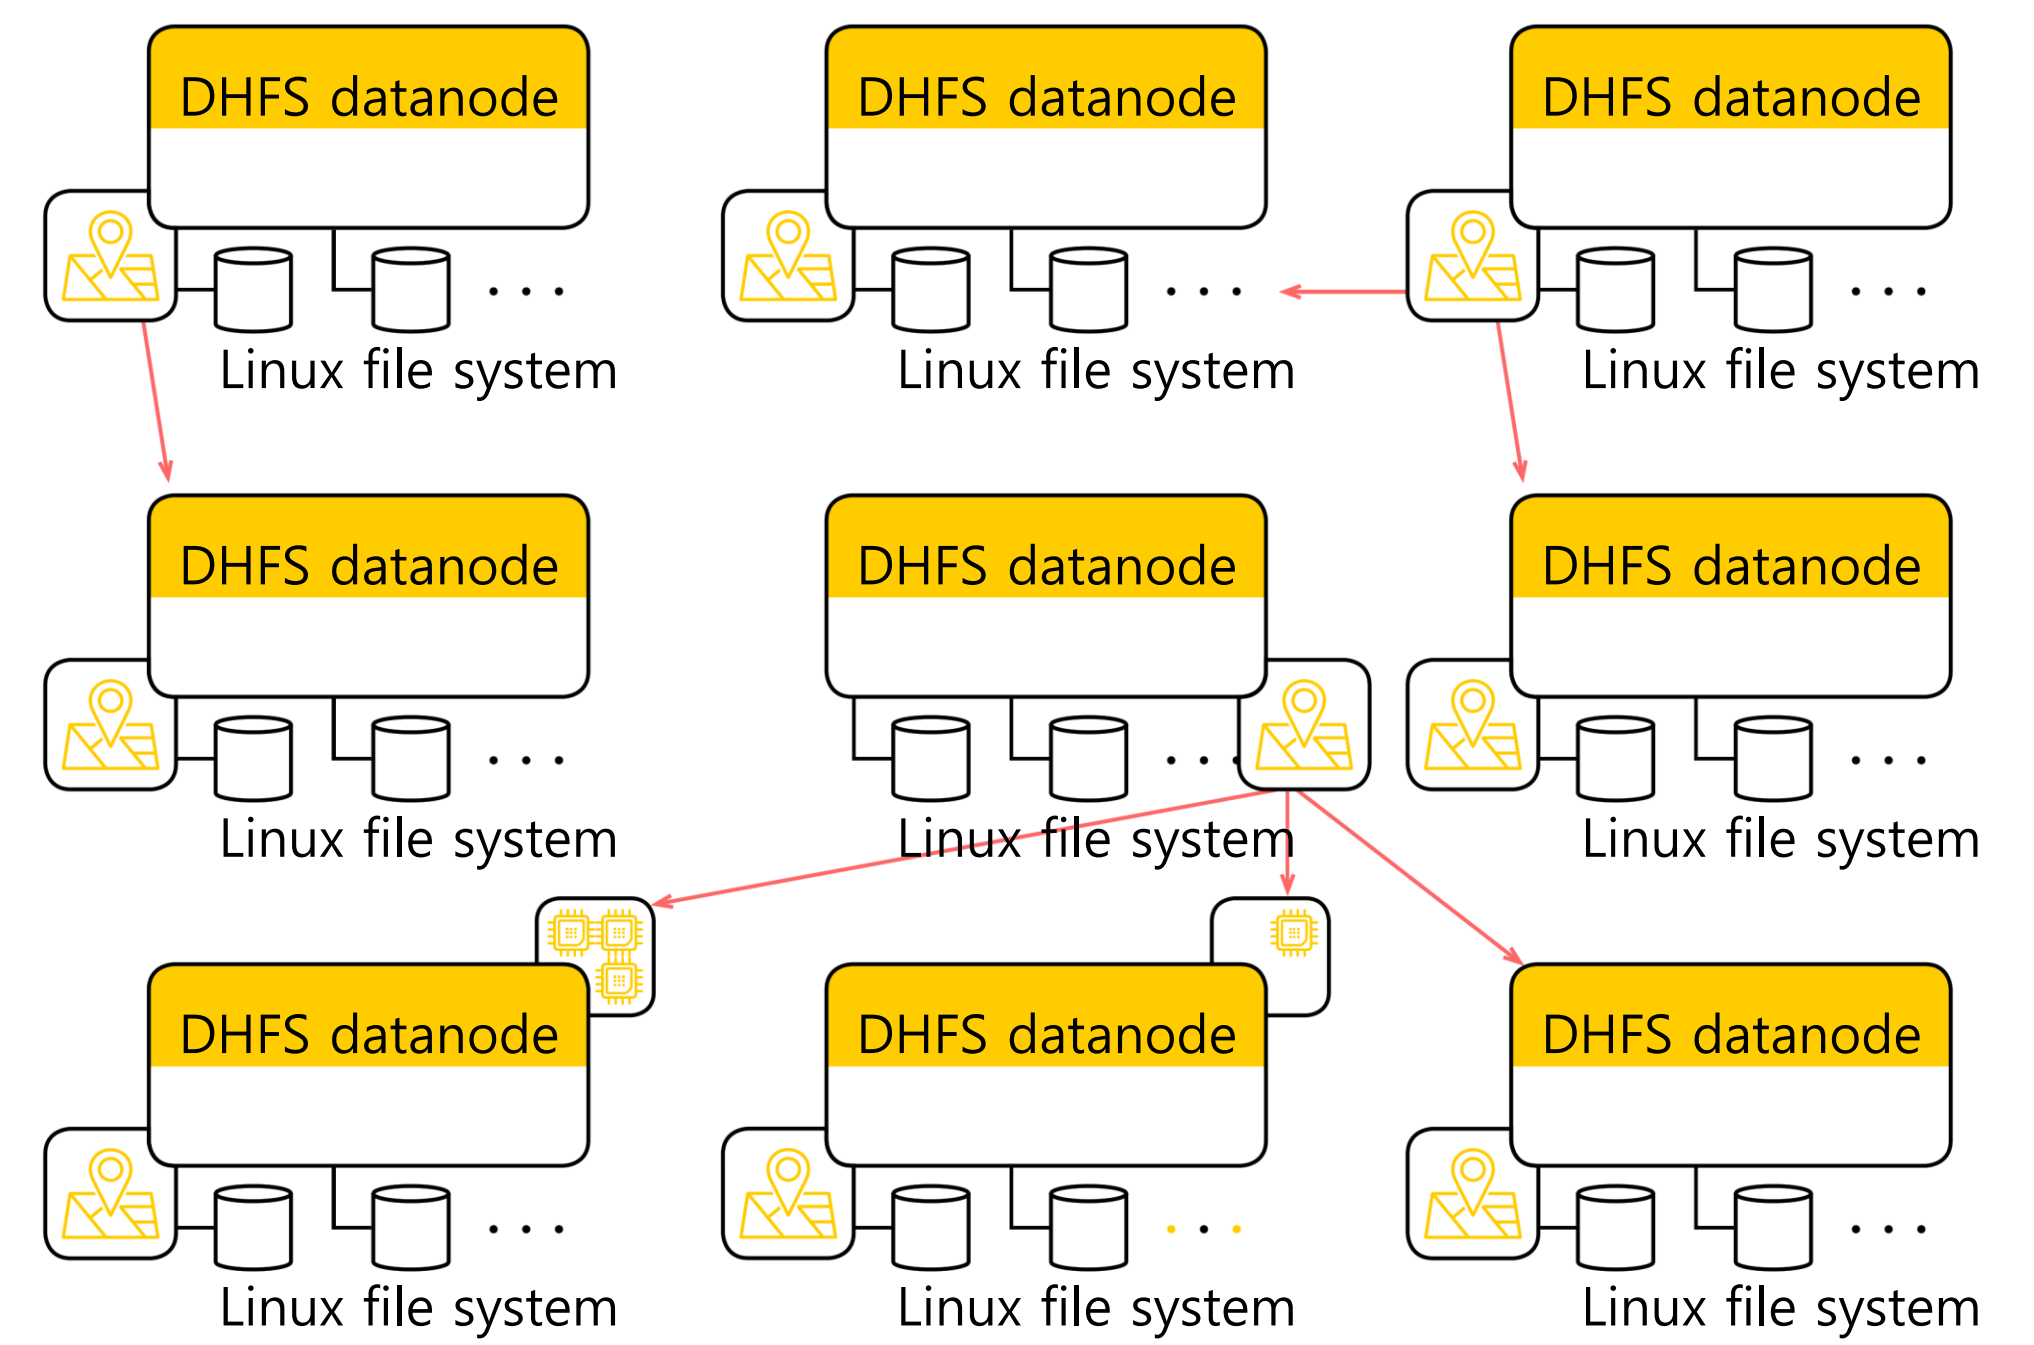


HDFS Read



데이터가 HDFS에 저장되어 있는 경우, 3개의 replica에서 모든 mapper로 데이터를 전달해야 한다. => 비효율적

Distributed Cache



Distributed Cache를 활용하면 각 machine에 한 번만 네트워크로 보내면 mapper가 서로 공유하기 때문에 효율적이다.

```
yarn jar $HADOOP_STREAMING_JAR ₩
  -files read_from_hdfs_mapper.py, hdfs:///user/adral/milano-grid.geojson ₩
  -mapper 'python map_side_mapper.py' ₩
  -numReduceTasks 0 ₩
  -input /data/telecommunication ₩
  -output telecom-joins
```

```
$ hdfs dfs -cat telecom-joins/part-00000 | head
South 0.0813626235113
South 0.141864254702
South 0.136587822758
...
```

map_side_mapper.py

```
geojson = json.load(open("milano-grid.geojson"))
grid = load_grid(geojson)
for line in sys.stdin:
    square_id, aggregate = line.split("\t", 1)
    square_id = int(square_id)
    time_interval, country, sms_in, sms_out, call_in, call_out, internet = aggregate.split("\t")
    if sms_in:
        sms_in = float(sms_in)
        print(grid[square_id], sms_in, sep="\t")
```

from load_grid function load_grid에서 north or south 지역을 계산

grid[square_id] = "North" if (min_y + max_y) / 2 > middle_point else "South"

HDFS read

```
Job Counters
  Launched map tasks=10
  Data-local map tasks=10
  Total time spent by all maps in occupied slots (ms)=311034
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=155517
  Total vcore-seconds taken by all map tasks=155517
  Total megabyte-seconds taken by all map tasks=636997632
```

local read; Distributed Cache


```
Job Counters
  Launched map tasks=10
  Data-local map tasks=10
  Total time spent by all maps in occupied slots (ms)=221296
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=110648
  Total vcore-seconds taken by all map tasks=110648
  Total megabyte-seconds taken by all map tasks=453214208
```


grid 데이터가 BIG 데이터인 경우, mapper 메모리에 데이터를 저장할 수 없어서 Map-Side Join은 사용불가

BIG

- » Square ID
- » Time Interval
- » Country Code
- » SMS-in Activity
- » SMS-out Activity
- » Call-in Activity
- » Call-out Activity
- » Internet Traffic Activity

small



BIG

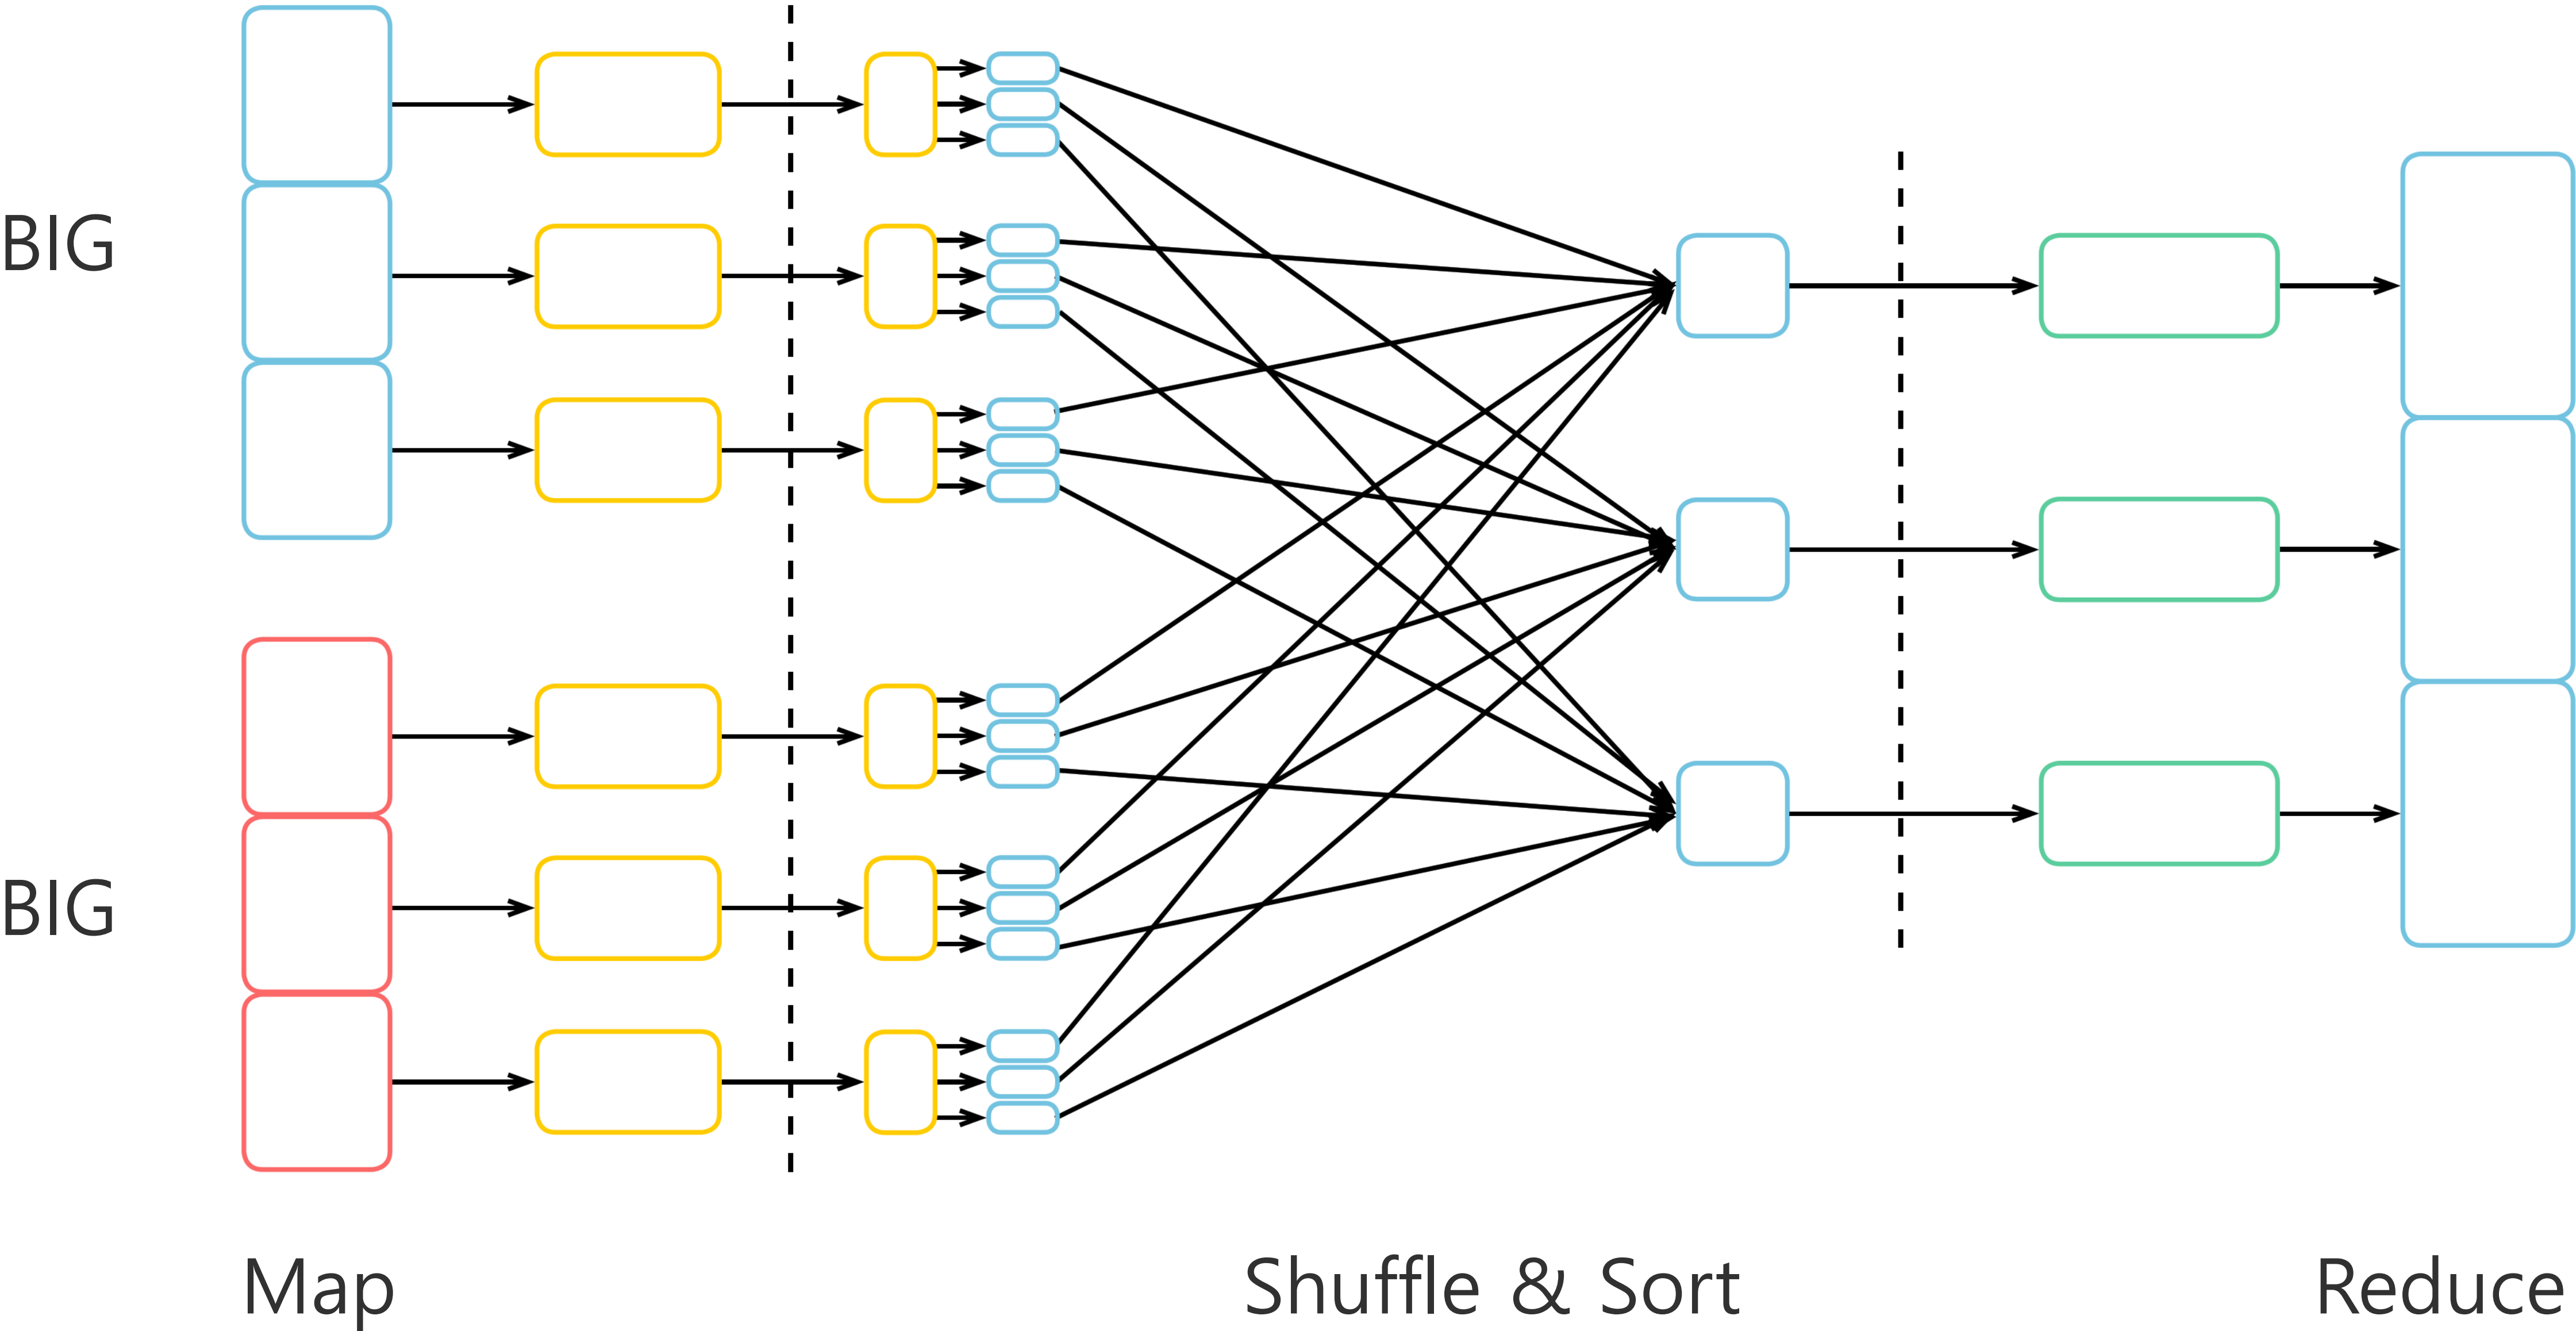
```
1 1383260400000 0 0.08136262351125882
1 1383260400000 39 0.14186425470242922
0.1567870050390246 0.16093793691701822
0.052274848528573205 11.028366381681026
1 1383261000000 0 0.13658782275823106
0.02730046487718618
1 1383261000000 33
0.026137424264286602
...
```

```
{'type': 'Polygon', 'coordinates':
[[[9.0114910478323, 45.35880131440966],
[9.014491488013135, 45.35880097314403],
[9.0144909480813, 45.35668565341486],
[9.011490619692509,
45.356685994655464], [9.0114910478323,
45.35880131440966]]]]}
...
```

Reduce-Side Join

Reducer에서 데이터를 join하는 경우

Mapper에서 선행 작업이 필요하다.



reduce_side_mapper.py

```
if "geojson" in os.environ["mapreduce_map_input_file"]:
    geojson = json.load(sys.stdin)
    grid = load_grid(geojson)
    for grid_id, ce11_type in grid.items():
        print(grid_id, "grid", ce11_type, sep="\t")
else:
    for line in sys.stdin:
        square_id, aggregate = line.split("\t", 1)
        square_id = int(square_id)
        time_interval, country, sms_in, sms_out, call_in, call_out, internet = aggregate.split("\t")
        if sms_in:
            sms_in = float(sms_in)
            print(square_id, "logs", sms_in, sep="\t")
```

os.environ["mapreduce_map_input_file"] : input 파일명

mapper에서 log 데이터와 geojson 데이터 2종류를 받는다.

구분해서 처리하기 위해 파일명으로 확인하고 output으로 grid 또는 logs 텍스트를 추가

Mapper만 먼저 확인

참고 : gejson이 BIG 데이터일 경우, distributed cache를 활용하지 않고 input으로 입력한다.

```
yarn jar $HADOOP_STREAMING_JAR ₩  
  -files reduce_side_mapper.py ₩  
  -mapper 'python reduce_side_mapper.py' ₩  
  -numReduceTasks 0 ₩  
  -input /data/telecommunication, /user/adral/geojson ₩  
  -output telecom-joins
```

```
$ hdfs dfs -text telecom-joins/part-00010 | head -3
```

```
1 grid South ← string  
2 grid South  
3 grid South
```

```
$ hdfs dfs -text telecom-joins/part-00000 | head -3
```

```
1 logs 0.0813626235113 ← numeric  
1 logs 0.141864254702  
1 logs 0.136587822758
```


partitioner

```
yarn jar $HADOOP_STREAMING_JAR \  
-D mapreduce.partition.keypartitioner.options="-k1,1" \  
-D stream.num.map.output.key.fields=2 \  
-files reduce_side_mapper_slice.py \  
-mapper 'python reduce_side_mapper_slice.py' \  
-numReduceTasks 5 \  
-input /data/telecommunication,/user/adral/geojson \  
-output telecom-joins \  
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
```

secondary sort를 위해 key fields를 2로 설정

comparator

100	grid	South
100	logs	0.00422994505598
1002	grid	South
1002	logs	0.0241862339965
1007	grid	South
1007	logs	0.0145776778024
1011	grid	South
1011	logs	0.0627696965595
1016	grid	South
1016	logs	0.0123509364406

grid 다음 logs 데이터가 나오는 형태로 정렬
(reducer 코딩을 위해서)

Secondary Sort

reduce_side_reducer.py

grid를 먼저 읽고 logs의 value를 더한다.

```
from __future__ import print_function
import sys
```

```
current_grid = None
grid_load = 0
grid_location = None
```

```
for line in sys.stdin:
    grid_id, label, value = line.strip("\n").split("\t", 2)
    if label == "grid":
        if current_grid:
            print(current_grid, grid_location, grid_load, sep="\t")
            current_grid = grid_id
            grid_load = 0
            grid_location = value
        else:
            counts = float(value)
            grid_load += counts
```

```
if current_grid != None:
    print(current_grid, grid_location, grid_load, sep="\t")
```



```
yarn jar $HADOOP_STREAMING_JAR ₩
-D mapreduce.partition.keypartitioner.options="-k1,1" ₩
-D stream.num.map.output.key.fields=2 ₩
-files reduce_side_mapper.py, reduce_side_reducer.py ₩
-mapper 'python reduce_side_mapper.py' ₩
-reducer 'python reduce_side_reducer.py' ₩
-numReduceTasks 5 ₩
-input /data/telecommunication,/user/adral/geojson ₩
-output telecom-joins ₩
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
```

100	South	55.723185988
1002	South	26.6296384356
1007	South	25.216401618
1011	South	33.9120375534
1016	South	29.0697003186
1020	South	27.635637365
1025	South	21.7622321062
1034	South	93.4964559323
1039	South	106.557543309
1043	South	130.640809358

Summary

- » you know how and when to use Map-Side Join
- » you know how and when to use Reduce-Side Join
- » you know how and when to use Secondary Sort

Telecommunications Analytics

Job Chaining

KeyFieldSelection 필요한 데이터만 Key, Value를 간단히 추출하는 법

-D mapreduce.fieldsel.map.output.key.value.fields.spec=

-mapper org.apache.Hadoop.mapred.lib.FieldSelectionMapReduce

```
yarn jar $HADOOP_STREAMING_JAR \  
-D mapreduce.fieldsel.map.output.key.value.fields.spec=0:3,5- \  
-mapper org.apache.hadoop.mapred.lib.FieldSelectionMapReduce \  
-numReduceTasks 0 \  
-input sms-call-internet-mi-2013-11-01.txt \  
-output telecom-joins
```

key index: 0 (Square ID)

```
yarn jar $HADOOP_STREAMING_JAR \  
-D mapreduce.fieldsel.map.output.key.value.fields.spec=0:3,5- \  
-mapper org.apache.hadoop.mapred.lib.FieldSelectionMapReduce \  
-numReduceTasks 0 \  
-input sms-call-internet-mi-2013-11-01.txt \  
-output telecom-joins
```

value index: 3,5-

```
$ head sms-call-internet-mi-2013-11-01.txt
```

```
11383260400000 0 0.08136262351125882  
1 1383260400000 39 0.14186425470242922 0.1567870050390246 0.16093793691701822  
0.052274848528573205 11.028366381681026  
1 1383261000000 0 0.13658782275823106 0.02730046487718618  
...
```

4번째는 지워짐

```
$ head telecom-joins/part-00000
```

```
1 0.08136262351125882  
1 0.14186425470242922 0.16093793691701822 0.052274848528573205 11.028366381681026  
1 0.13658782275823106 0.02730046487718618  
1 0.026137424264286602  
...
```

Job Chaining(파이프라인) 구성



Job Chaining 활용 프레임워크



HADOOPY



MrJob

Job Chaining(파이프라인) 구성 예시

FieldSelection

MapSide Join




Job의 상태 확인

Job이 정상적으로 종료되면 return code가 0이다.

```
yarn jar ...FieldSelectionMapReduce...  
application_return_code=$?
```

```
if [ $application_return_code != "0" ]  
then  
    echo "FieldSelection phase was NOT successful"  
    exit $application_return_code  
fi
```

Job Chaining(파이프라인) 구성 예시



Job의 상태 확인

현재 application ID 확인방법

```
$ yarn application -list
```

Queue	Application-Id	Application-Name	Application-Type	User
	State	Final-State	Progress	Tracking-URL

```
application_1491639197451_0559 / select age from (select a.Age as age,...
1000(Stage-1)      MAPREDUCE      s201701 root.users.s201701      RUNNING
UNDEFINED          90.74% http://virtual-node1.atp-fivt.org:45272
```

```
...
application_1491639197451_0565streamjob7453411855907589438.jar
MAPREDUCE      adral      root.users.adral      RUNNING      UNDEFINED
5%  http://virtual-node3.atp-fivt.org:46751
```

...

Job의 상태 확인 방법 1

\$ yarn application -status application_1491639197451_0565

Application Report :

Application-Id : application_1491639197451_0565

Application-Name : streamjob7453411855907589438.jar

Application-Type : MAPREDUCE

User : adral

Queue : root.users.adral

Start-Time : 1492272089050

Finish-Time : 1492272181541

Progress : 100%

State : FINISHED

Final-State : SUCCEEDED

Tracking-URL : http://virtual-master.atp-fivt.org:19888/jobhistory/job/job_1491639197451_0565

RPC Port : 42533

AM Host : virtual-node3.atp-fivt.org

Aggregate Resource Allocation : 544105 MB-seconds, 176 vcore-seconds

Log Aggregation Status : SUCCEEDED

Diagnostics:

Job의 상태 확인 방법 2

```
$ hdfs dfs -test -e telecom-joins/field-selection/_SUCCESS
```

-e : exits

OUTPUT 폴더의 /_SUCCESS file은 job이 잘 끝났음을 의미한다.

```
$ hdfs dfs -ls telecom-joins/field-selection
```

Found 11 items


```
-rw-r--r-- 3 ...      0 2017-04-15 18:30 telecom-joins/field-selection/_SUCCESS
-rw-r--r-- 3 ... 8516829 2017-04-15 18:28 telecom-joins/field-selection/part-00000
-rw-r--r-- 3 ... 8555052 2017-04-15 18:29 telecom-joins/field-selection/part-00001
-rw-r--r-- 3 ... 8380546 2017-04-15 18:29 telecom-joins/field-selection/part-00002
-rw-r--r-- 3 ... 8546802 2017-04-15 18:29 telecom-joins/field-selection/part-00003
-rw-r--r-- 3 ... 8293343 2017-04-15 18:29 telecom-joins/field-selection/part-00004
-rw-r--r-- 3 ... 8348375 2017-04-15 18:29 telecom-joins/field-selection/part-00005
...
```

Telecommunications Analytics

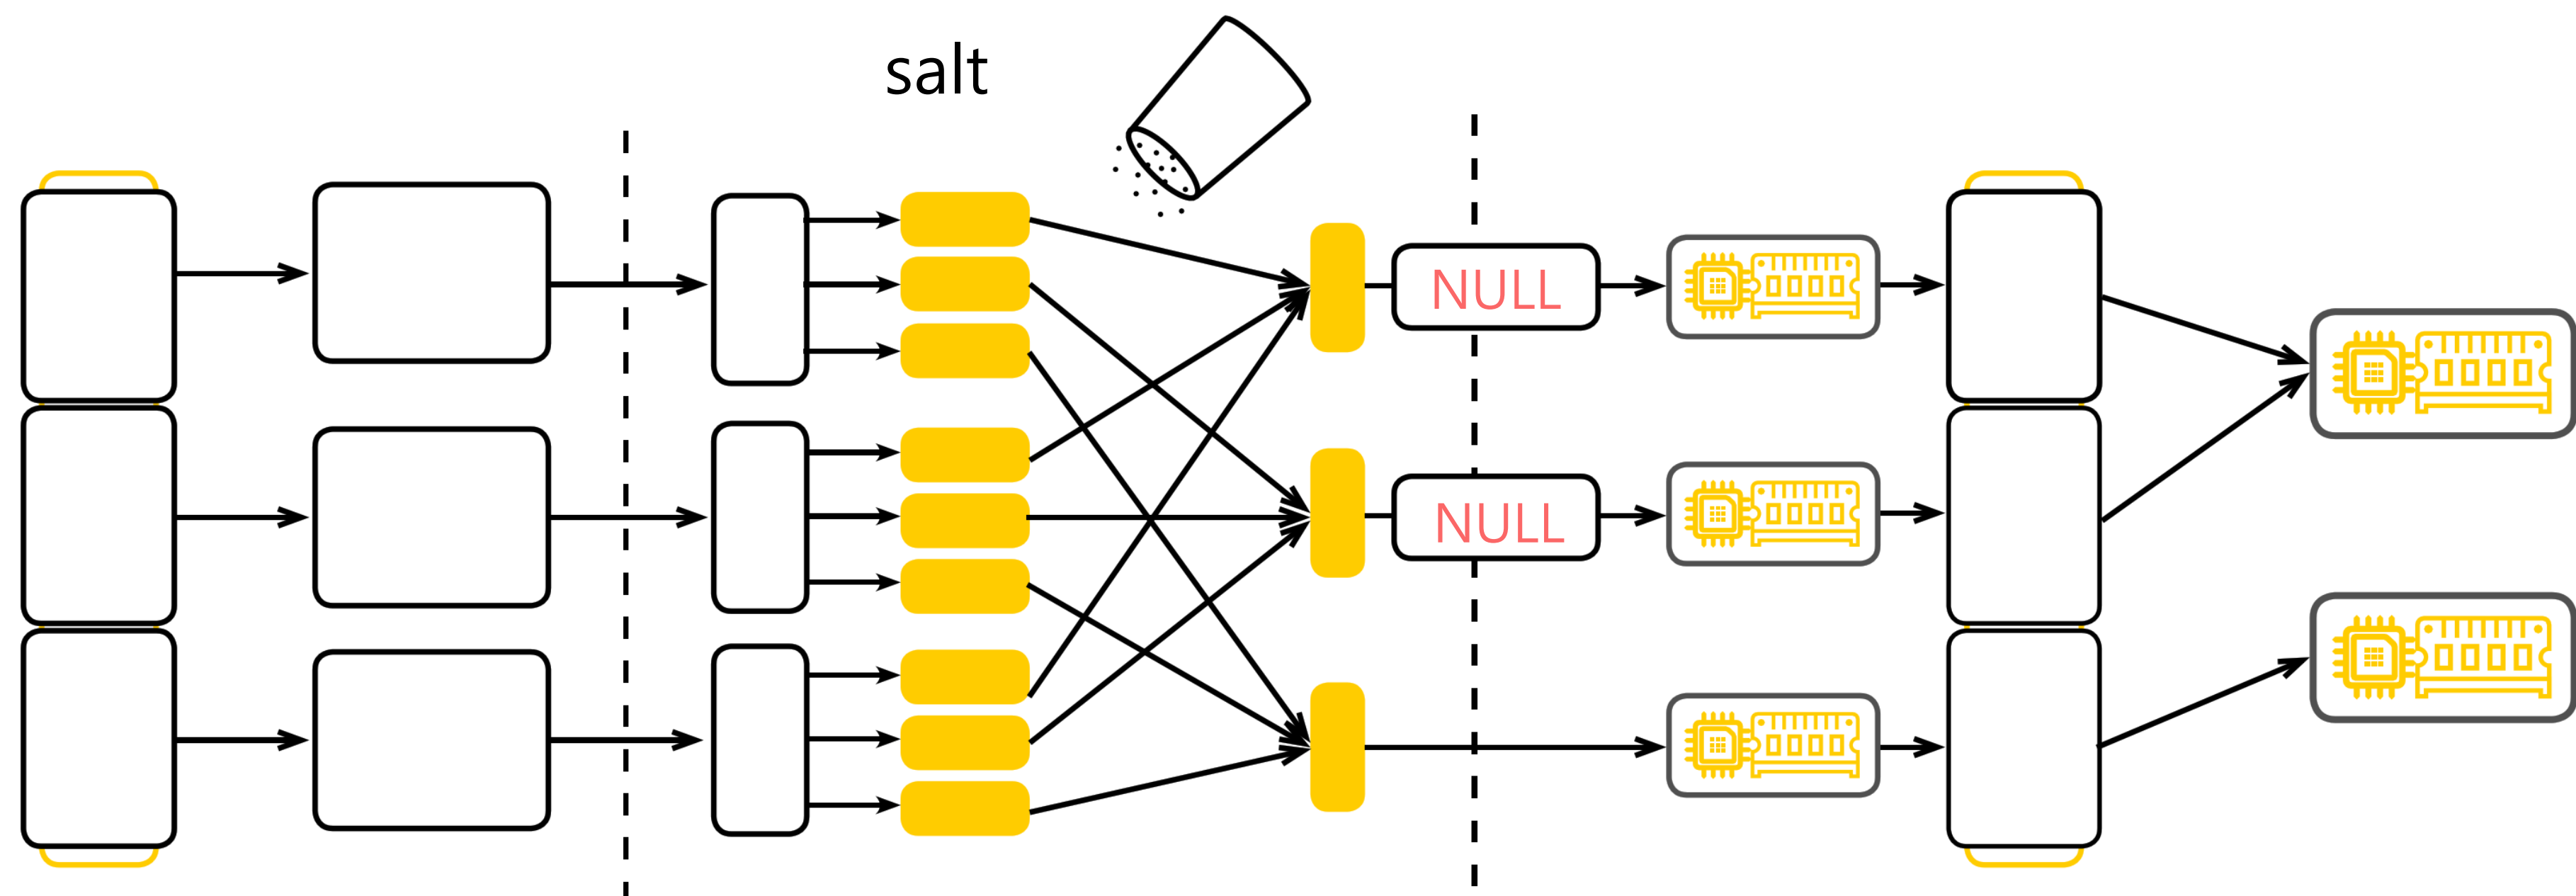
Data Skew, Salting

Data Skew란 특정 key에 대한 value가 너무 많을 때를 말한다.

=> 특정 reducer에 데이터가 많이 몰려 JOB 수행이 오래 걸린다.

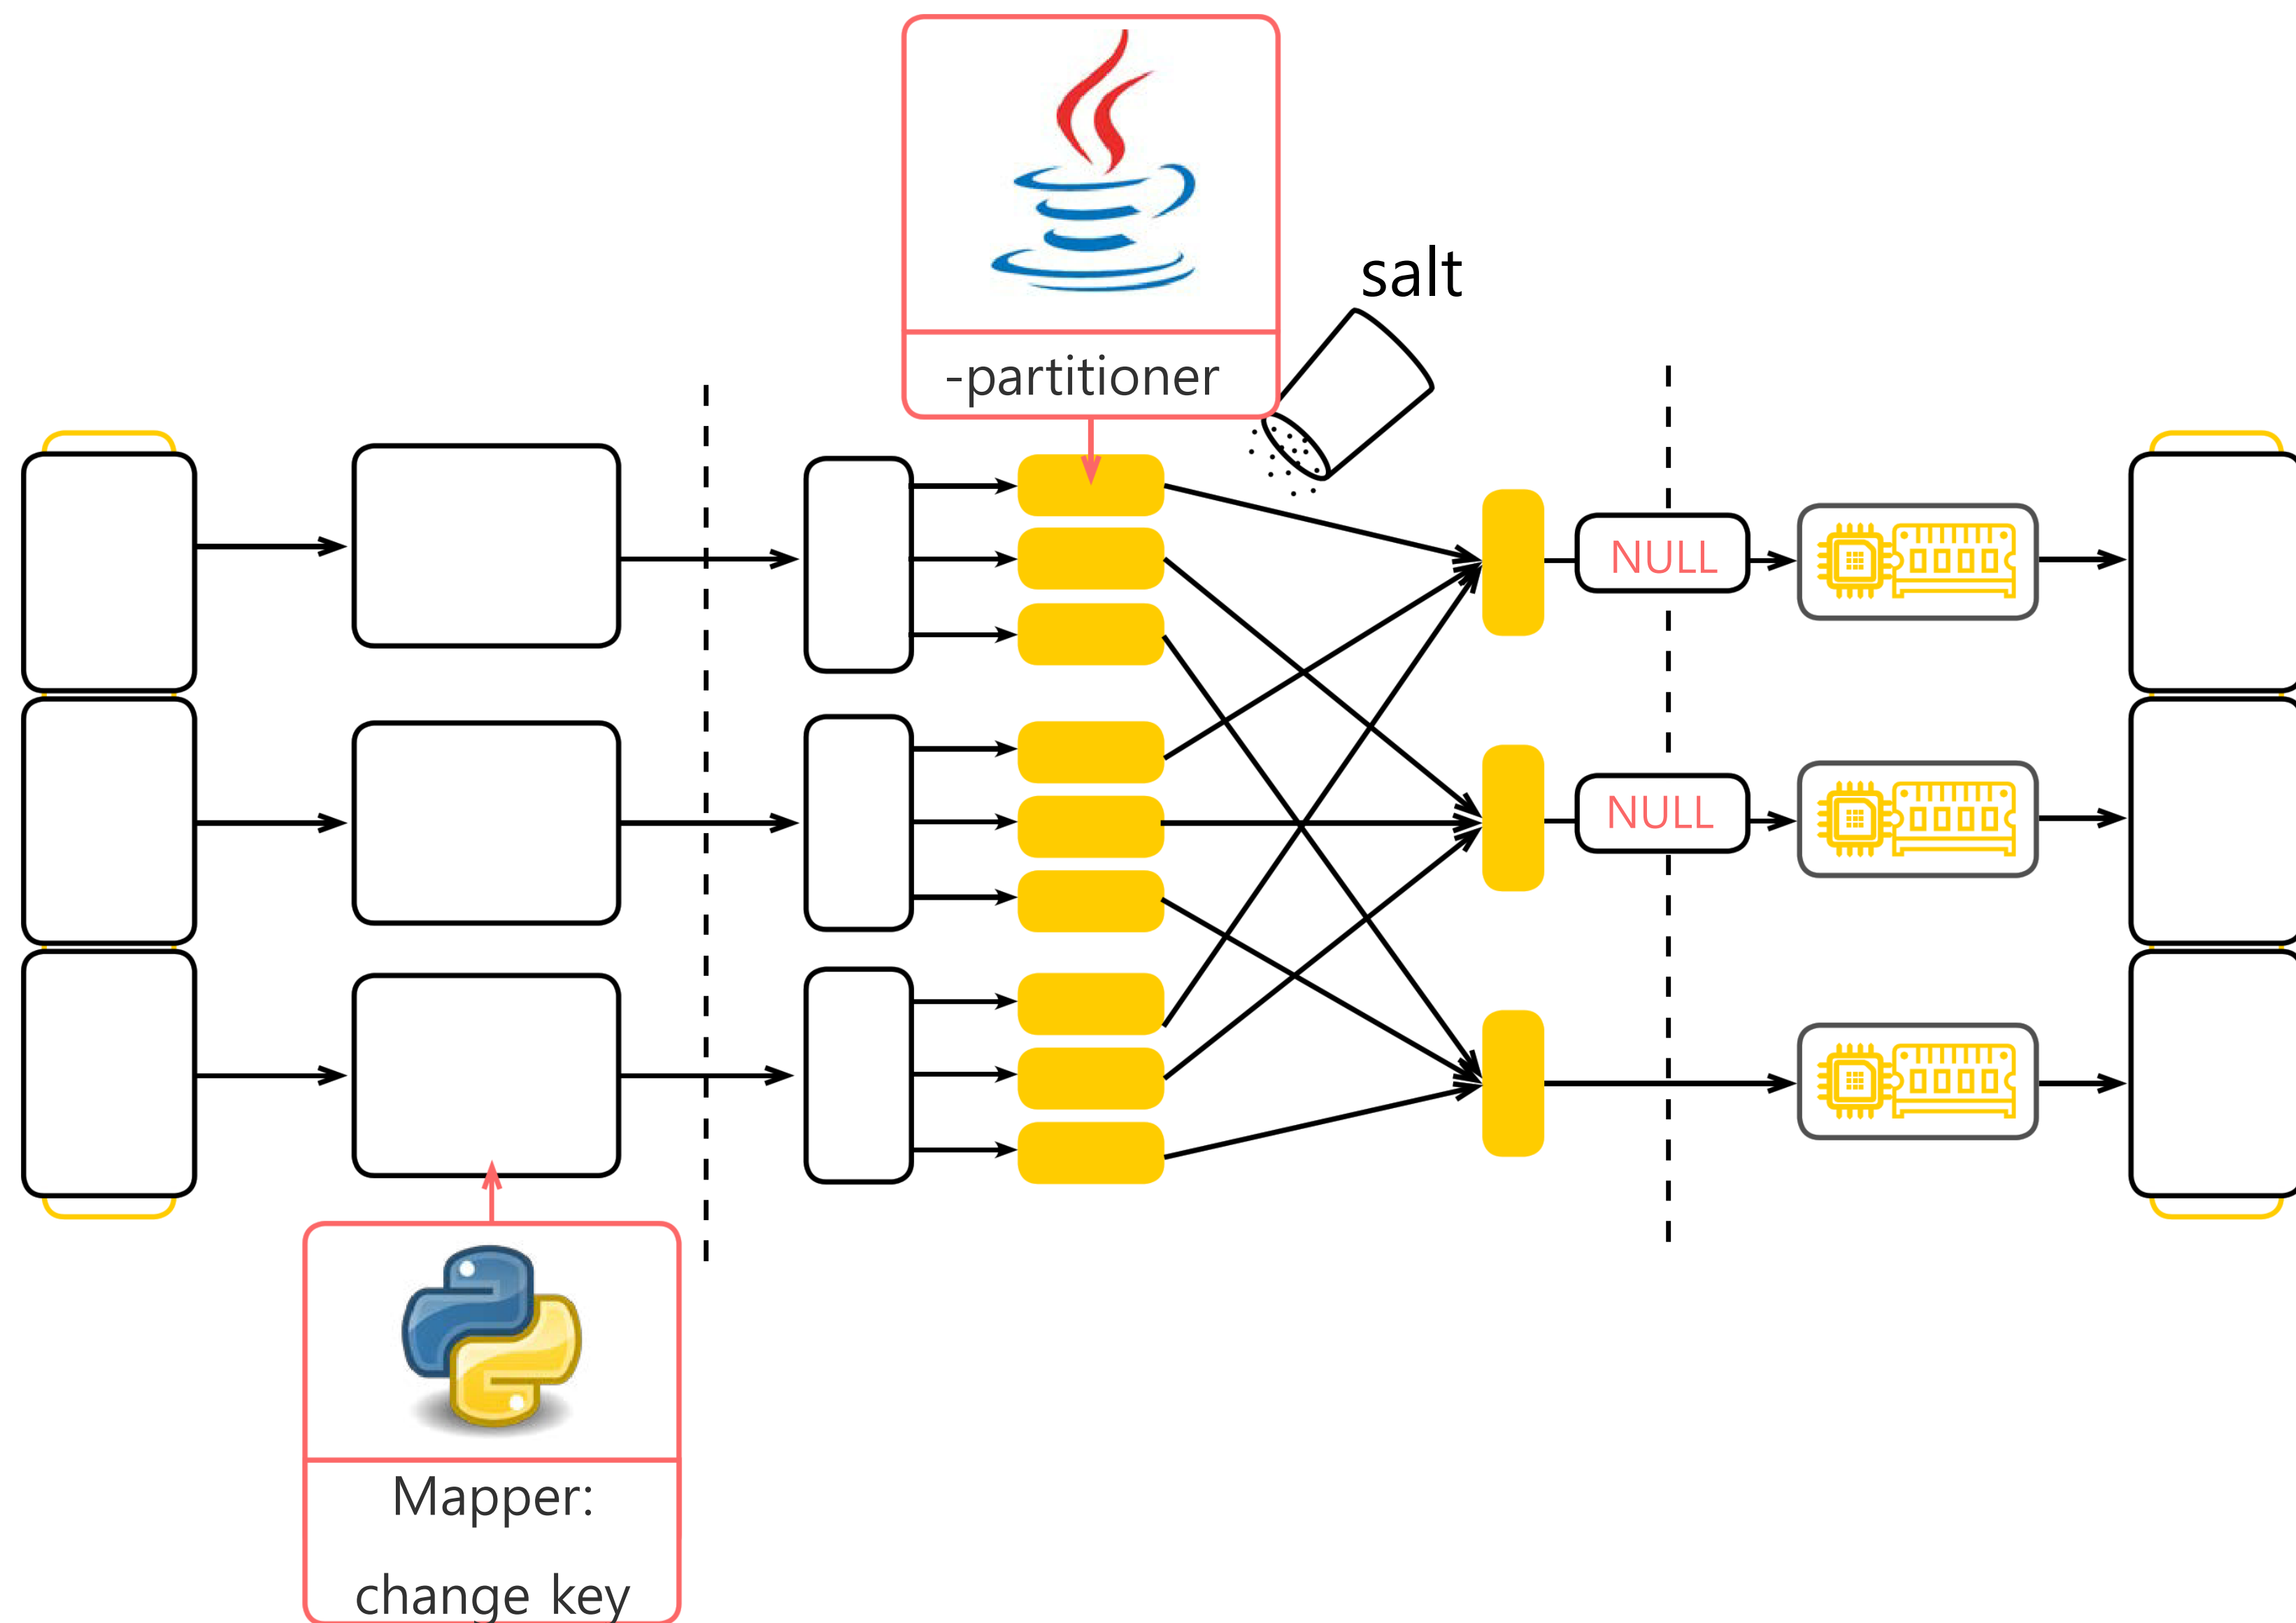


Salting은 같은 Key의 데이터를 임의로 여러 reducer로 분산시켜 작업하는 것



Salting 방법 2가지

- Java로 partitioner 코드를 수정하는 방법
- Mapper에서 Key를 임의로 수정하는 방법



Mapper

```
from random import random
geojson = json.load(open("milano-grid.geojson"))
grid = load_grid(geojson)
for line in sys.stdin:
    square_id, value_to_aggregate = line.rstrip("\n").split("\t", 1)
    square_id = int(square_id)
grid_location = "null" if random() < 0.9 else grid[square_id]
    if value_to_aggregate:
        print(grid_location, value_to_aggregate, sep="\t")
```

참고 : 테스트를 위해 10%의 데이터를 모두 Null 키를 만든다.

```
from random import randrange
grid_location = "null_{0}".format(randrange(100)) if random() < 0.9 else grid[square_id]
```

null_+임의값 : Null 데이터를 분배하기 위해 Null key에 임의의 값을 더한다.

...


null_58 40989.56529872355

null_67 40775.58025775422

null_76 42430.98650098723

...

Job Chaining을 활용해 새로운 Job의 mapper로 null key를 원래대로 만든다.



```
for line in sys.stdin:
    key, value = line.rstrip("\n").split("\t", 1)
    key = "null" if "null_" in key else key
    print("DoubleValueSum:{}".format(key), value, sep="\t")
```



-reducer aggregate

South 164302.58197312435
null 4145425.004916422
North 296659.74407499237